

# The Picnic Signature Scheme

## Design Document

Melissa Chase, David Derler, Steven Goldfeder,  
Daniel Kales, Jonathan Katz, Vladimir Kolesnikov, Claudio Orlandi,  
Sebastian Ramacher, Christian Rechberger,  
Daniel Slamanig, Xiao Wang, Greg Zaverucha

September 30, 2020  
Version 3.0

# Contents

<b>1</b>	<b>Introduction</b>	<b>4</b>
1.1	The Picnic Design Team . . . . .	4
1.2	Acknowledgments . . . . .	5
<b>2</b>	<b>Background</b>	<b>5</b>
2.1	Commitments . . . . .	5
2.2	Zero-Knowledge Proofs and $\Sigma$ -Protocols . . . . .	6
2.3	Non-Interactive Zero-Knowledge Proofs of Knowledge . . . . .	7
2.4	Signature Schemes . . . . .	9
2.5	Fiat-Shamir Transform . . . . .	10
2.6	Unruh Transform . . . . .	10
2.7	(2,3)-Decomposition of Circuits . . . . .	11
2.8	ZKB++ . . . . .	14
2.9	KKW . . . . .	21
2.9.1	The MPC Protocol . . . . .	21
2.9.2	The Proof Protocol . . . . .	22
2.9.3	Further Optimizations for Picnic3 . . . . .	26
2.10	LowMC . . . . .	30
<b>3</b>	<b>The Picnic Signature Schemes</b>	<b>31</b>
3.1	Efficient Instantiation of Unruh's Transform . . . . .	32
3.2	Seed Generation . . . . .	34
3.3	Random Tapes . . . . .	34
3.4	Challenge Generation . . . . .	34
3.5	Function $G$ . . . . .	34
<b>4</b>	<b>Choice of Parameters</b>	<b>34</b>
4.1	Choice of LowMC and SHAKE . . . . .	35
4.2	LowMC Parameters . . . . .	37
4.3	MPC Parameters . . . . .	39
4.4	Alternative Parameters . . . . .	40
4.5	Recommended Parameters . . . . .	41
<b>5</b>	<b>Formal Security Analysis</b>	<b>42</b>
5.1	Security Analysis of ZKB++ . . . . .	42
5.2	Security Analysis of Picnic-FS . . . . .	44

5.3	Security Analysis of Picnic-UR . . . . .	46
5.4	Strong Unforgeability of Picnic-FS and Picnic-UR . . . . .	54
<b>6</b>	<b>Formal Security Analysis of Picnic3</b>	<b>55</b>
6.1	Proof of Security of the Underlying MPC Protocol . . . . .	55
6.2	Security Proof of the Signature Scheme . . . . .	56
6.3	Tree-Based Optimizations . . . . .	65
6.3.1	Seed Tree . . . . .	66
6.3.2	Use in Picnic3 . . . . .	67
6.4	QROM Security . . . . .	68
6.5	Computationally Unique Responses . . . . .	71
<b>7</b>	<b>Analysis with Respect to Known Attacks</b>	<b>73</b>
7.1	Usage and Security Margin of LowMC . . . . .	73
7.2	Attacks in the Single-User Setting . . . . .	75
7.3	Attacks in the Multi-User Setting . . . . .	75
7.4	Multi-Target Attacks . . . . .	77
<b>8</b>	<b>Expected Security Strength</b>	<b>79</b>
8.1	LowMC Parameter Selection . . . . .	81
8.2	Hash Function Security . . . . .	81
<b>9</b>	<b>Advantages and Limitations</b>	<b>82</b>
9.1	Compatibility with Existing Protocols . . . . .	82
9.2	TLS and X.509 Compatibility . . . . .	83
9.3	Hardware Security Module Compatibility . . . . .	83
<b>10</b>	<b>Additional Security Properties</b>	<b>87</b>
10.1	Side-Channel Attacks . . . . .	87
10.2	Security Impact of Using Weak Ephemeral Values . . . . .	89
10.3	Parameter Integrity . . . . .	90
<b>11</b>	<b>Efficiency and Memory Usage</b>	<b>91</b>
11.1	Description of the Benchmark Platforms . . . . .	91
11.1.1	Platform A . . . . .	91
11.1.2	Platform B . . . . .	92
11.2	Description of the Benchmarking Methodology . . . . .	92
11.3	Benchmark Results: Sizes . . . . .	92
11.4	Benchmark Results: Timings . . . . .	93

11.5	Memory Requirements . . . . .	100
11.5.1	Reference Implementation Detailed Memory Usage . .	100
11.5.2	Optimized Implementation Detailed Memory Usage . .	100
11.6	Size of Precomputed Constants and Data . . . . .	100
<b>A</b>	<b>Change History</b>	<b>112</b>

# 1 Introduction

Picnic is a signature scheme that is designed to provide security against attacks by quantum computers, in addition to attacks by classical computers. The scheme uses a zero-knowledge proof system and is based on symmetric key primitives like hash functions and block ciphers with conjectured post-quantum security. In particular, Picnic does not rely on number-theoretic or algebraic hardness assumptions.

In this document we present the building blocks of the Picnic signature scheme in Section 2. In Section 3 we present several variants of the Picnic signature scheme and various optimizations to the building blocks. We specify the parameters for some of the building blocks in Section 4 (and leave complete details of the parameters to the specification document). In Section 5 we include the formal security proofs for the proposed instantiations of Picnic. Section 7 presents an analysis of the algorithm with respect to known attacks and Section 8 provides a thorough description of the expected security strength. Finally, Section 9 discuss advantages and limitations, Section 10 discusses additional security properties, and Section 11 presents results on efficiency and memory usage of the Picnic scheme.

**Source Materials** Parts of this document are taken or adapted from research papers [CDG<sup>+</sup>17, KKW18, DKP<sup>+</sup>19a, KZ20b] by the authors and the Picnic Specification document [Tea19a]. The Picnic website [Tea19b] lists these sources and related talks, with links to the papers themselves.

## 1.1 The Picnic Design Team

Picnic was designed collaboratively by the following group of people.

Melissa Chase, Microsoft  
David Derler, DFINITY  
Steven Goldfeder, Cornell Tech  
Jonathan Katz, George Mason University  
Daniel Kales, Graz University of Technology  
Vladimir Kolesnikov, Georgia Tech  
Claudio Orlandi, Aarhus University  
Sebastian Ramacher, Graz University of Technology  
Christian Rechberger, Graz University of Technology & DTU  
Daniel Slamanig, AIT Austrian Institute of Technology

Xiao Wang, Northwestern University  
Greg Zaverucha, Microsoft

## 1.2 Acknowledgments

We are grateful for help from Christian Paquin (integration of Picnic in OQS and OpenSSL), Larry Joy (HSM demo), Alexander Grass and Angela Promitzer (contributions to the optimized implementation). We are also grateful to Serge Fehr and his co-authors Jelle Don, Christian Majenz and Christian Schaffner, for generalizing their QROM results [DFMS19a] to apply to Picnic2/Picnic3.

## 2 Background

In this section we review some background material relevant to the Picnic design.

### 2.1 Commitments

**Definition 2.1** (Commitment Scheme). A (non-interactive) commitment scheme consists of algorithms **Com** and **Open** with the following properties:

**Com**( $M$ ) : On input a message  $M \in \{0, 1\}^*$ , the commitment algorithm outputs  $(C, D) \leftarrow \text{Com}(M; R)$ , where  $R$  is a random value used when forming the commitment.  $C$  is the commitment string, while  $D$  is the decommitment string which is kept secret until opening time.

**Open**( $C, D$ ) : On input  $C, D$ , the verification algorithm either outputs a message  $M$  or  $\perp$ .

Computationally secure commitments must satisfy the following properties

**Correctness.** If **Com**( $M$ ) outputs  $(C, D)$  then **Open**( $C, D$ ) =  $M$ .

**Hiding.** For every message pair  $M, M'$  the probability ensembles  $\{C : (C, D) \leftarrow \text{Com}(M)\}_{\kappa \in \mathbb{N}}$  and  $\{C : (C, D) \leftarrow \text{Com}(M')\}_{\kappa \in \mathbb{N}}$  are computationally indistinguishable for security parameter  $\kappa$ .

**Binding.** We say that an adversary  $\mathcal{A}$  wins if it outputs  $C, D, D'$  such that  $\text{Open}(C, D) = M$ ,  $\text{Open}(C, D') = M'$  and  $M \neq M'$ . We require that for all efficient algorithms  $\mathcal{A}$  (running in time polynomial in  $\kappa$ ), the probability that  $\mathcal{A}$  wins is a negligible function of  $\kappa$ .

Our implementation uses hash-based commitments, which requires modeling the hash function as a random oracle in our security analysis. Let  $H$  be a cryptographic hash function. The commitment scheme works as follows:

$\text{Com}(M)$  : Sample  $R \leftarrow^R \{0, 1\}^\kappa$  and set  $C \leftarrow H(R, M)$  and return  $(C, (R, M))$ ;

$\text{Open}(C, D)$  : Parse  $D$  as  $(R, M)$  and return  $M$  if  $H(R, M) = C$ , and return  $\perp$  otherwise.

## 2.2 Zero-Knowledge Proofs and $\Sigma$ -Protocols

A sigma protocol is a three-flow protocol between a prover and verifier, used to prove knowledge of a secret. A well-known class of sigma protocols are the so-called generalized Schnorr proofs, which allow the prover to prove knowledge of a discrete logarithm, and that it satisfies certain properties. In the present work we use a sigma protocol that allows one to prove knowledge of an input to an arbitrary binary circuit. Sigma protocols are usually zero-knowledge proofs, which informally means that the proof protocol does not reveal any information about the secret. We describe interactive protocols, but will show later how to make them non-interactive (so that signatures are non-interactive). Let  $L$  be an **NP**-language with associated witness relation  $R$  so that  $L = \{x \mid \exists w : R(x, w) = 1\}$ . A  $\Sigma$ -protocol for language  $L$  is defined as follows.

**Definition 2.2** ( $\Sigma$ -Protocol). A  $\Sigma$ -protocol for relation  $R$  is an interactive three-move protocol between a PPT prover  $P = (\text{Commit}, \text{Prove})$  and a PPT verifier  $V = (\text{Challenge}, \text{Verify})$ , where  $P$  makes the first move and transcripts are of the form  $(a, e, z) \in A \times E \times Z$ , where  $a$  is output by **Commit**,  $e$  is output by **Challenge** and  $z$  is output by **Prove**. Additionally,  $\Sigma$  protocols satisfy the following properties

**Completeness.** A  $\Sigma$ -protocol for language  $L$  is complete, if for all security parameters  $\kappa$ , and for all  $(x, w) \in R$ , it holds that

$$\Pr[\langle P(1^\kappa, x, w), V(1^\kappa, x) \rangle = 1] = 1.$$

**$s$ -Special Soundness.** A  $\Sigma$ -protocol for language  $L$  is  $s$ -special sound, if there exists a PPT extractor  $E$  so that for all  $x$ , and for all sets of accepting transcripts  $\{(\mathbf{a}, \mathbf{e}_i, \mathbf{z}_i)\}_{i \in [s]}$  with respect to  $x$  where  $\forall i, j \in [s], i \neq j : \mathbf{e}_i \neq \mathbf{e}_j$ , generated by any algorithm with polynomial runtime in  $\kappa$ , it holds that

$$\Pr \left[ w \leftarrow E(1^\kappa, x, \{(\mathbf{a}, \mathbf{e}_i, \mathbf{z}_i)\}_{i \in [s]}) \quad : \quad (x, w) \in R \right] \geq 1 - \epsilon(\kappa).$$

We can also consider a computational variant which says that if  $\mathcal{A}$  is a PPT algorithm then the probability that it can produce an accepting transcript from which  $E$  fails to extract a valid witness is negligible.

**Special Honest-Verifier Zero-Knowledge.** A  $\Sigma$ -protocol is special honest-verifier zero-knowledge, if there exists a PPT simulator  $S$  so that for every  $x \in L$  and every challenge  $\mathbf{e}$  from the challenge space, it holds that a transcript  $(\mathbf{a}, \mathbf{e}, \mathbf{z})$ , where  $(\mathbf{a}, \mathbf{z}) \leftarrow S(1^\kappa, x, \mathbf{e})$  is computationally indistinguishable from a transcript resulting from an honest execution of the protocol.

The  $s$ -special soundness property gives an immediate bound for the soundness of the protocol: if no witness exists then (ignoring a negligible error) the prover can successfully answer at most  $(s - 1)/t$  of the possible challenges, where  $t = |\mathbf{E}|$  is the size of the challenge space. If this value is too large, it is possible to reduce the soundness error using parallel repetition (see [Dam10, CDS94] for details).

**Lemma 2.3.** *Let  $\Pi$  be a  $\Sigma$ -protocol with  $s$ -special soundness. Then the  $\ell$ -fold parallel repetition of  $\Pi$ , denoted  $\Pi^\ell$ , is also a  $\Sigma$ -protocol with  $s^\ell$ -special soundness.*

## 2.3 Non-Interactive Zero-Knowledge Proofs of Knowledge

For our signatures, we use non-interactive zero-knowledge proofs of knowledge, in which the proof is a single message. Here we define zero-knowledge and the necessary notion of simulation-extractability. We present the definition the random oracle model against classical adversaries and the definition in the quantum random oracle model against quantum adversaries.<sup>1</sup>

---

<sup>1</sup>These definition roughly combine the ROM definitions of [BPW12] and the QROM definitions of [Unr15].



Zero-knowledge says that there is a simulator which can produce proofs that are indistinguishable from those produced by  $P$  without knowing the witnesses to an adversary who is given access to a simulated version of the random oracle.

**Definition 2.4** (Zero-Knowledge). A protocol  $P, V$  for relation  $R$  is zero-knowledge against (quantum) adversaries in the (quantum) random oracle model if there exist PPT algorithms  $\mathbf{S}_R, \mathbf{Sim}$  such that for all (quantum) PPT adversaries  $\mathcal{A}$

$$|\Pr[b \leftarrow \mathcal{A}^{R(\cdot), P(\cdot, \cdot)}(1^\lambda) : b = 1] - \Pr[b \leftarrow \mathcal{A}^{\mathbf{S}_R(\cdot), \mathbf{Sim}_P(\cdot, \cdot)}(1^\lambda) : b = 1]|$$

is negligible in  $\lambda$ , where  $R$  is a random function to which  $\mathcal{A}$  can provide (quantum) inputs,  $\mathbf{Sim}_P$  takes a pair  $(x, w) \in R$  and calls  $\mathbf{Sim}(x)$ , and  $\mathbf{Sim}, \mathbf{S}_R$  share state.

Simulation-extractability says that an adversary cannot produce a new proof for a statement for which he does not know the witness, even if he is allowed to see proofs produced by someone else for statements of his choice. More formally, we say that even when an adversary is given access to the zero-knowledge simulator to obtain proofs for (true or false) statements of its choice, whenever it produces a new proof (not produced by the simulator) that is accepted by the verifier, there is an extractor that can look at the implementation of the adversary and extract a valid witness for that statement.

**Definition 2.5** (Simulation extractability). A protocol  $P, V$  for relation  $R$  satisfies simulation extractability against (quantum) adversaries in the (quantum) random oracle model if there exist PPT algorithms  $\mathbf{S}_R, \mathbf{Sim}$  satisfying the zero-knowledge definition and a PPT (quantum) extractor such that for all (quantum) PPT adversaries  $\mathcal{A}$

$$\Pr[(x, \pi) \leftarrow \mathcal{A}^{\mathbf{S}_R(\cdot), \mathbf{Sim}(\cdot)}(1^\lambda); w \leftarrow E(\mathcal{A}, x, \pi) : \\ V^{\mathbf{S}_R}(x, \pi) = 1 \wedge (x, \pi) \notin \mathcal{Q} \wedge (x, w) \notin R]$$

is negligible in  $\lambda$ , where  $\mathbf{S}_R, \mathbf{Sim}$ , and  $E$  share state,  $\mathcal{Q}$  is the list of  $\mathcal{A}$ 's queries to  $\mathbf{Sim}$  and the resulting responses, and passing  $\mathcal{A}$  as input to  $E$  means  $E$  is given access to a (quantum) implementation of  $\mathcal{A}$ .

## 2.4 Signature Schemes

In the following we recall a standard definition of signature schemes along with two widely used security notions.

**Definition 2.6** (Signature Scheme). A signature scheme  $\Sigma$  is a triple  $(\text{Gen}, \text{Sign}, \text{Verify})$  of PPT algorithms, which are defined as follows:

$\text{Gen}(1^\kappa)$  : This algorithm takes a security parameter  $\kappa$  as input and outputs a secret (signing) key  $\text{sk}$  and a public (verification) key  $\text{pk}$  with associated message space  $\mathcal{M}$  (we may omit to make the message space  $\mathcal{M}$  explicit).

$\text{Sign}(\text{sk}, m)$  : This algorithm takes a secret key  $\text{sk}$  and a message  $m \in \mathcal{M}$  as input and outputs a signature  $\sigma$ .

$\text{Verify}(\text{pk}, m, \sigma)$  : This algorithm takes a public key  $\text{pk}$ , a message  $m \in \mathcal{M}$  and a signature  $\sigma$  as input and outputs a bit  $b \in \{0, 1\}$ .

Besides the usual correctness property,  $\Sigma$  needs to provide some unforgeability notion. We consider two notions, namely existential unforgeability under adaptively chosen message attacks (EUF-CMA security) and its strong variant (sEUF-CMA security), which we define below.

**Definition 2.7** (EUF-CMA). A signature scheme  $\Sigma$  is EUF-CMA secure, if for all PPT adversaries  $\mathcal{A}$  there is a negligible function  $\varepsilon(\cdot)$  such that

$$\Pr \left[ (\text{sk}, \text{pk}) \leftarrow \text{Gen}(1^\kappa), (m^*, \sigma^*) \leftarrow \mathcal{A}^{\text{Sign}(\text{sk}, \cdot)}(\text{pk}) : \right. \\ \left. \text{Verify}(\text{pk}, m^*, \sigma^*) = 1 \wedge (m^*, \cdot) \notin \mathcal{Q}^{\text{Sign}} \right] \leq \varepsilon(\kappa),$$

where the environment keeps track of the queries and responses to and from the signing oracle via  $\mathcal{Q}^{\text{Sign}}$ .

**Definition 2.8** (sEUF-CMA). A signature scheme  $\Sigma$  is strongly EUF-CMA (sEUF-CMA) secure, if for all PPT adversaries  $\mathcal{A}$  there is a negligible function  $\varepsilon(\cdot)$  such that

$$\Pr \left[ (\text{sk}, \text{pk}) \leftarrow \text{Gen}(1^\kappa), (m^*, \sigma^*) \leftarrow \mathcal{A}^{\text{Sign}(\text{sk}, \cdot)}(\text{pk}) : \right. \\ \left. \text{Verify}(\text{pk}, m^*, \sigma^*) = 1 \wedge (m^*, \sigma^*) \notin \mathcal{Q}^{\text{Sign}} \right] \leq \varepsilon(\kappa),$$

where the environment keeps track of the queries and responses to and from the signing oracle via  $\mathcal{Q}^{\text{Sign}}$ .

## 2.5 Fiat-Shamir Transform

The Fiat-Shamir (FS) transform [FS86] is an elegant way to construct signature schemes from  $\Sigma$ -protocols. The basic idea is similar to constructing NIZK proofs from  $\Sigma$ -protocols, but the challenge  $e$  is generated by hashing the prover's first message  $\mathbf{a}$  and the message  $m$  to be signed, i.e., define a modified challenge algorithm **Challenge'** that outputs  $e \leftarrow H(\mathbf{a}, m)$ . Then, the prover can locally obtain the challenge after computing the initial message. Starting a verifier  $\mathbf{V}' = (\mathbf{Challenge}', \mathbf{Verify})$  on the same initial message would yield the same challenge. The prover outputs  $(\mathbf{a}, \mathbf{z})$  as the challenge.

More formally, by using the hash function  $H : \mathbf{A} \times \mathbf{X} \rightarrow \mathbf{E}$ , which we model as a random oracle, we obtain the non-interactive PPT algorithms  $(\mathbf{Prove}_H, \mathbf{Verify}_H)$ , defined as follows:

**Prove** $_H(1^\kappa, (x, m), w)$  : Start  $\mathbf{P}$  on input  $(1^\kappa, x, w)$ , obtain the first message  $\mathbf{a}$ , answer with  $\mathbf{e} \leftarrow H(\mathbf{a}, m)$ , and finally obtain  $\mathbf{z}$ . Return  $\pi \leftarrow (\mathbf{a}, \mathbf{z})$ .

**Verify** $_H(1^\kappa, (x, m), \pi)$  : Parse  $\pi$  as  $(\mathbf{a}, \mathbf{z})$ . Start  $\mathbf{V}'$  on  $(1^\kappa, x)$ , send  $(\mathbf{a}, m)$  as the first message to the verifier. When  $\mathbf{V}'$  outputs  $\mathbf{e}$ , reply with  $\mathbf{z}$  and output 1 if  $\mathbf{V}'$  accepts and 0 otherwise.

## 2.6 Unruh Transform

Similar to the Fiat-Shamir transform, Unruh's transform [Unr12, Unr15, Unr16] allows one to construct NIZK proofs and signature schemes from  $\Sigma$ -protocols. In contrast to the FS transform, Unruh's transform can be proven secure in the QROM (quantum random oracle model), strengthening the security guarantee against quantum adversaries.

At a high level, Unruh's transform works as follows: Given a 2-special-sound  $\Sigma$ -protocol, integers  $t$  and  $M$ , a statement  $x$  and a random permutation  $G$ , the prover will repeat the first phase of the  $\Sigma$ -protocol  $t$  times. Then, for each of the  $t$  runs, it produces proofs to  $M$  different randomly selected challenges. The prover applies  $G$  to each of the so-obtained responses. The prover then selects the responses to publish for each round of the  $\Sigma$ -protocol by querying the random oracle on the message to be signed, all first rounds of the  $\Sigma$ -protocol and the outputs of  $G$  on all responses.

For a more formal treatment, we keep  $t$  as above and let  $M \in [2, |\mathbf{E}|]$ . Let  $H : \{0, 1\}^* \rightarrow [M]^t$  be a hash function we model as a random oracle.

We obtain the non-interactive PPT algorithms  $(\text{Prove}_H, \text{Verify}_H)$  defined as follows:

$\text{Prove}_H(1^\kappa, (x, m), w) :$

1. For  $i \in [t]$ :
  - (a) Start  $\mathbf{P}$  on  $(1^\kappa, x, w)$  and obtain first message  $\mathbf{a}_i$ .
  - (b) For  $j \in [M]$ , set  $\mathbf{e}_{i,j} \xleftarrow{R} \mathbf{E} \setminus \{\mathbf{e}_{i,1}, \dots, \mathbf{e}_{i,j-1}\}$  and obtain response  $\mathbf{z}_{i,j}$  for challenge  $\mathbf{e}_{i,j}$ .
2. For  $i, j \in [t] \times [M]$ , set  $g_{i,j} \leftarrow G(\mathbf{z}_{i,j})$ .
3. Let  $(J_1, \dots, J_t) \leftarrow H(m, (\mathbf{a}_i)_{i \in [t]}, (\mathbf{e}_{i,j})_{(i,j) \in [t] \times [M]}, (g_{i,j})_{(i,j) \in [t] \times [M]})$
4. Return  $\pi \leftarrow ((\mathbf{a}_i)_{i \in [t]}, (\mathbf{e}_{i,j})_{(i,j) \in [t] \times [M]}, (g_{i,j})_{(i,j) \in [t] \times [M]}, (\mathbf{z}_{i,J_i})_{i \in [t]})$

$\text{Verify}_H(1^\kappa, (x, m), \pi) :$  Parse  $\pi$  as  $((\mathbf{a}_i)_{i \in [t]}, (\mathbf{e}_{i,j})_{(i,j) \in [t] \times [M]}, (g_{i,j})_{(i,j) \in [t] \times [M]}, (\mathbf{z}_i)_{i \in [t]})$ .

1. Let  $(J_1, \dots, J_t) \leftarrow H(m, (\mathbf{a}_i)_{i \in [t]}, (\mathbf{e}_{i,j})_{(i,j) \in [t] \times [M]}, (g_{i,j})_{(i,j) \in [t] \times [M]})$
2. For  $i \in [t]$  check that all  $\mathbf{e}_{i,1}, \dots, \mathbf{e}_{i,M}$  are pairwise distinct.
3. For  $i \in [t]$  check whether  $V$  accepts the proof with respect to  $x$ , first message  $\mathbf{a}_i$ , challenge  $\mathbf{e}_{i,J_i}$  and response  $\mathbf{z}_i$ .
4. For  $i \in [t]$  check  $g_{i,J_i} = G(\mathbf{z}_i)$ .
5. Output 1 if all checks succeeded and 0 otherwise.

We discuss a specialization of Unruh's transform to our  $\Sigma$ -protocol in Section 3.1.

## 2.7 (2,3)-Decomposition of Circuits

A circuit decomposition is a protocol for jointly computing a circuit, similar to an MPC protocol, but with greater efficiency. In a (2,3)-decomposition there are three players and the protocol has 2-privacy, i.e., it remains secure even if two of the three players are corrupted.

**Definition 2.9** ((2,3)-decomposition). Let  $f(\cdot)$  be a function that is computed by an  $n$ -gate circuit  $\phi$  such that  $f(x) = \phi(x) = y$ . Let  $k_1, k_2$ , and  $k_3$  be tapes of length  $\kappa$  chosen uniformly at random from  $\{0, 1\}^\kappa$  corresponding to

players  $P_1, P_2$  and  $P_3$ , respectively. Consider the following set of functions,  $\mathcal{D}$ :

$$\begin{aligned} (\text{view}_1^{(0)}, \text{view}_2^{(0)}, \text{view}_3^{(0)}) &\leftarrow \text{Share}(x, k_1, k_2, k_3) \\ \text{view}_i^{(j+1)} &\leftarrow \text{Update}(\text{view}_i^{(j)}, \text{view}_{i+1}^{(j)}, k_i, k_{i+1}) \\ y_i &\leftarrow \text{Output}(\text{View}_i) \\ y &\leftarrow \text{Reconstruct}(y_1, y_2, y_3) \end{aligned}$$

such that **Share** is a potentially randomized function that takes  $x$  as input and outputs the initial view for each player containing the secret share of  $x_i$  of  $x$  - i.e.  $\text{view}_i^{(0)} = x_i$ . The function **Update** computes the wire values for the next gate and updates the view accordingly. The function **Output** <sub>$i$</sub>  takes as input the final view,  $\text{View}_i \equiv \text{view}_i^{(n)}$  after all gates have been computed and outputs player  $P_i$ 's *output share*, denoted  $y_i$ .

Correctness requires that reconstructing a (2,3)-decomposed evaluation of a circuit  $\phi$  yields the same value as directly evaluating  $\phi$  on the input value. The 2-privacy property requires that revealing the values from two shares reveals nothing about the input value. More formally, these two properties are defined as follows: We define the experiment  $\text{EXP}_{\text{decomp}}^{(\phi, x)}$  in Experiment 1, which runs the decomposition over a circuit  $\phi$  on input  $x$ : We say that  $\mathcal{D}$  is

$\text{EXP}_{\text{decomp}}^{(\phi, x)}$ :

1. First run the **Share** function on  $x$ :  $\text{view}_1^{(0)}, \text{view}_2^{(0)}, \text{view}_3^{(0)} \leftarrow \text{Share}(x, k_1, k_2, k_3)$
2. For each of the three views, call the update function successively for every gate in the circuit:  $\text{view}_i^{(j)} = \text{Update}(\text{view}_i^{(j-1)}, \text{view}_{i+1}^{(j-1)}, k_i, k_{i+1})$  for  $i \in [1, 3], j \in [1, n]$
3. From the final views, compute the output share of each view:  $y_i \leftarrow \text{output}(\text{View}_i)$

### Experiment 1: Decomposition Experiment

a (2, 3)-*decomposition* of  $\phi$  if the following two properties hold when running  $\text{EXP}_{\text{decomp}}^{(\phi, x)}$ :

**Correctness.** For all circuits  $\phi$ , for all inputs  $x$  and for the  $y_i$ 's produced by  $\text{EXP}_{\text{decomp}}^{(\phi, x)}$ ,

$$\Pr[\phi(x) = \text{Reconstruct}(y_1, y_2, y_3)] = 1.$$

**2-Privacy.** Let  $\mathcal{D}$  be correct. Then for all  $e \in \{1, 2, 3\}$  there exists a PPT simulator  $\mathcal{S}_e$  such that for any probabilistic polynomial-time (PPT) algorithm  $\mathcal{A}$ , for all circuits  $\phi$ , for all inputs  $x$ , and for the distribution of views and  $k_i$ 's produced by  $\text{EXP}_{\text{decomp}}^{(\phi, x)}$  we have that

$$\left| \Pr[\mathcal{A}(x, y, k_e, \text{View}_e, k_{e+1}, \text{View}_{e+1}, y_{e+2}) = 1] - \Pr[\mathcal{A}(x, y, \mathcal{S}_e(\phi, y)) = 1] \right|$$

is negligible.

We now discuss the (2,3)-decomposition used by ZKB++. Let  $R$  be an arbitrary finite ring and  $\phi$  a function such that  $\phi : R^m \rightarrow R^\ell$  can be expressed by an  $n$ -gate arithmetic circuit over the ring using addition by constant, multiplication by constant, addition and multiplication gates. A (2, 3)-decomposition of  $\phi$  is given by the following functions. In the notation below, arithmetic operations are done in  $R^s$  where the operands are elements of  $R^s$ ):

- $(x_1, x_2, x_3) \leftarrow \text{Share}(x, k_1, k_2, k_3)$  samples random  $x_1, x_2, x_3 \in R^m$  such that  $x_1 + x_2 + x_3 = x$ .
- $y_i \leftarrow \text{Output}_i(\text{view}_i^{(n)})$  selects the  $\ell$  output wires of the circuit as stored in the view  $\text{view}_i^{(n)}$ .
- $y \leftarrow \text{Reconstruct}(y_1, y_2, y_3) = y_1 + y_2 + y_3$
- $\text{view}_i^{(j+1)} \leftarrow \text{Update}_i^{(j)}(\text{view}_i^{(j)}, \text{view}_{i+1}^{(j)}, k_i, k_{i+1})$  computes  $P_i$ 's view of the output wire of gate  $g_j$  and appends it to the view. Notice that it takes as input the views and random tapes of both party  $P_i$  as well as party  $P_{i+1}$ . We use  $w_k$  to refer to the  $k$ -th wire, and we use  $w_k^{(i)}$  to refer to the value of  $w_k$  in party  $P_i$ 's view. The update operation depends on the type of gate  $g_j$ .

The gate-specific operations are defined as follows.

**Addition by Constant** ( $w_b = w_a + k$ ):

$$w_b^{(i)} = \begin{cases} w_a^{(i)} + k & \text{if } i = 1, \\ w_a^{(i)} & \text{otherwise.} \end{cases}$$

**Multiplication by Constant** ( $w_b = w_a \cdot k$ ):

$$w_b^{(i)} = k \cdot w_a^{(i)}$$

**Binary Addition** ( $w_c = w_a + w_b$ ):

$$w_c^{(i)} = w_a^{(i)} + w_b^{(i)}$$

**Binary Multiplication** ( $w_c = w_a \cdot w_b$ ):

$$w_c^{(i)} = w_a^{(i)} \cdot w_b^{(i)} + w_a^{(i+1)} \cdot w_b^{(i)} + w_a^{(i)} \cdot w_b^{(i+1)} + R_i(c) - R_{i+1}(c),$$

where  $R_i(c)$  is the  $c$ -th output of a pseudorandom generator seeded with  $k_i$ .

Note that with the exception of the constant addition gate, the gates are symmetric for all players. Also note that  $P_i$  can compute all gate types locally with the exception of binary multiplication gates as this requires inputs from  $P_{i+1}$ . In other words, for every operation except binary multiplication, the **Update** function does not use the inputs from the second party, i.e.,  $\text{view}_{i+1}^{(j)}$  and  $k_{i+1}$ .

While we do not give the details here, [GMO16a] shows that this decomposition meets the correctness and 2-privacy requirements of Definition 2.9. In other words, for every operation except binary multiplication, the **Update** function does not use the inputs from the second party, i.e.,  $\text{view}_{i+1}^{(j)}$  and  $R_{i+1}$ .

## 2.8 ZKB++

ZKB++, an optimized version of ZKBoo [GMO16a], is a proof system for zero-knowledge proofs on arbitrary circuits. ZKBoo and ZKB++ build on the MPC-in-the-head paradigm of Ishai *et al.* [IKOS09], that we describe

only informally here. The multiparty computation protocol (MPC) will implement the relation, and the input is the witness. For example, the MPC could compute  $y = \text{SHA-256}(x)$  where players each have a share of  $x$  and  $y$  is public. The idea is to have the prover simulate a multiparty computation protocol “in their head”, commit to the state and transcripts of all players, then have the verifier “corrupt” a random subset of the simulated players by seeing their complete state. The verifier then checks that the corrupted players performed the correct computation, and if so, he has some assurance that the output is correct. Iterating this for many rounds then gives the verifier high assurance.

ZKBOO generalizes the idea of [IKOS09] by replacing MPC with circuit decompositions. In Scheme 1 and Scheme 2 we present the prover and the verifier of the ZKB++  $\Sigma$ -protocol.

**P.Commit :** 1. For each iteration  $i \in [t]$ : Sample random seeds  $k_1^{(i)}, k_2^{(i)}, k_3^{(i)}$  obtain view  $\text{View}_j^{(i)}$  and output share  $y_j^{(i)}$ . For each player  $P_j$  compute

- (a)  $(x_1^{(i)}, x_2^{(i)}, x_3^{(i)}) \leftarrow \text{Share}(x, k_1^{(i)}, k_2^{(i)}, k_3^{(i)})$
- (b)  $\text{View}_j^{(i)} \leftarrow \text{Update}(\dots \text{Update}(x_j^{(i)}, x_{j+1}^{(i)}, k_j^{(i)}, k_{j+1}^{(i)}) \dots)$
- (c)  $y_j^{(i)} \leftarrow \text{Output}(\text{View}_j^{(i)})$
- (d) Commit  $C_j^{(i)} \leftarrow \text{Com}(k_j^{(i)}, x_j^{(i)}, \text{View}_j^{(i)}, y_j^{(i)})$ , and let  $\mathbf{a}^{(i)} \leftarrow (y_1^{(i)}, y_2^{(i)}, y_3^{(i)}, C_1^{(i)}, C_2^{(i)}, C_3^{(i)})$ .

2. Return  $(\mathbf{a}^{(i)})_{i \in [t]}$ .

**P.Prove :** On input of a challenge  $(\mathbf{e}^{(i)})_{i \in [t]}$ , set for each iteration  $i \in [1, t]$

$$\mathbf{z}^{(i)} \leftarrow \begin{cases} (\text{View}_2^{(i)}, k_1^{(i)}, k_2^{(i)}) & \text{if } \mathbf{e}^{(i)} = 1, \\ (\text{View}_3^{(i)}, k_2^{(i)}, k_3^{(i)}, x_3^{(i)}) & \text{if } \mathbf{e}^{(i)} = 2, \\ (\text{View}_1^{(i)}, k_3^{(i)}, k_1^{(i)}, x_3^{(i)}) & \text{if } \mathbf{e}^{(i)} = 3. \end{cases}$$

and return  $(\mathbf{z}^{(i)})_{i \in [t]}$ .

**Scheme 1:** The prover of the ZKB++  $\Sigma$ -protocol.

We now discuss various instantiation aspects of ZKB++ that make the



**V.Challenge :** Store  $(\mathbf{a}^{(i)})_{i \in [t]}$  and return  $(\mathbf{e}^{(i)})_{i \in [t]} \xleftarrow{R} \mathbf{E}^t$ .

**V.Verify :** 1. For each iteration  $i \in [t]$  reconstruct the views, input and output shares that were not explicitly given as part of the proof response  $\mathbf{z}^{(i)}$ :

(a) Set

$$x_{\mathbf{e}^{(i)}}^{(i)} \leftarrow \begin{cases} R_{\mathbf{e}^{(i)}}(0) & \text{if } \mathbf{e}^{(i)} \neq 3, \\ x_3^{(i)} \text{ given as part of } \mathbf{z}^{(i)} & \text{if } \mathbf{e}^{(i)} = 3. \end{cases}$$

$$x_{\mathbf{e}^{(i)}+1}^{(i)} \leftarrow \begin{cases} R_{\mathbf{e}^{(i)}+1}(0) & \text{if } \mathbf{e}^{(i)} \neq 2, \\ x_3^{(i)} \text{ given as part of } \mathbf{z}^{(i)} & \text{if } \mathbf{e}^{(i)} = 2. \end{cases}$$

(b) Obtain  $\text{View}_{\mathbf{e}^{(i)}+1}^{(i)}$  from  $\mathbf{z}^{(i)}$ .

(c)  $\text{View}_e^{(i)} \leftarrow \text{Update}(\dots \text{Update}(x_{\mathbf{e}^{(i)}}^{(i)}, x_{e+1}^{(i)}, k_e^{(i)}, k_{e+1}^{(i)}) \dots)$

(d)  $y_{\mathbf{e}^{(i)}}^{(i)} \leftarrow \text{Output}(\text{View}_{\mathbf{e}^{(i)}}^{(i)})$

(e)  $y_{\mathbf{e}^{(i)}+1}^{(i)} \leftarrow \text{Output}(\text{View}_{\mathbf{e}^{(i)}+1}^{(i)})$

(f)  $y_{\mathbf{e}^{(i)}+2}^{(i)} \leftarrow y - y_{\mathbf{e}^{(i)}}^{(i)} - y_{\mathbf{e}^{(i)}+1}^{(i)}$

2. Re-compute the commitments for views  $\text{View}_{\mathbf{e}^{(i)}}^{(i)}$  and  $\text{View}_{\mathbf{e}^{(i)}}^{(i)}$ . For  $j \in \{\mathbf{e}^{(i)}, \mathbf{e}^{(i)} + 1\}$ :

$$C_j^{(i)} \leftarrow \text{Com}(k_j^{(i)}, x_j^{(i)}, \text{View}_j^{(i)}, y_j^{(i)})$$

3. Set  $\mathbf{a}'^{(i)} \leftarrow (y_1^{(i)}, y_2^{(i)}, y_3^{(i)}, C_1^{(i)}, C_2^{(i)}, C_3^{(i)})$  taking  $C_{\mathbf{e}^{(i)}+2}^{(i)}$  from  $\mathbf{a}^{(i)}$ .

4. If  $\mathbf{a}'^{(i)} = \mathbf{a}^{(i)}$  for all  $i \in [t]$ , output **Accept**, otherwise **Reject**.

**Scheme 2:** The verifier of the ZKB++  $\Sigma$ -protocol.

optimizations with respect to ZKBOO possible. To highlight the difference, we also present the Fiat-Shamir transformed ZKBOO proof system in Scheme 3 and the Fiat-Shamir transformed ZKB++ in Scheme 4.

**The Share Function.** We make the **Share** function sample the shares pseudorandomly as:

$$(x_1, x_2, x_3) \leftarrow \text{Share}(x, k_1, k_2, k_3) := \\ x_1 = R_1(0), \quad x_2 = R_2(0), \quad x_3 = x - x_1 - x_2.$$

$R_i$  is a pseudorandom generator seeded with  $k_i$ . We specify the **Share** function in this manner as it will lead to more compact proofs. Moving now to the ZKBOO protocol, for each round, the prover is required to “open” two views. In order to verify the proof, the verifier must be given both the random tape and the input share for each opened view. If these values are generated independently of one another, then the prover will have to explicitly include both of them in the proof. However, with our sampling method, in **View**<sub>1</sub> and **View**<sub>2</sub>, the prover only needs to include  $k_i$ , as  $x_i$  can be deterministically computed by the verifier.

The exact savings depends on which views the prover must open, and thus depends on the challenge.

**Not Including Input Shares.** Since the input shares are generated pseudorandomly using the seed  $k_i$ , we do not need to include them in the view when  $e = 1$ . However, if  $e = 2$  or  $e = 3$ , we still need to send one input share for the third view for which the input share cannot be derived from the seed. Thus we explicitly specify the input share when required and do not include it in **View** <sub>$i$</sub> <sup>( $j$ )</sup>.

**No Additional Randomness for Commitments.** Since the first input to the commitment is the seed value  $k_i$  for the random tape, the protocol input to the commitment doubles as a randomization value, ensuring that commitments are hiding. To simplify security analysis, we in fact choose two different random oracles  $H', H''$ . We use  $H'(k_i)$  as the seed to generate the random tape used to generate the input shares and views, and we use  $H''(k_i)$  as input to the commitment. In the random oracle model then, this produces two independent random values; as  $H''(k_i)$  for the unopened view only appears as input to the commitment, this effectively replaces the randomness needed for the commitment scheme in the RO model. (Since one already needs the RO model to make the proofs non-interactive, there is no extra

**Prove<sub>H</sub>**( $1^\kappa, y, x$ ):

1. For each iteration  $i \in [1, t]$ : Sample random tapes  $k_1^{(i)}, k_2^{(i)}, k_3^{(i)}$  and run the decomposition to get an output view  $\text{View}_j^{(i)}$  and output share  $y_j^{(i)}$ . In particular, for each player  $P_j$ :
  - (a)  $(x_1^{(i)}, x_2^{(i)}, x_3^{(i)}) \leftarrow \text{Share}(x, k_1^{(i)}, k_2^{(i)}, k_3^{(i)})$
  - (b)  $\text{View}_j^{(i)} \leftarrow \text{Update}(\dots \text{Update}(x_j^{(i)}, x_{j+1}^{(i)}, k_j^{(i)}, k_{j+1}^{(i)}) \dots)$
  - (c)  $y_j^{(i)} \leftarrow \text{Output}(\text{View}_j^{(i)})$
  - (d) Commit  $C_j^{(i)} \leftarrow \text{Com}(k_j^{(i)}, \text{View}_j^{(i)})$ , and let  $\mathbf{a}^{(i)} \leftarrow (y_1^{(i)}, y_2^{(i)}, y_3^{(i)}, C_1^{(i)}, C_2^{(i)}, C_3^{(i)})$ .
2. Compute the challenge:  $\mathbf{e} \leftarrow H(\mathbf{a}^{(1)}, \dots, \mathbf{a}^{(t)})$ . Interpret the challenge such that for  $i \in [1, t]$ ,  $\mathbf{e}^{(i)} \in \{1, 2, 3\}$
3. For each iteration  $i \in [1, t]$ , let  $\mathbf{z}^{(i)} = (D_{\mathbf{e}^{(i)}}^{(i)}, D_{\mathbf{e}^{(i)}+1}^{(i)})$ .
4. Output  $\pi \leftarrow [(\mathbf{a}^{(1)}, \mathbf{z}^{(1)}), (\mathbf{a}^{(2)}, \mathbf{z}^{(2)}), \dots, (\mathbf{a}^{(t)}, \mathbf{z}^{(t)})]$

**Verify<sub>H</sub>**( $1^\kappa, y, \pi$ ):

1. Parse  $\pi$  as  $[(\mathbf{a}^{(1)}, \mathbf{z}^{(1)}), (\mathbf{a}^{(2)}, \mathbf{z}^{(2)}), \dots, (\mathbf{a}^{(t)}, \mathbf{z}^{(t)})]$ .
2. Compute the challenge:  $\mathbf{e}' \leftarrow H(\mathbf{a}^{(1)}, \dots, \mathbf{a}^{(t)})$ . Interpret the challenge such that for  $i \in [1, t]$ ,  $\mathbf{e}'^{(i)} \in \{1, 2, 3\}$ .
3. For each iteration  $i \in [1, t]$ : If there exists  $j \in \{\mathbf{e}'^{(i)}, \mathbf{e}'^{(i)} + 1\}$  such that  $\text{Open}(C_j^{(i)}, D_j^{(i)}) = \perp$ , output **Reject**. Otherwise, for all  $j \in \{\mathbf{e}'^{(i)}, \mathbf{e}'^{(i)} + 1\}$ , set  $\{k_j^{(i)}, \text{View}_j^{(i)}\} \leftarrow \text{Open}(C_j^{(i)}, D_j^{(i)})$ .
4. For each iteration  $i \in [1, t]$ : If  $\text{Reconstruct}(y_1^{(i)}, y_2^{(i)}, y_3^{(i)}) \neq y$ , output **Reject**. If there exists  $j \in \{\mathbf{e}'^{(i)}, \mathbf{e}'^{(i)} + 1\}$  such that  $y_j^{(i)} \neq \text{Output}(\text{View}_j^{(i)})$ , output **Reject**. For each wire value  $w_j^{(\mathbf{e})} \in \text{View}_e$ , if  $w_j^{(\mathbf{e})} \neq \text{Update}(\text{view}_e^{(j-1)}, \text{view}_{e+1}^{(j-1)}, k_e, k_{e+1})$  output **Reject**.
5. Output **Accept**.

**Scheme 3:** The ZKBOO non-interactive proof system.

**Prove<sub>H</sub>**( $1^\kappa, y, x$ ): 1. For each iteration  $i \in [t]$ : Sample random tapes  $k_1^{(i)}, k_2^{(i)}, k_3^{(i)}$  and obtain output view  $\text{View}_j^{(i)}$  and output share  $y_j^{(i)}$ . For each player  $P_j$  compute

- (a)  $(x_1^{(i)}, x_2^{(i)}, x_3^{(i)}) \leftarrow \text{Share}(x, k_1^{(i)}, k_2^{(i)}, k_3^{(i)})$
- (b)  $\text{View}_j^{(i)} \leftarrow \text{Update}(\dots \text{Update}(x_j^{(i)}, x_{j+1}^{(i)}, k_j^{(i)}, k_{j+1}^{(i)}) \dots)$
- (c)  $y_j^{(i)} \leftarrow \text{Output}(\text{View}_j^{(i)})$
- (d) Commit  $C_j^{(i)} \leftarrow \text{Com}(k_j^{(i)}, x_j^{(i)}, \text{View}_j^{(i)}, y_j^{(i)})$ , and let  $\mathbf{a}^{(i)} \leftarrow (y_1^{(i)}, y_2^{(i)}, y_3^{(i)}, C_1^{(i)}, C_2^{(i)}, C_3^{(i)})$ .

2. Compute the challenge:  $\mathbf{e} \leftarrow H(\mathbf{a}^{(1)}, \dots, \mathbf{a}^{(t)})$ .

3. For each iteration  $i \in [1, t]$  set:  $\mathbf{z}^{(i)} \leftarrow (\text{View}_2^{(i)}, k_1^{(i)}, k_2^{(i)})$  if  $\mathbf{e}^{(i)} = 1$  and  $\mathbf{z}^{(i)} \leftarrow (\text{View}_{\mathbf{e}^{(i)}+1}^{(i)}, k_{\mathbf{e}^{(i)}}^{(i)}, k_{\mathbf{e}^{(i)}+1}^{(i)}, x_3^{(i)})$  otherwise, and return  $\pi \leftarrow (\mathbf{e}, \mathbf{z}_{i \in [t]}^{(i)})$ .

**Verify<sub>H</sub>**( $1^\kappa, y, \pi$ ): Parse  $\pi$  as  $(\mathbf{e}, \mathbf{z}_{i \in [t]}^{(i)})$ .

1. For each iteration  $i \in [t]$  reconstruct the views, input and output shares that were not explicitly given as part of the proof  $\mathbf{z}^{(i)}$ :
  - (a) Set  $x_{\mathbf{e}^{(i)}}^{(i)} \leftarrow R_{\mathbf{e}^{(i)}}(0)$  if  $\mathbf{e}^{(i)} \neq 3$ , otherwise obtain  $x_3^{(i)}$  from  $\mathbf{z}^{(i)}$ . Set  $x_{\mathbf{e}^{(i)}+1}^{(i)} \leftarrow R_{\mathbf{e}^{(i)}+1}(0)$  if  $\mathbf{e}^{(i)} \neq 2$ , otherwise obtain  $x_3^{(i)}$  from  $\mathbf{z}^{(i)}$ .
  - (b) Obtain  $\text{View}_{\mathbf{e}^{(i)}+1}^{(i)}$  from  $\mathbf{z}^{(i)}$ .
  - (c)  $\text{View}_e^{(i)} \leftarrow \text{Update}(\dots \text{Update}(x_{\mathbf{e}^{(i)}}^{(i)}, x_{e+1}^{(i)}, k_e^{(i)}, k_{e+1}^{(i)}) \dots)$
  - (d)  $y_{\mathbf{e}^{(i)}}^{(i)} \leftarrow \text{Output}(\text{View}_{\mathbf{e}^{(i)}}^{(i)})$ ,  $y_{\mathbf{e}^{(i)}+1}^{(i)} \leftarrow \text{Output}(\text{View}_{\mathbf{e}^{(i)}+1}^{(i)})$
  - (e)  $y_{\mathbf{e}^{(i)}+2}^{(i)} \leftarrow y + y_{\mathbf{e}^{(i)}}^{(i)} + y_{\mathbf{e}^{(i)}+1}^{(i)}$
2. Re-compute the commitments for views  $\text{View}_{\mathbf{e}^{(i)}}^{(i)}$  and  $\text{View}_{\mathbf{e}^{(i)}+1}^{(i)}$ . For  $j \in \{\mathbf{e}^{(i)}, \mathbf{e}^{(i)} + 1\}$ :  $C_j^{(i)} \leftarrow \text{Com}(k_j^{(i)}, x_j^{(i)}, \text{View}_j^{(i)}, y_j^{(i)})$ .
3. Set  $\mathbf{a}'^{(i)} \leftarrow (y_1^{(i)}, y_2^{(i)}, y_3^{(i)}, C_1^{(i)}, C_2^{(i)}, C_3^{(i)})$  taking  $C_{\mathbf{e}^{(i)}+2}^{(i)}$  from  $\mathbf{a}^{(i)}$ .
4. Re-compute the challenge:  $\mathbf{e}' \leftarrow H(\mathbf{a}'^{(1)}, \dots, \mathbf{a}'^{(t)})$ . If  $\mathbf{e} = \mathbf{e}'$  output **Accept**, otherwise **Reject**.

**Scheme 4:** The Fiat-Shamir transformed ZKB++ protocol.

assumption here.) Hence we will use the following hash-based commitment scheme:

**Com**( $M$ ) : Set  $C \leftarrow H(M)$  and return  $(C, M)$ ;

**Open**( $C, D$ ) : Return  $M$  if  $H(D) = C$ , and return  $\perp$  otherwise.

We will prove in Section 5.1 that our proof system is secure when using this scheme.

**Not Including the Output Shares.** The output shares  $y_i$  are included in the proof as part of **a**. Moreover, for the two views that are opened, those output shares are included a second time. First, we do not need to send two of the output shares twice. We actually do not need to send any output shares at all as they can be deterministically computed from the rest of the proof as follows:

For the two views that are given as part of the proof, the output share can be recomputed from the remaining parts of the view. Essentially, the output share is just the value on the output wires. Given the random tapes and the communicated bits from the binary multiplication gates, all wires for both views can be recomputed.

For the third view, recall that the **Reconstruct** function simply adds the three output shares to obtain  $y$ . But the verifier is given  $y$ , and can thus instead recompute the third output share. In particular, given  $y_i$ ,  $y_{i+1}$  and  $y$ , the verifier can compute:  $y_{i+2} = y - y_i - y_{i+1}$ . Thus we explicitly specify the output share when required and do not include it in  $\text{View}_i^{(j)}$ .

When applying the Fiat-Shamir and Unruh transforms to ZKB++ to obtain a signature scheme, we can also perform the following modifications.

**Not Including Commitments.** It is unnecessary to send all three commitments to the verifier. Since for the two views that are opened, the verifier can recompute the commitment. Only for the third view that the verifier is not given the commitment needs to be explicitly sent.

**Security.** One can observe that all optimizations except “*No Additional Randomness for Commitments*” are equivalence transformations, and, therefore, do not impact the security of the overall ZKB++ proof system. In Section 5.1, we formally confirm that using no additional randomness for the commitments does not impact the security of the ZKB++ proof system.

## 2.9 KKW

KKW [KKW18] is another variant of ZKBoo, focusing on minimizing the signature size. It is very similar to ZKBoo and ZKB++ but uses an MPC protocol with more than three parties.

### 2.9.1 The MPC Protocol

Recall that the prover needs to simulate a set of parties,  $S_1, \dots, S_n$ . In KKW, all simulated parties run an  $n$ -party protocol  $\Pi$  in the preprocessing model, secure against semi-honest corruption of all-but-one of the parties.

The protocol  $\Pi$  maintains the invariant that, for each wire in the circuit, the parties hold an  $n$ -out-of- $n$  secret sharing of a random mask along with the (public) masked value of the wire. Specifically, if we let  $z_\alpha$  denote the value of wire  $\alpha$  in the circuit  $C$  when evaluated on input  $w$ , then the parties will hold  $[\lambda_\alpha]$  (for uniform  $\lambda_\alpha \in \{0, 1\}$ ) along with the value  $\hat{z}_\alpha \stackrel{\text{def}}{=} z_\alpha \oplus \lambda_\alpha$ .

**Preprocessing phase.** In the preprocessing phase, shares are set up among the parties as follows. For each wire  $\alpha$  that is either an input wire of the circuit or the output wire of an AND gate, the parties are given  $[\lambda_\alpha]$ , where  $\lambda_\alpha \in \{0, 1\}$  is uniform. For an XOR gate with input wires  $\alpha, \beta$  and output wire  $\gamma$ , define  $\lambda_\gamma \stackrel{\text{def}}{=} \lambda_\alpha \oplus \lambda_\beta$ ; note the parties can compute  $[\lambda_\gamma]$  locally. Finally, for each AND gate with input wires  $\alpha, \beta$ , the parties are given  $[\lambda_{\alpha, \beta}]$ , where  $\lambda_{\alpha, \beta} \stackrel{\text{def}}{=} \lambda_\alpha \cdot \lambda_\beta$ .

One observation is that the shares of the  $\{\lambda_\alpha\}$  are uniform, and so can be generated by having each party  $S_i$  apply a pseudorandom generator to a short, random seed  $\mathbf{seed}_i$  given to that party, and then (implicitly) defining the  $\{\lambda_\alpha\}$  based on the resulting shares. All-but-one of the shares of the  $\{\lambda_{\alpha, \beta}\}$  can also be generated in this way, but the final share is constrained by the values of  $\lambda_\alpha, \lambda_\beta$ . To ensure that the shares of the  $\{\lambda_{\alpha, \beta}\}$  are correct,  $S_n$  can be given an additional  $|C|$  “correction bits” that determine its share of  $\lambda_{\alpha, \beta}$  for each AND gate with input wires  $\alpha, \beta$ .

To summarize: each  $S_i$  is given a  $\kappa$ -bit seed  $\mathbf{seed}_i \in \{0, 1\}^\kappa$ , and  $S_n$  is additionally given  $|C|$  bits denoted by  $\mathbf{aux}_n$ . We refer to this information as the *state* of the parties, and denote the state of  $S_i$  by  $\mathbf{state}_i$ . In the online phase of the protocol, each party  $S_i$  uses  $\mathbf{seed}_i$  to generate its shares of the  $\{\lambda_\alpha\}$ ; for  $1 \leq i \leq n - 1$ , party  $S_i$  also uses  $\mathbf{seed}_i$  to generate its shares of the  $\{\lambda_{\alpha, \beta}\}$ . Party  $S_n$  uses  $\mathbf{aux}_n$  as its shares of the  $\{\lambda_{\alpha, \beta}\}$ .

**Protocol execution.** Note that in this setting, where all parties are semi-honest, we can perform public reconstruction of a shared value  $[x]$  by simply having each party broadcast its share.

We assume the parties begin the protocol holding a masked value  $\hat{z}_\alpha$  for each input wire  $\alpha$ . (In this context these masked values will be provided to the parties by the prover who is simulating execution of the protocol.) These masked values, along with the corresponding  $\{\lambda_\alpha\}$ , define an effective input to the protocol. During the online phase of the protocol, the parties inductively compute  $\hat{z}_\alpha$  for all wires in the circuit. Specifically, for each gate of the circuit with input wires  $\alpha, \beta$  and output wire  $\gamma$ , where the parties already hold  $\hat{z}_\alpha, \hat{z}_\beta$ , the parties do:

- If the gate is an XOR gate, then the parties locally compute

$$\hat{z}_\gamma := \hat{z}_\alpha \oplus \hat{z}_\beta .$$

- If the gate is an AND gate, the parties locally compute

$$[s] := \hat{z}_\alpha[\lambda_\beta] \oplus \hat{z}_\beta[\lambda_\alpha] \oplus [\lambda_{\alpha,\beta}] \oplus [\lambda_\gamma],$$

and then publicly reconstruct  $s$ . Finally, they compute  $\hat{z}_\gamma := s \oplus \hat{z}_\alpha \hat{z}_\beta$ .

One can verify that  $\hat{z}_\gamma = z_\gamma \oplus \lambda_\gamma$ .

Once the parties have computed  $\hat{z}_\alpha$  for the output wire  $\alpha$ , the output value  $z_\alpha$  is computed by publicly reconstructing  $\lambda_\alpha$  and then setting  $z_\alpha := \hat{z}_\alpha \oplus \lambda_\alpha$ .

We remark that the online phase of this protocol is deterministic. Also observe that all communication is due to share reconstruction: for a circuit with  $|C|$  AND gates, at most  $|C| + 1$  share reconstructions are needed.

### 2.9.2 The Proof Protocol

The high-level idea is to have the prover run  $M$  emulations of the preprocessing phase and their online phases, and commit to all emulations. The verifier selects  $M - \tau$  of them and checks the preprocessing phase; the online phase of the remaining  $\tau$  executions will be checked by revealing the view of all-but-one parties. The protocol is shown in Figure 1 and Figure 2. In the following, we will discuss more details of KKW.

**Checking the preprocessing phase.** Recall from the previous section that, following the preprocessing phase, the state of party  $S_i$  for  $1 \leq i \leq n-1$  is a seed  $\mathbf{seed}_i$ , while the state of party  $S_n$  is a seed  $\mathbf{seed}_n$  along with a  $|C|$ -bit string  $\mathbf{aux}_n$ . We improve the communication complexity in several ways:

### KKW Prover Algorithm

**Inputs:** Both parties have a circuit  $C$ ; the prover also holds  $w$  with  $C(w) = 1$ . Values  $M, n, \tau$  are parameters of the protocol.

**Round 1** For each  $j \in [M]$ , the prover does:

1. Choose uniform  $\text{seed}_j^* \in \{0, 1\}^\kappa$  and use it to generate values  $\text{seed}_{j,1}, r_{j,1}, \dots, \text{seed}_{j,n}, r_{j,n}$ . Also compute  $\text{aux}_j \in \{0, 1\}^{|C|}$  as described in the text. For  $i = 1, \dots, n-1$ , let  $\text{state}_{j,i} := \text{seed}_{j,i}$ ; let  $\text{state}_{j,n} := \text{seed}_{j,n} \parallel \text{aux}_j$ .
2. For  $i \in [n]$ , compute  $\text{com}_{j,i} := \text{Com}(\text{state}_{j,i}; r_{j,i})$ .
3. The prover simulates the online phase of the  $n$ -party protocol  $\Pi$  (as described in the text) using  $\{\text{state}_{j,i}\}_i$ , beginning by computing the masked inputs  $\{\hat{z}_{j,\alpha}\}$  (based on  $w$  and the  $\{\lambda_{j,\alpha}\}$  defined by the preprocessing). Let  $\text{msgs}_{j,i}$  denote the messages broadcast by  $S_i$  in this protocol execution.
4. Let  $h_j := H(\text{com}_{j,1}, \dots, \text{com}_{j,n})$  and  $h'_j := H(\{\hat{z}_{j,\alpha}\}, \text{msgs}_{j,1}, \dots, \text{msgs}_{j,n})$ .

Compute  $h := H(h_1, \dots, h_M)$  and  $h' := H(h'_1, \dots, h'_M)$  and send  $h^* := H(h, h')$  to the verifier.

**Round 2** The verifier chooses a uniform  $\tau$ -sized set  $\mathcal{C} \subset [M]$  and  $\mathcal{P} = \{p_j\}_{j \in [\tau]}$  where each  $p_j \in [n]$  is uniform. Send  $(\mathcal{C}, \mathcal{P})$  to the prover.

**Round 3** For each  $j \in [M] \setminus \mathcal{C}$ , the prover sends  $\text{seed}_j^*, h'_j$  to the verifier.

For each  $j \in \mathcal{C}$ , the prover sends  $\{\text{state}_{j,i}, r_{j,i}\}_{i \neq p_j}, \text{com}_{j,p_j}, \{\hat{z}_{j,\alpha}\}$ , and  $\text{msgs}_{j,p_j}$  to the verifier.

Figure 1: Prover algorithm of the KKW proof protocol.

1. The prover computes  $H(\text{com}_1, \dots, \text{com}_n)$ , and then sends the hash of the results from all  $m$  executions; thus, it sends just a single hash value to the verifier.



### KKW Verifier Algorithm

**Verification** The verifier accepts iff all the following checks succeed:

1. For every  $j \in \mathcal{C}$ ,  $i \neq p_j$ , the verifier uses  $\mathbf{state}_{j,i}$  and  $r_{j,i}$  to compute  $\mathbf{com}_{j,i}$ . It then computes  $h_j := H(\mathbf{com}_{j,1}, \dots, \mathbf{com}_{j,n})$ .
2. For  $j \in [M] \setminus \mathcal{C}$ , the verifier uses  $\mathbf{seed}_j^*$  to compute  $h_j$  as an honest prover would. It then computes  $h := H(h_1, \dots, h_M)$ .
3. For each  $j \in \mathcal{C}$ , the verifier simulates an execution of  $\Pi$  among the  $\{S_i\}_{i \neq p_j}$  using  $\{\mathbf{state}_{j,i}\}_{i \neq p_j}$ , masked input-wire values  $\{\hat{z}_\alpha\}$ , and  $\mathbf{msgs}_{j,p_j}$ . This yields  $\{\mathbf{msgs}_i\}_{i \neq p_j}$  and an output bit  $b$ . The verifier checks that  $b \stackrel{?}{=} 1$  and computes  $h'_j := H(\{\hat{z}_{j,\alpha}\}, \mathbf{msgs}_{j,1}, \dots, \mathbf{msgs}_{j,n})$  as well as  $h' := H(h'_1, \dots, h'_m)$ .
4. The verifier checks that  $H(h, h') \stackrel{?}{=} h^*$ .

Figure 2: Verifier algorithm of the KKW proof protocol.

2. When opening a challenged execution, it is unnecessary for the prover to send  $\mathbf{aux}_n$  since the correct value of  $\mathbf{aux}_n$  can be computed from  $\mathbf{seed}_1, \dots, \mathbf{seed}_n$ .
3. By generating the  $\{\mathbf{seed}_i\}$  and the  $\{r_i\}$  from a “root” seed  $\mathbf{seed}^* \in \{0, 1\}^\kappa$ , the prover can open a challenged execution of the preprocessing phase by simply sending  $\mathbf{seed}^*$ .

**Checking the online execution.** An execution of the protocol proceeds gate-by-gate, with the processing of each AND gate requiring reconstruction of one shared value. Although the communication complexity of share reconstruction in the protocol is  $n$  bits (one bit per party), for the purposes of checking, we do not need the prover to send  $n$  bits per gate in order to prove consistent execution. This is because the verifier only needs to obtain the protocol messages sent by the (single) *unopened* party in order to check the execution of the  $n - 1$  opened parties. Thus, it suffices for the prover to send just a single bit per AND gate.

In addition to the protocol messages sent by the unopened party, the prover also needs to reveal the state (from the preprocessing phase) of every

opened party. For each opened party  $S_i$ ,  $i \neq n$ , this involves just  $O(\kappa)$  bits; if  $S_n$  is opened then this requires  $|C| + O(\kappa)$  bits due to  $\mathbf{aux}_n$ . In either case the marginal communication complexity per AND gate is *independent* of the number of parties  $n$ .

**Reducing the number of random seeds.** In the  $c$ -th emulation of the preprocessing phase, the prover generates  $n$  seeds  $\mathbf{seed}_{c,1}, \dots, \mathbf{seed}_{c,n}$  from a root seed  $\mathbf{seed}_c^*$ , commits to the  $n$  generated seeds, and then sends  $n - 1$  of those seeds to the verifier. The second step requires  $(n - 1) \cdot \kappa$  bits of communication.

Motivated by the NNL scheme for stateless revocation [NNL01], we observe that we can reduce the communication by generating the seeds in a more structured way. Namely, imagine labeling the root of a binary tree of depth  $\log n$  with  $\mathbf{seed}_c^*$ , and then inductively labeling the children of each node with the output of a pseudorandom generator applied to the node's label. The  $\{\mathbf{seed}_{c,i}\}_{i \in [n]}$  will be the labels of the  $n$  leaves of the tree. To reveal  $\{\mathbf{seed}_{c,i}\}_{i \neq p}$ , it suffices to reveal the labels on the siblings of the path from the root of the tree to leaf  $p$ . Those labels allow the verifier to reconstruct  $\{\mathbf{seed}_{c,i}\}_{i \neq p}$  while still hiding  $\mathbf{seed}_{c,p}$ . Applying this optimization reduces the communication complexity to  $O(\kappa \cdot \log n)$  for revealing the seeds used by the  $n - 1$  opened parties.

We can, in fact, apply the same idea to the root seeds  $\{\mathbf{seed}_j^*\}_{j=1}^m$  used for the different emulations of the preprocessing phase; this reduces the communication required to reveal all-but-one of those seeds in Round 3 from  $(m - 1) \cdot \kappa$  bits to  $O(\kappa \cdot \log m)$  bits. Further, we are not limited to revealing all-but-one of the leaf labels; more generally, the scheme just described supports revealing all-but- $\tau$  of the leaf labels using communication at most  $O(\kappa \cdot \tau \log \frac{m}{\tau})$  bits (cf. [NNL01]).

**Reducing the size of commitments with Merkle trees.**  $M - \tau$  of the commitments  $h'_1, \dots, h'_M$  are sent as part of the proof for the instances where  $j \notin C$ . These are necessary for verification, when recomputing the challenge the prover provides  $M - \tau$  of the commitments, and the verifier recomputes the other  $\tau$ . We can reduce the proof size by having the prover commit to  $h'_1, \dots, h'_M$  using a Merkle tree and hash the root when computing the challenge. Then, the verifier is given the root and enough information to confirm that the  $\tau$  commitments he recomputes are in fact committed to by the root. In this way, only  $O(\tau \log(M/\tau))$  hash values are communicated, which is much less than  $M - \tau$  because  $\tau$  is much smaller than  $M$  in our

parameter sets.

**Reducing the size of commitment openings.** As in ZKB++, we can reduce the size of openings for commitments by not using additional randomness when forming the commitments  $\text{com}_{i,j}$  and  $h'_j$ . Intuitively, since other values in the commitment have sufficient entropy from an attacker’s perspective, having the commitment does not provide additional useful information. This is analyzed more formally in Section 6.2, where our security proof for the signature scheme uses commitments without additional randomness.

### 2.9.3 Further Optimizations for Picnic3

In KKW and similar MPC protocols, linear operations can be computed without any interaction between the parties. This is great for the traditional, interactive MPC setting, but for MPC-in-the-head these linear operations can contribute a nontrivial amount to the overall runtime. In the following, we give optimizations for evaluation of functions that have an expensive linear component, such as LowMC. The optimizations here are part of what separates the previous Picnic2 signatures from Picnic3. More detail and an algorithmic description of these optimizations is given in [KZ20b], here we give an overview.

For these next optimizations, it may be helpful to have the structure of LowMC in mind. Referring to Section 2.10, note that each round of LowMC consists of the S-box layer, which consists primarily of AND operations, followed by the linear layer and key addition layers, which contain no AND operations (on secret shared inputs). Since Picnic3 only uses LowMC instances with a full S-box layer, we assume this case in our presentation.

**Improved Preprocessing in KKW** First, we recall the preprocessing phase in KKW and how its output is used in the online phase. For each wire  $\alpha$  in the circuit, each party  $i$  holds a share of the masking value  $\lambda_\alpha$ , which is used to hide the plain value of  $\alpha$  as  $\hat{z}_\alpha = \alpha \oplus \lambda_\alpha$ . Due to this representation, XOR gates can be computed locally; however, for AND gates, we need to set up *correction values* in the form of correlated random values during the preprocessing phase. Specifically, for an AND gate with input wires  $\alpha, \beta$  and output wire  $\gamma$  the preprocessing phase calculates the value of  $\lambda_{\alpha\beta} = \lambda_\alpha \lambda_\beta$ , and shares this value between all parties. This is done by picking a random share for the first  $n - 1$  parties from their random tape, and the last party’s share is chosen so that the sum of all shares is equal

to  $\lambda_{\alpha\beta}$ . The last party's shares are then stored on the auxiliary tape, which is an output of the preprocessing phase together with the random seeds for all parties. The value  $\lambda_\gamma$  is a *fresh random value* masking the output of an AND gate and is chosen by picking a random share for each party from their random tape. These preprocessing values are then used in the calculation of the online phase to compute the share of the broadcasted values. In detail, each party computes and broadcasts  $[s] = \hat{z}_\alpha[\lambda_\beta] \oplus \hat{z}_\beta[\lambda_\alpha] \oplus [\lambda_{\alpha\beta}] \oplus [\lambda_\gamma]$ , which can then be combined by all parties to calculate  $s$  which is used to fix the masked value  $\hat{z}_\gamma = \hat{z}_\alpha \hat{z}_\beta \oplus s$ .

A direct implementation of the KKW preprocessing phase expands the seeds for each player and then executes the circuit for all of them, fixing wrong values of  $\lambda_{\alpha\beta}$  for each AND gate. However, observe that this correction value **corr** (correcting the last party's share  $[\lambda_{\alpha\beta}]$ ) does not actually depend on the information of the masks of a single party, but only on their sum (**corr** =  $\lambda_{\alpha\beta} \oplus \lambda_\alpha \lambda_\beta$ ). Because this sum is a linear operation, instead of computing the linear layer (which for some ciphers such as LowMC is quite expensive) for each party and then computing the sum of the resulting mask shares to get the combined input mask for each AND gate, we make use of the linearity and exchange the order of these two operations. This results in an  $n$ -fold reduction of the linear layer costs for the preprocessing, where  $n$  is the number of parties in the MPC protocol. For KKW-based Picnic, this optimization speeds up the signing process by a factor of about 1.5, and the verification process by more than a factor of 2 when compared to the direct implementation. Similar techniques have been used in other MPC applications, such as the inner product optimizations used by, e.g., SecureML [MZ17] or ABY3 [MR18].

**Optimized Sampling of Masks** Notice how the preprocessing optimization allows the prover to calculate the linear layer only once during the preprocessing phase, regardless of the number of parties. However, for the online phase, since each party needs to broadcast their share of  $s$ , the prover needs to calculate the input masks to the AND gate  $[\lambda_\alpha], [\lambda_\beta]$ , which are a linear combination of the output shares  $[\lambda_\gamma]$  of the previous S-box layer. The prover needs to compute the linear layer once for each party to get its input masks to the AND gates based on the previous round's output masks. This calculation of  $N$  LowMC linear layers accounted for about 75% of the runtime of the online simulation in Picnic2. In a similar fashion to the preprocessing phase, we also want to avoid the per-party computation of these masks in

the online phase. We will now introduce a variant of the preprocessing phase that is still secure, but allows us to skip the computation of the expensive linear layer in the online phase.

The basic idea is that instead of having to compute the inputs to the S-box layer, which are a linear combination of the output masks of the previous S-box layer, we want to be able to read them from the parties' random tapes instead. In the original preprocessing phase, the shares of the output masks  $\lambda_\gamma$  are sampled from the respective party's random tape. Instead of selecting the shares of the masks of the S-box *output* in this way, we now sample the shares of the *input* to the next S-box layer from the random tape instead.

This change can be iteratively repeated starting from the last round, where during preprocessing, we first choose the output masks and calculate back up to the S-box layer using the invertibility of the linear layer. Then, random masks are sampled from the random tape for the inputs to the S-box layer, and based on these random masks, the correction values for all AND gates are calculated. This process is repeated for the next rounds until we arrive at the first round. There, in contrast to the other rounds, the input masks are not chosen randomly from the tape, but instead, the input masks for the plaintext are set to zero since it is a public input. This means that the input masks for the first S-box layer are only based on the first round key addition. Since the key-matrices of LowMC are also invertible, we can apply this technique here as well and sample masks for the first roundkey  $k_0$  from the tapes instead of sampling masks for the master key  $k$ . Note that even with these changes, we can still use the preprocessing optimization, as the correction values still only depend on the sum of the masks.

The main improvement manifests in the online phase, where, instead of having to read the output masks from the random tape and calculating the linear layer for each party to get the input masks for the next S-box layer, we can simply read the input masks for each party from the random tape. The impact of this optimization is dependent on the specific linear layer that is used and is especially noticeable for the case of LowMC and its large, random linear matrices. Our experiments showed a speedup by a factor of  $\approx 1.63$  for signing. In general, this improvement reduces the cost of the linear parts of the online phase from  $N + 1$  evaluations to 1. While this reduces the impact of  $N$ , the number of MPC parties, on the runtime, there are still other parts of the signature algorithm that depend on  $N$ , like the expansion of random tapes and the non-linear gates. Therefore, choosing parameter sets with a reduced number of parties, as discussed in Section 4.3 is still substantial when

applied to this optimized MPC protocol.

**Security implications of this change.** Due to the linear layer in LowMC (and many other block cipher designs) being invertible, this new location for choosing the random masks is equivalent in terms of security. First, observe that when selecting  $\lambda_\alpha$  uniformly at random, after adding the key masks  $\lambda_k$ , the value  $\lambda_\alpha \oplus \lambda_k$  is still uniformly random. Second, since the linear transformation is invertible, the application of the linear matrix to the state is a bijection. Therefore, if we apply this (inverse) linear layer transformation to a uniformly random  $n$ -bit state, it again results in a uniformly random  $n$ -bit state. This argument shows that moving the random variables which are sampled from the tape from the output of the S-box layer to the input of the following S-box layer has no impact on security since the masks at the output of the S-box layer are still uniformly random, like in the original protocol.

**Removing the Final Broadcast** At the end of the simulation of the MPC protocol, all parties hold a share of the masking value  $\lambda_\alpha$  for the masked output  $\hat{z}_\alpha$ . To reconstruct the output, all parties broadcast  $[\lambda_\alpha]$ , which is then combined and used to unmask the output  $\alpha$ . We will now detail a modification to the MPC protocol that removes the need to broadcast these output masks while not leaking any more information than the original protocol.

The last round of the evaluation of the block cipher involves the S-box layer, where some correction bits  $[s]$  are broadcast, and then, after the final linear layer and key addition, the shares of the output masks are broadcast. Since we reveal the value of the output shares anyway, we now define the value of the shares of the output mask  $[\lambda_\alpha]$  as 0. Calculating backwards from this definition, we need to choose  $[\lambda_\gamma]$ , the share of the output mask for the S-box layer, in a way that it is canceled with the key mask  $[\lambda_k]$ . We can express this change by using an equivalent last round key  $k' = L_r^{-1} \cdot k$ , making it more obvious that these changes result in an output mask of 0. For a function with  $n$  bits of output, this reduces the size of the signature by  $\tau n$  bits.

**Security implications of this change.** After this change, parties now broadcast the value  $[s'] = \hat{z}_\alpha[\lambda_\beta] \oplus \hat{z}_\beta[\lambda_\alpha] \oplus [\lambda_{\alpha\beta}] \oplus [\lambda_{k'}]$ , compared to the two previous values of  $[s] = \hat{z}_\alpha[\lambda_\beta] \oplus \hat{z}_\beta[\lambda_\alpha] \oplus [\lambda_{\alpha\beta}] \oplus [\lambda_\gamma]$  and  $[\lambda_\alpha] = [\lambda_k] \oplus [L(\lambda_\gamma, \dots)]$ , where  $[L(\lambda_\gamma, \dots)]$  is some linear combination of the  $n$  mask values at the output of the S-box layer. From this information an outside observer can exactly calculate the new broadcast value  $[s']$  by simply applying

the inverse linear layer to the output mask shares  $[\lambda_\alpha]$  and adding these masks to the broadcasted values  $[s]$ . Therefore, this new variant of revealing the output does not leak any more information than before.

## 2.10 LowMC

LowMC [ARS<sup>+</sup>15a, ARS<sup>+</sup>16] is a very parameterizable symmetric encryption scheme design enabling instantiation with low AND depth and low multiplicative complexity. Given any blocksize, a choice for the number of S-boxes per round, and security expectations in terms of time and data complexity, instantiations can be created minimizing the AND depth, the number of ANDs, or the number of ANDs per encrypted bit. Table 1 lists the choices for the parameters for security levels L1, L3, L5.

The description of LOWMC is possible independently of the choice of parameters using a partial specification of the S-box and arithmetic in vector spaces over  $\mathbb{F}_2$ . In particular, let  $n$  be the blocksize,  $m$  be the number of S-boxes,  $k$  the key size, and  $r$  the number of rounds, we choose round constants  $C_i \xleftarrow{R} \mathbb{F}_2^n$  for  $i \in [1, r]$ , full rank matrices  $K_i \xleftarrow{R} \mathbb{F}_2^{n \times k}$  and regular matrices  $L_i \xleftarrow{R} \mathbb{F}_2^{n \times n}$  independently during the instance generation and keep them fixed. Keys for LOWMC are generated by sampling from  $\mathbb{F}_2^k$  uniformly at random.

LOWMC encryption starts with key whitening which is followed by several rounds of encryption. A single round of LOWMC is composed of an S-box layer, a linear layer, addition with constants and addition of the round key, i.e.

$$\begin{aligned} \text{LOWMCROUND}(i) &= \text{KEYADDITION}(i) \\ &\quad \circ \text{CONSTANTADDITION}(i) \\ &\quad \circ \text{LINEARLAYER}(i) \circ \text{SBOXLAYER}. \end{aligned}$$

SBOXLAYER is an  $m$ -fold parallel application of the same 3-bit S-box on the first  $3 \cdot m$  bits of the state. The S-box is defined as  $S(a, b, c) = (a \oplus bc, a \oplus b \oplus ac, a \oplus b \oplus c \oplus ab)$ .

The other layers only consist of  $\mathbb{F}_2$ -vector space arithmetic. LINEARLAYER( $i$ ) multiplies the state with the linear layer matrix  $L_i$ , CONSTANTADDITION( $i$ ) adds the round constant  $C_i$  to the state, and KEYADDITION( $i$ ) adds the round key to the state, where the round key is generated by multiplying the master key with the key matrix  $K_i$ .

Algorithm 1 gives a full description of the encryption algorithm.

---

**Algorithm 1** LowMC encryption for key matrices  $K_i \in \mathbb{F}_2^{n \times k}$  for  $i \in [0, r]$ , linear layer matrices  $L_i \in \mathbb{F}_2^{n \times n}$  and round constants  $C_i \in \mathbb{F}_2^n$  for  $i \in [1, r]$ .

---

**Require:** plaintext  $p \in \mathbb{F}_2^n$  and key  $y \in \mathbb{F}_2^k$

```

 $s \leftarrow K_0 \cdot y + p$ 
for  $i \in [1, r]$  do
   $s \leftarrow Sbox(s)$ 
   $s \leftarrow L_i \cdot s$ 
   $s \leftarrow C_i + s$ 
   $s \leftarrow K_i \cdot y + s$ 
end for
return  $s$ 

```

---

LowMC is very flexible in the choice of parameters: the block size  $n$ , the key size  $k$ , the number of 3-bit S-boxes  $m$  in the substitution layer and the allowed data complexity  $d$  of attacks can independently be chosen. To reduce the multiplicative complexity, the number of S-boxes applied in parallel can be reduced, leaving part of the substitution layer as the identity mapping. The number of rounds  $r$  needed to achieve the security goal is then determined as a function of all these parameters. We discuss concrete choices of the parameters in Section 4.2.

### 3 The Picnic Signature Schemes

We consider several schemes, all obtained by transforming an interactive zero-knowledge protocol into a (non-interactive) signature. Different signature schemes are obtained by varying the zero-knowledge protocol used and the transformation that is applied. **Picnic-FS** uses ZKB++ with the Fiat-Shamir transform, **Picnic-UR** uses ZKB++ with the Unruh transform, and **Picnic3** uses the proof protocol of [KKW18, KZ20b] with the Fiat-Shamir transform.<sup>2</sup> Finally, there is also the **Picnic-full** variant, which is identical to **Picnic-FS** except for the choice of LowMC parameters. In Scheme 5 we provide a general description of our schemes.

---

<sup>2</sup>Picnic3 replaces Picnic2, a very similar scheme from earlier versions of this document.



For all schemes, the public key contains values  $u, y$ , and the signer proves knowledge of a pre-image  $x$  of  $y$  with respect to a one-way function  $f_u$ . In all cases, we instantiate the one way function with LOWMC; specifically, if  $F$  denotes the LowMC block cipher then we choose a uniform domain element  $u$  and define the one-way function  $f_u$  via  $f_u(x) = F_x(u)$ . (See Section 2.10 for an overview of LowMC and Section 4.2 for a discussion of our parameter choices.) One-wayness of  $f_u$  follows from the assumption that LowMC is a secure block cipher (i.e., pseudorandom permutation); Section 7 discusses the difficulty of inverting  $f_u$  using known attacks.

<b>Gen</b> ( $1^\kappa$ ) :	Choose $u \xleftarrow{R} \mathbf{K}_\kappa$ , $x \xleftarrow{R} \mathbf{D}_\kappa$ , compute $y \leftarrow f_u(x)$ , set $\mathbf{pk} \leftarrow (y, u)$ and $\mathbf{sk} \leftarrow (\mathbf{pk}, x)$ and return $(\mathbf{sk}, \mathbf{pk})$ .
<b>Sign</b> ( $\mathbf{sk}, m$ ) :	Parse $\mathbf{sk}$ as $(\mathbf{pk}, x)$ , compute $p = (r, s) \leftarrow \text{Prove}_H((y, u), x)$ and return $\sigma \leftarrow p$ , where internally the challenge is computed as $c \leftarrow H(r, \mathbf{pk}  m)$ .
<b>Verify</b> ( $\mathbf{pk}, m, \sigma$ ) :	Parse $\mathbf{pk}$ as $(y, u)$ , and $\sigma$ as $p = (r, s)$ . Return 1 if the following holds, and 0 otherwise: $\text{Verify}_H((y, u), p) = 1,$ where internally the challenge is computed as $c \leftarrow H(r, \mathbf{pk}  m)$ .

**Scheme 5:** Generic description the Picnic-FS and Picnic2-FS signature schemes. Scheme Picnic-UR is similar, except **Prove** and **Verify** are different.

### 3.1 Efficient Instantiation of Unruh’s Transform

Although Unruh’s transformation was only designed for  $\Sigma$ -protocols with 2-special soundness, we can easily modify it for ZKB++ (which has 3-special soundness).

Since ZKB++ has 3-special soundness, we would need at least three responses for each iteration. Moreover, since there are only three possible challenges in ZKB++, we run Unruh’s transform with  $\mathbf{E} = \{1, 2, 3\}$  and  $M = 3$ —i.e., every possible challenge and response. If we apply this naively, we obtain the following protocol:

**Prove**<sub>H</sub>( $1^\kappa, (x, m), w$ ) :

1. For  $i \in [t]$ :
  - (a) Start **P** on  $(1^\kappa, x, w)$  and obtain first message  $\mathbf{a}_i$ .

- (b) For all  $\mathbf{e}_{i,j} = j \in \mathbf{E}$ , obtain response  $\mathbf{z}_{i,j}$  for challenge  $\mathbf{e}_{i,j}$ .
- 2. For  $(i, j) \in [t] \times \mathbf{E}$ , set  $g_{i,j} \leftarrow G(\mathbf{z}_{i,j})$ .
- 3. Let  $(J_1, \dots, J_t) \leftarrow H(m, (\mathbf{a}_i)_{i \in [t]}, (g_{i,1}, \dots, g_{i,3})_{i \in [t]})$
- 4. Return  $\pi \leftarrow ((\mathbf{a}_i)_{i \in [t]}, (g_{i,1}, \dots, g_{i,3})_{(i,j) \in [t]}, (\mathbf{z}_{i,J_i})_{i \in [t]})$

As we no longer randomly select the challenges, we can omit them as input to the hash function and do not need to include them in the proof.

To instantiate the function  $G$  in the protocol, Unruh shows that one does not need a random oracle that is actually a permutation. Instead, as long as the domain and codomain of  $G$  are the same size (and large enough), it can be used, since it is indistinguishable from a random permutation. So let  $G : \{0, 1\}^{|\mathbf{z}_{i,j}|} \rightarrow \{0, 1\}^{|\mathbf{z}_{i,j}|}$  be a hash function modeled as a random oracle. The size of the response changes depending on what the challenge is. If the challenge is 0, the response is slightly smaller as it does not need to include the extra input share. So more precisely, this is actually two hash functions,  $G_0$  used for the 0-challenge response and  $G_{1,2}$  used for the other two challenges. In our specification document we define  $G$  precisely.

**Optimization 1: Making Use of Overlapping Responses.** We can make use of the structure of the ZKB++ proofs to achieve a very significant reduction in the proof size. Although we refer to three separate challenges, in the case of the ZKB++ protocol, there is a large overlap between the contents of the responses corresponding to these challenges. In particular, there are only three distinct views in the ZKB++ protocol, two of which are opened for a given challenge.

Instead of computing a permutation of each *response*,  $\mathbf{z}_{i,j}$ , we can compute a permutation of each *view*,  $v_{i,j}$ . For each  $i \in \{1, \dots, t\}$ , and for each  $j \in \mathbf{E}$ , the prover computes  $g_{i,j} = G(v_{i,j})$ .

The verifier checks the permuted value for each of the two views in the response. In particular, for challenge  $j \in \{1, 2, 3\}$ , the verifier will need to check that  $g_{i,j} = G(v_{i,j})$  and  $g_{i,j+1} = G(v_{i,j+1})$ .

**Optimization 2: Omit Re-Computable Values.** Moreover, since  $G$  is a public function, we do not need to include  $G(v_{i,j})$  in the transcript if we have included  $v_{i,j}$  in the response. Thus for the two views (corresponding to a single challenge) that the prover sends as part of the proof, we do not need to include the permutations of those views. We only need to include  $G(v_{i,(j+2)})$ , where  $v_{i,(j+2)}$  is the view that the prover does not open for the given challenge.

### 3.2 Seed Generation

We generate seeds for the random tapes using SHAKE with the private key, message, and public key as input and requesting the required number of output bytes from the XOF. Complete details are given in the specification document. Our current implementation and specification use deterministic signatures as a default to facilitate testing, however the specification shows how to randomize signatures by including additional entropy in the derivation process. The specification recommends randomizing signatures, especially when side-channel or fault attacks are a concern.

### 3.3 Random Tapes

We generate random tapes using the SHAKE XOF as a KDF to expand the seed to the required number of output bytes. Complete details are given in the specification document.

### 3.4 Challenge Generation

For both the FS and Unruh transform the challenge is computed with a hash function. For Picnic-FS and Picnic-UR the function is  $H : \{0, 1\}^* \rightarrow \{0, 1, 2\}^t$  (implemented using SHAKE) and rejection sampling: we split the output bits in pairs of two bits and reject all pairs with both bits set. For Picnic3 the range of the function is a set and a list of integers whose size depends on the parameter set. The same rejection sampling method is used.

### 3.5 Function $G$

As explained in Section 3.1, the function  $G$  used in Picnic-UR may be implemented with a hash function with the same domain and range. We implement  $G(x)$  with SHAKE, where the requested number of output bits is  $|x|$ .

## 4 Choice of Parameters

In this section we explain our parameter selection and rationale.

## 4.1 Choice of LowMC and SHAKE

The signature size depends on constants that are close to the security expectation. The only exceptions are the number of binary multiplication gates, and the size of the ring, which both depend on the choice of the primitive. In this section we compare LowMC to existing standardized primitives and to other primitives with a low number of multiplications.

**Standardized Primitives.** The smallest known Boolean circuit for AES-128 needs 5440 AND gates, AES-192 needs 6528 AND gates, and AES-256 needs 7616 AND gates [BMP13]. An AES circuit in  $\mathbb{F}_{2^4}$  might be more efficient in our setting, as in this case the number of multiplications is lower than 1000 [CGP<sup>+</sup>12]. This results in an impact on the signature size that is equivalent to 4000 AND gates. In [dDOS19] it was shown that generalizing KKW from binary to arithmetic circuits (over  $\mathbb{F}_{2^8}$  for AES), and computing S-Boxes using techniques from the MPC literature, one could build a Picnic-like signature using AES, at a fraction of the signature size resulting from a direct application of KKW to a binary AES circuit. However, signature sizes are still 2.5, 3.2 and 2.9 times larger than Picnic3 at security levels L1, L3 and L5. At levels L3 and L5 one must use either two calls to AES, or Rijndael, since the block size of AES is fixed to 128 bits. Even though collision resistance is not required, hash functions like SHA-256 are a popular choice for proof-of-concept implementations. The number of AND gates of a single call to the SHA-256 compression function is about 25000 and a single call to the permutation underlying SHA-3 is 38400.

**Lightweight Ciphers.** Most early designs in this domain focused on small area when implemented in hardware where an XOR gate is by a small factor larger than an AND or NAND gate. Notable designs with a low number of AND gates at the 128-bit security level are the block ciphers Noekeon [DPVAR00] (2048 ANDs) and Fantomas [GLSV14] (2112 ANDs). Furthermore, one should mention Prince [BCG<sup>+</sup>12] (1920 ANDs), or the stream cipher Trivium [DP08] (1536 AND gates to compute 128 output bits, with 80-bit security).

**Custom Ciphers with a Low Number of Multiplications.** Motivated by applications in SHE/FHE schemes, MPC protocols and SNARKs, recently a trend to design symmetric encryption primitives with a low number of multiplications or a low multiplicative depth started to evolve. This is a trend we can take advantage of.

We start with the LowMC [ARS<sup>+</sup>15a] block cipher family. In the most recent version of the design [ARS<sup>+</sup>16], the number of AND gates can be below 500 for 80-bit security, below 800 for 128-bit security, and below 1400 for 256-bit security. The stream cipher Kreyvium [CCF<sup>+</sup>16] (similarly to Trivium) needs 1536 AND gates to compute 128 output bits, but offers a higher security level of 128 bits. Even though FLIP [MJSC16] was designed to have especially low depth, it needs hundreds of AND gates per bit and is hence not competitive in our setting.

Last but not least there are the block ciphers and hash functions around MiMC [AGR<sup>+</sup>16] which need less than  $2 \cdot s$  multiplications for  $s$ -bit security in a field of size close to  $2^s$ . Note that MiMC is the only design in this category which aims at minimizing multiplications in a field larger than  $\mathbb{F}_2$ . However, since the size of the signature depends on both the number of multiplications and the size of the field, this leads to a factor  $2s^2$  which, for all arguably secure instantiations of MiMC, is already larger than the number of AND gates in the AES circuit.

LowMC has two important advantages over other designs: It has the lowest number of AND gates for every security level: The closest competitor Kreyvium needs about twice as many AND gates and only exists for the 128-bit security level. The fact that it allows for an easy parameterization of the security level is another advantage. We hence use LowMC for our concrete proposal.

The category of ciphers optimized for MPC or FHE applications continues to grow, and new designs have been proposed since this survey was written in 2017. However, being new, these designs do not yet have significant cryptanalysis supporting their security, and do not offer significant performance benefits in the context of Picnic.

**Hashing with SHAKE.** Keccak is a family of cryptographic primitives including hash functions and extensible output functions (XOF). It was selected as the successor to SHA-2 and was standardized as SHA3 [NIS15]. SHAKE is an XOF constructed from SHA3. Since both SHA3 and SHAKE have a large number of AND gates (as described above), we do not use them for Picnic key generation.

However, other parts of the ZKB++ and KKW protocols require a hash function. We will use SHAKE in two modes

1. as a hash function with a fixed output length

2. as a key derivation function, where we expand a fixed length seed into a larger pseudorandom value, by requesting larger, variable sized SHAKE outputs.

For the Picnic3 parameter sets, in our current implementation hashing with SHAKE accounts for roughly 57% of signing runtime and 72% of verification time (at security level L1). If in the future, a faster hash function becomes widely viewed as providing sufficient security (an example candidate is KangarooTwelve [BDP<sup>+</sup>18]), it may be worth considering using it with Picnic3. See [KZ20b] for more discussion and supporting benchmarks. Similarly, this suggests that hardware support for SHAKE can significantly improve timings for Picnic3. Additionally, some informal experiments with hardware accelerated SHAKE on IBM’s z architecture showed a significant speedup in signing and verification times (roughly 1.5x).

## 4.2 LowMC Parameters

To minimize the number of AND gates for a given key length  $k$  and data complexity  $d$ , we want to minimize  $r \cdot m$  (where  $r$  is the number of rounds and  $m$  is the number of S-boxes). One strategy would be to set  $m = 1$ , and to look for an  $n$  that minimizes  $r$ . Examples of such an approach are already given in the document describing version 2 of the LowMC design [ARS<sup>+</sup>16]. In our setting, this approach may not lead to the best results, as discussed below. Whenever we want to instantiate our signature scheme with LowMC with  $\kappa$ -bit quantum security, we set  $k = n \approx 2 \cdot \kappa$ . This matches AES, and the security levels in the NIST call for proposals. We also search for instances with data complexity  $d = 1$ , since for a given key the adversary only ever sees a single plaintext-ciphertext pair (namely, the public key in the Picnic scheme).<sup>3</sup>

Table 1 gives the LowMC instances we use in Picnic. Instances in the first half of the table use a partial S-box layer, and those in the the second half use a full S-box layer.

The original search for LowMC instances in Picnic [CDG<sup>+</sup>17] used the two following constraints: 1) the key and block size must be 128, 192, and 256 bits, and 2) the S-box layer should be small, since each S-box increases the number of AND gates, directly increasing the signature size. For these

---

<sup>3</sup> $d$  is given in units of  $\log_2(n)$ , where  $n$  is the number of pairs. Thus setting  $d = 1$  corresponds to two pairs, which is exactly what we need for our signature schemes.

reasons, the search started with a single S-box per round, which gives the smallest signatures, and then increased the number of S-boxes until a good balance between signature size and runtime was found. Instances with a full S-box layer were overlooked.

When LowMC has a full S-box layer, while each round has a large number of AND gates, far fewer rounds are required to provide equivalent security<sup>4</sup>. For example, at L1, four rounds are the minimum necessary according to the analysis of [ARS<sup>+</sup>15b] (and third-party cryptanalysis [DLMW15, DEM16, RST18, GKRS20])<sup>5</sup> and even when using five rounds (adding an additional round of security margin), the total number of ANDs is about the same as the baseline instances (see [KZ20b, Table 5], 7-18% more with 5 rounds, and about 17% less with 4 rounds). About 40-50% of the data in a signature size grows with the number of AND gates (see Tables 11 and 12 in [KZ20b]), so a small change in the number of AND gates will not have a large impact on signature size.

Of course it is also not possible to have a full S-box layer when  $n = 128$ ; since each S-box operates on three bits of state, the state size must be a multiple of three. Therefore, at L1, we choose to use a 129-bit state. At L3, since 192 is divisible by three, no special accommodations are required, and at L5, we chose a 255-bit state and key.

Using a full S-box layer impacts implementations in two other ways. First, the optimization technique of Dinur et al. [DKP<sup>+</sup>19b] no longer applies, since it requires  $s \ll n$ , and as the number of S-boxes increases, the performance improvement decreases. This means that each round in a full S-box instance is considerably more expensive, but also much simpler to implement efficiently. The second impact is the size of precomputed constants required to implement LowMC decreases; recall from Section 2.10 that each additional round requires a linear layer matrix ( $n^2$  bits), a round constant ( $n$  bits) and a key matrix ( $n^2$  bits). The amount of pre-computed data used by the partial and full S-box instances is compared in [KZ20b, Table 5], the full instances are found to use roughly half the amount of precomputed data. For these reasons, we expect FPGA implementations of Picnic [KRR<sup>+</sup>20] would especially benefit from using LowMC instances with a full S-box layer.

---

<sup>4</sup>For an intuition of why, consider attacks that can utilize the identity part of the S-box to skip the first and last round of the cipher.

<sup>5</sup>To determine the number of rounds for a given  $n, s$  and  $d$ , use `determine_rounds.py` from the LowMC source package [Tie17], which is updated to include all known cryptanalysis.

Parameters	Blocksize	S-boxes	Keysize	Rounds
	$n$	$m$	$k$	$r$
L1	128	10	128	20
L3	192	10	192	30
L5	256	10	256	38
L1-full	129	43	129	4
L3-full	192	64	192	4
L5-full	255	85	255	4

Table 1: Parameters for LOWMC targeting security levels L1, L3 and L5. The first half of the table contains parameters with a partial S-box layer ( $3m < n$ ), and the second half contains parameters with a full S-box layer ( $3m = n$ ). All parameters are computed for data complexity  $d = 1$ .

Picnic3 uses LowMC parameters with a full S-box layer exclusively. In Version 3 of the Picnic specification, an additional group of parameter sets were added where Picnic-FS uses a full S-box layer (called `picnic-L1-full`, `picnic-L3-full`, and `picnic-L5-full`).

### 4.3 MPC Parameters

**Parameters for ZKB++.** A single repetition of ZKB++ has a soundness error of  $2/3$ , which means that we need to perform parallel repetitions to achieve the desired soundness error. Hence we need 219 parallel repetitions for 128-bit classical security ( $(3/2)^{219} \geq 2^{128}$ ). For 128-bit PQ security, we must set our repetition count to  $t := 438$ . This is double the repetition count required for classical security due to Grover’s algorithm [Gro96]. The required number of repetitions for the L1, L3 and L5 security levels are given in Table 2.

level	# parallel repetitions
L1	219
L3	329
L5	418

Table 2: Number of parallel repetitions required at each security level.

**Parameters for KKW.** In the KKW protocol, there are multiple param-



eters to be chosen, including the number of emulated parties  $n$ , the number of total executions  $M$ , and the number of online executions checked by the verifier  $\tau$ . The soundness error  $\epsilon$ , depends on the choice of  $(M, n, \tau)$ , and is computed with the following equation.

$$\epsilon(M, n, \tau) = \max_{M-\tau \leq k \leq M} \left\{ \frac{\binom{k}{M-\tau}}{\binom{M}{M-\tau} \cdot n^{k-M+\tau}} \right\}.$$

Therefore, there are many choices of  $(M, n, \tau)$  for each security level, and each choice implies different trade offs in computation time and signature size. While the size is easy to estimate given  $(M, n, \tau)$ , the runtime is not, and we found the only reliable way is benchmark an optimized implementation and compare. This was done thoroughly in [KZ20b], we provide a brief discussion here.

Previously, our specification fixed  $n = 64$ , since generally larger  $n$  gives smaller signatures, and it allows implementations to pack the bit shares of each party into a 64-bit word, and operate on them with machine word operations. However, emulating 64 parties requires a large amount of hashing as each party in each MPC instance requires random tapes, and creates commitments. After benchmarking multiple options, we chose  $n = 16$ .

Our choices of  $M$  and  $\tau$  move away from the smallest possible signature size (i.e., the smallest  $\tau$ ), order to reduce  $M$ . We found that for a small increase in size,  $M$  is significantly reduced, which significantly reduces the CPU cost of signing and verification.

## 4.4 Alternative Parameters

In this section we list some possible alternatives to the parameters we chose, and discuss the tradeoffs. These apply to the Picnic3 parameters, since the KKW protocol has more parameters than the ZKB++ protocol. Some alternative parameters for LowMC, that apply to all parameter sets, were discussed in Section 4.2 and [CDG<sup>+</sup>17].

**Speed-Size Tradeoffs** As detailed in [KZ20b], there are multiple parameters that can be adjusted to reduce the time required for signing and verification, at the expense of larger signatures. One could view the endpoints of the curve as Picnic3 and Picnic-FS, and gradually increase signature size while decreasing runtimes.

We can choose a smaller  $M$ , provided we increase  $\tau$  to maintain the soundness error  $\epsilon$ . Larger  $\tau$  increases the signature size. An alternative size/speed tradeoff can be made by changing  $n$ , the number of parties. Using larger  $n$  increases the computational cost, but can let us decrease  $\tau$  to have shorter signatures (up to a point, e.g., at L1 we can increase  $n$  to 3060 and have signatures that are about 10KB, but further increasing  $n$  does not decrease signature size further).

**Using a 5-round protocol** The KKW interactive proof protocol can naturally be done as a five round protocol (see [KKW18, Figure 1]). In the first round the prover commits to the offline portion of many MPC instances, the verifier selects a subset, then the prover commits to the online simulations (using the selected subset), and the verifier selects one party from each instance to remain unopened. The potential benefit of this approach is that the prover needs to perform the online MPC simulation for a much smaller number of instances (roughly  $\tau$  instead of  $M$ ).

However, when making this non-interactive, the soundness error is given by a different formula (since both challenges must be sufficiently large), and we must use different values of  $M$  and  $\tau$ . We investigated a five-round variant of Picnic in [KZ20a], and concluded that performance is better with the three-round option.

## 4.5 Recommended Parameters

There are twelve parameter sets in the Picnic specification, and we provide the following guidance on which to use.

- For applications where speed is more important than signature size:  
Use Picnic-full (Picnic-L1-full, Picnic-L3-full, and Picnic-L5-full).
- For applications where signature size is more important than speed:  
Use Picnic3 (Picnic3-L1, Picnic3-L3, and Picnic3-L5).

The other six parameter sets Picnic\*-FS and Picnic\*-UR are moved to historical status, and not recommended for use or standardization.

**Rationale** The UR parameter sets (using the Unruh transform) were originally defined for their QROM analysis. Since then our understanding of the Fiat-Shamir transform security in the QROM has progressed, and we now have a QROM security proof for Picnic3 and Picnic-full, which are much

faster and shorter than the UR variant. Both proofs are non-tight, giving only asymptotic guarantees, but the analysis of the UR parameter sets was also non-tight.

The original **Picnic-\*\*-FS** parameter sets are nearly the same as their **Picnic-full** counterparts, except for the LowMC instance, but the latter is significantly faster with shorter signatures. Having a common set of LowMC instances across all six recommended parameter sets makes implementations simpler (since they need fewer constants), and helps to focus cryptanalysis.

## 5 Formal Security Analysis

This section contains formal security analyses of the **Picnic-FS** and **Picnic-UR** schemes (**Picnic3** is analyzed in §6). We begin in Section 5.1 by analyzing the (interactive) **ZKB++** protocol. We then separately analyze **Picnic-FS** and **Picnic-UR**, which differ in the transform applied to **ZKB++**, in the sense of unforgeability. As we discuss in Section 5.4, the scheme can actually be shown to satisfy *strong* unforgeability. The **Picnic-FS** scheme was also proven secure in the QROM in [DFMS19a], see the discussion in Section 6.4.

### 5.1 Security Analysis of **ZKB++**

First, we observe that not including output shares and commitments are what we call equivalence transformations: there is a transformation (which anyone can compute) which takes a **ZKBoo** proof and removes output shares and commitments, or which takes a proof without the output shares and commitments and produces a proof which does again include these values. Thus removing these values does not reveal any more information or make it any easier to forge proofs.

The other modification we make is to generate the initial shares pseudorandomly. Note that this cannot make it easier to forge proofs, because we are only reducing the options the prover has in choosing the shares. On the other hand, we note that the 2-privacy simulator for the decomposition works even if the initial random shares are generated pseudorandomly, so the zero-knowledge proof still goes through. Again, after this step, removing the input shares is an equivalence transformation that has no effect on security.

Thus, we only have to show that including no additional randomness in the commitments preserves completeness, 3-special soundness, and special

honest-verifier zero-knowledge of ZKB++.

First, we observe that completeness is clearly not impacted and the corollary below follows from this and [GMO16b, Proof of Proposition 2].

**Corollary 5.1.** *The modified version of ZKBoo—where the commitments no longer contain additional randomness—is complete, i.e., ZKB++ is complete.*

Second, under the observation that removing the randomness in the commitments does not impact the binding property of the commitments we can derive the following corollary from [GMO16b, Proof of Proposition 2].

**Corollary 5.2.** *The modified version of ZKBoo—where the commitments no longer contain additional randomness—is 3-special sound, i.e., ZKB++ is 3-special sound.*

What remains is to prove the following theorem.

**Theorem 5.3.** *The modified version of ZKBoo—where the commitments no longer contain additional randomness—is special honest-verifier zero-knowledge in the random oracle model, i.e., ZKB++ is special honest-verifier zero-knowledge in the (quantum) random oracle model.*

Before we prove the theorem, we recall that—as the challenge can be determined a priori in the proof for special honest-verifier zero-knowledge—we can use the 2-privacy simulator of the (2,3)-decomposition underlying ZKB++ (cf. Section 2.7 for details) to produce satisfying transcripts for the two views which need to be opened according to the challenge. Now, in the original proof [GMO16b, Proof of Proposition 2], the hiding property of the commitment which is not required to be opened ensures that the simulation works out. We have to argue that this still holds when no additional randomness is included in the commitments. Since we already use the random oracle heuristic for our signature scheme, we also rely on the random oracle heuristic for the subsequent proof.

*Proof (Sketch).* Recall the modification discussed in Section 2.8. We consider the random oracle model, in which  $H', H''$  are independent random oracles. Then, consider the seed  $k_i$  for each the unopened view: this  $k_i$  is only used as input to  $H', H''$ . So the initial prover is distributed identically to another prover which replaces  $H'(k_i), H''(k_i)$  for the unopened views with random values  $Z, Z'$ . Now,  $Z'$  is a random value which is only used as input to the commitment, so we can apply the same HVZK argument as ZKBoo.  $\square$

## 5.2 Security Analysis of Picnic-FS

In this section we prove security of Picnic-FS in the ROM. For discussion of security of Picnic-FS in the QROM, see Section 6.4. If we view ZKB++ as a canonical identification scheme that is secure against passive adversaries one just needs to keep in mind that most definitions are tailored to (2-)special soundness, and the 3-special soundness of ZKB++ requires an additional rewind. In particular, an adapted version of the proof of [Kat10, Theorem 8.2], which considers this additional rewind, attests the security of Picnic-FS. We obtain the following:

**Corollary 5.4.** *Picnic-FS instantiated with ZKB++ and a secure one-way function yields an EUF-CMA secure signature scheme in the ROM.*

However, we actually aim for a stronger result, i.e., sEUF-CMA, which also excludes malleability of the signatures. To show that Picnic-FS also satisfies this property, we need to view ZKB++ as a  $\Sigma$ -protocol which is transformed to its non-interactive counterpart via the FS transform and show that this protocol is actually simulation extractable. We base our argumentation upon the argumentation of [FKMV12a] to confirm simulation extractability. What we have to do is to show that the FS transformed ZKB++ is zero-knowledge and provides quasi-unique responses. We do so by proving two lemmas. Combining those lemmas with [FKMV12a, Theorem 2 and Theorem 3] then yields simulation extractability as a corollary.

**Lemma 5.5.** *Let  $Q_H$  be the number of queries to the random oracle  $H$ ,  $Q_S$  be the overall queries to the simulator, and let the commitments be instantiated via a RO  $H'$  with output space  $\{0, 1\}^\rho$  and the committed values having min entropy  $\nu$ . Then the probability  $\epsilon(\kappa)$  for all PPT adversaries  $\mathcal{A}$  to break zero-knowledge of  $\kappa$  parallel executions of the FS transformed ZKB++ is bounded by  $\epsilon(\kappa) \leq s/2^\nu + (Q_S \cdot Q_H)/2^{3 \cdot \rho}$ .*

The subsequent proof is similar to the general results for  $\Sigma$ -protocols from [FKMV12a], yet we have to account for the additional challenge that the simulator only outputs transcripts which are statistically close to original transcripts (which is in contrast to the identically distributed transcripts in [FKMV12a]). Furthermore, we also provide concrete bounds.

*Proof.* We bound the probability of any PPT adversary  $\mathcal{A}$  to win the zero-knowledge game by showing that the simulation of the proof oracle is sta-

tistically close to the real proof oracle. For our proof let the environment maintain a list  $H$  where all entries are initially set to  $\perp$ .

**Game 0:** The zero-knowledge game where the proofs are honestly computed, and the ROs are simulated honestly.

**Game 1:** As Game 0, but whenever the adversary requests a proof for some tuple  $(x, w)$  we choose  $\mathbf{e} \xleftarrow{R} \{0, 1, 2\}^\kappa$  before computing  $\mathbf{a}$  and  $\mathbf{z}$ . If  $H[(\mathbf{a}, x)] \neq \perp$  we abort and call that event  $E$ . Otherwise, we set  $H[(\mathbf{a}, x)] \leftarrow \mathbf{e}$ .

*Transition - Game 0  $\rightarrow$  Game 1:* Game 0 and Game 1 proceed identically unless  $E$  happens. The message  $\mathbf{a}$  includes 3 RO commitments with respect to  $H'$ , i.e., a lower bound for the min-entropy is  $3 \cdot \rho$ . We have that  $|\Pr[S_0] - \Pr[S_1]| \leq (Q_S \cdot Q_H)/2^{3 \cdot \rho}$ .

**Game 2:** As Game 1, but we compute the commitments in  $\mathbf{a}$  so that the commitments which will never be opened according to  $\mathbf{e}$  contain random values.

*Transition - Game 1  $\rightarrow$  Game 2:* The statistical difference between Game 1 and Game 2 can be upper bounded by  $|\Pr[S_1] - \Pr[S_2]| \leq \kappa \cdot 1/2^\nu$  (for compactness we collapsed the  $s$  game changes into a single game).

**Game 3:** As Game 2, but we use the HVZK simulator to obtain  $(\mathbf{a}, \mathbf{e}, \mathbf{z})$ .

*Transition - Game 2  $\rightarrow$  Game 3:* This change is conceptual, i.e.,  $\Pr[S_2] = \Pr[S_3]$ .

In Game 0, we sample from the first distribution of the zero-knowledge game, whereas we sample from the second one in Game 3; the distinguishing bounds shown above conclude the proof.  $\square$

**Lemma 5.6.** *Let the commitments be instantiated via a RO  $H'$  with output space  $\{0, 1\}^\rho$  and let  $Q_{H'}$  be the number of queries to  $H'$ , then the probability to break quasi-unique responses is bounded by  $Q_{H'}^2/2^\rho$ .*

*Proof.* To break quasi-unique responses, the adversary would need to come up with two valid proofs  $(\mathbf{a}, \mathbf{e}, \mathbf{z})$  and  $(\mathbf{a}, \mathbf{e}, \mathbf{z}')$ . The last message  $\mathbf{z}$  (resp  $\mathbf{z}'$ ) only contains openings to commitments, meaning that breaking quasi unique responses implies finding a collision for at least one of the commitments. The probability for this to happen is upper bounded by  $Q_{H'}^2/2^\rho$  which concludes the proof.  $\square$

Combining Lemma 5.5 and Lemma 5.6 with [FKMV12a, Theorem 2 and Theorem 3] yields the following corollary.

**Corollary 5.7.** *The FS transformed ZKB++ protocol is simulation extractable.*

### 5.3 Security Analysis of Picnic-UR

Here we prove that the proof system we get by applying our modified Unruh transform to ZKB++ is both zero knowledge and simulation extractable in the quantum random oracle model.

Before we begin, we note that the quantum random oracle model is highly non-trivial, and a lot of the techniques used in standard random oracle proofs do not apply. The adversary is a quantum algorithm that may query the oracle on quantum inputs which are a superposition of states and receive superposition of outputs. If we try to measure those states, we change the outcome, so we do not for example have the same ability to view the adversary's input and program the responses that we would in the standard ROM.

Here we rely on lemmas from Unruh's work on quantum-secure Fiat-Shamir like proofs [Unr15]. We follow his proof strategy as closely as possible, modifying it to account for the optimizations we made and the fact that we have only 3-special soundness in our underlying  $\Sigma$ -protocol.

**Zero-Knowledge.** This proof very closely follows the proof from [Unr15]. The main difference is that we also use the random oracle to form our commitments, which is addressed in the transition from game 2 to game 3 below.

Consider the simulator described in Figure 3. From this point on we assume for simplicity of notation that  $\text{View}_3$  includes  $x_3$ .

We proceed via a series of games.

**Game 1:** This is the real game in the quantum random oracle model. Let  $H_{com}$  be the random oracle used for forming the commitments,  $H_{chal}$  be the random oracle used for forming the challenge, and  $G$  be the additional random permutation.

**Game 2:** We change the prover so that it first chooses random  $e^* = e^{*(1)}, \dots, e^{*(t)}$ , and then on step 2, it programs  $H_{chal}(a^{(1)}, \dots, a^{(t)}, h^{(1)}, \dots, h^{(t)}) = e^*$ .

$p \leftarrow \text{Sim}(x)$ : In the simulator, we follow Unruh, and replace the initial state (before programming) of the random oracles with random polynomials of degree  $2q - 1$  where  $q$  is an upper bound on the number of queries the adversary makes.

1. For  $i \in [1, t]$ , choose random  $e^{(i)} \leftarrow \{1, 2, 3\}$ . Let  $e$  be the corresponding binary string.
2. For each iteration  $r_i, i \in [1, t]$ : Sample random seeds  $k_{e^{(i)}}^{(i)}, k_{e^{(i)}+1}^{(i)}$  and run the circuit decomposition simulator to generate  $\text{View}_{e^{(i)}}^{(i)}, \text{View}_{e^{(i)}+1}^{(i)}$ , output shares  $y_1^{(i)}, y_2^{(i)}, y_3^{(i)}$ , and if  $e^{(i)} = 1$   $x_3^{(i)}$ .

For  $j = e^{(i)}, e^{(i)} + 1$  commit  $[C_j^{(i)}, D_j^{(i)}] \leftarrow [H(k_j^{(i)}, \text{View}_j^{(i)}), k_j^{(i)} || \text{View}_j^{(i)}]$ , and compute  $g_j^{(i)} = G(k_j^{(i)}, \text{View}_j^{(i)})$ .

Choose random  $C_{e^{(i)}+2}, g_{e^{(i)}+2}^{(i)}$

Let  $a^{(i)} = (y_1^{(i)}, y_2^{(i)}, y_3^{(i)}, C_1^{(i)}, C_2^{(i)}, C_3^{(i)})$ . And  $h^{(i)} = (g_1^{(i)}, g_2^{(i)}, g_3^{(i)})$ .

2. Set the challenge: program  $H(a^{(1)}, \dots, a^{(t)}) := e$ .
3. For each iteration  $r_i, i \in [1, t]$ : let  $b^{(i)} = (y_{e^{(i)}+2}^{(i)}, C_{e^{(i)}+2}^{(i)})$  and set

$$z^{(i)} \leftarrow \begin{cases} (\text{View}_2^{(i)}, k_1^{(i)}, k_2^{(i)}) & \text{if } e^{(i)} = 1, \\ (\text{View}_3^{(i)}, k_2^{(i)}, k_3^{(i)}, x_3^{(i)}) & \text{if } e^{(i)} = 2, \\ (\text{View}_1^{(i)}, k_3^{(i)}, k_1^{(i)}, x_3^{(i)}) & \text{if } e^{(i)} = 3. \end{cases}$$

4. Output  $p \leftarrow [e, (b^{(1)}, z^{(1)}), (b^{(2)}, z^{(2)}), \dots, (b^{(t)}, z^{(t)})]$ .

Figure 3: The zero knowledge simulator

Note that each the  $a^{(1)}, \dots, a^{(t)}, h^{(1)}, \dots, h^{(t)}$  has sufficient collision-entropy, since it includes  $\{h^{(i)} = (g_1^{(i)}, g_2^{(i)}, g_3^{(i)})\}$ , the output of a permutation on input whose first  $k$  bits are chosen at random (the  $k_j^{(i)}$ ), so we can apply Corollary 11 from [Unr15] (using a hybrid argument)



to argue that Game 1 and Game 2 are indistinguishable.

**Game 3:** We replace the output of each  $H_{com}(k_{e*(i)}, \text{View}_{e*(i)})$  and  $G(k_{e*(i)}, \text{View}_{e*(i)})$  with a pair of random values.

First, note that  $H_{com}$  and  $G$  are always called (by the honest party) on the same inputs, so we will consider them as a single random oracle with a longer output space, which we refer to as  $H$  for this proof.

Now, to show that Games 2 and 3 are indistinguishable, we proceed via a series of hybrids, where the  $i$ -th hybrid replaces the first  $i$  such outputs with random values.

To show that the  $i$ -th and  $i + 1$ -st hybrid are indistinguishable, we rely on Lemma 9 from [Unr15]. This lemma says the following: For any quantum  $A_0, A_1$  which make  $q_0, q_1$  queries to  $H$  respectively and classical  $A_C$ , all three of which may share state, let  $P_C$  be the probability if we choose a random function  $H$  and a random output  $B$ , then run  $A_0^H$  followed by  $A_C$  to generate  $x$ , and then run  $A_1^H(x, B)$ , that for a random  $j$ , the  $j$ -th query  $A_1^H$  makes is measured as  $x' = x$ . Then as long as the output of  $A_C$  has collision-entropy at least  $k$ , the advantage with which  $A_1^H$ , when run after  $A_0, A_C$  as described, distinguishes  $(x, B)$  from  $(x, H(x))$  is at most  $(4 + \sqrt{2})\sqrt{q_0}2^{-k/4} + 2q_1\sqrt{P_C}$ .

In other words, if we can divide our game into three such algorithms and argue that the  $A_1$  queries  $H$  on something that collapses to  $x$  with only negligible probability, then we can conclude that the two games are indistinguishable. Let  $A_0$  run the game up until just before the  $i$  th iteration in the proof generation. Let  $A_C$  be the process which chooses  $k_1^{(i)}, k_2^{(i)}, k_3^{(i)}$  and generates  $\text{View}_1^{(i)}, \text{View}_2^{(i)}, \text{View}_3^{(i)}$ , and outputs  $x = k_{e*(i)}, \text{View}_{e*(i)}$ . (Note that this has collision entropy  $|k_{e*(i)}|$  which is sufficient.) Let  $A_1$  be the process which runs the rest of the proof, and then runs the adversary on the response.

Now we just have to argue that the probability that we make a measurement of  $A_1$ 's  $j$ -th query to  $H$  and get  $x$  is negligible. To do this, we reduce to the security of the PRG used to generate the random tapes (and hence the views). Note that besides the one RO query,  $k_{e*(i)}$  is only used as input to the PRG. So, suppose there exists a quantum adversary  $A$  for which the resulting  $A_1$  has non-negligible probability of making an  $H$ -query that collapses to  $x$ . Then we can construct a

quantum attacker for the PRG: we run the above  $A_0, A_C$ , but instead of choosing  $k_{e*(i)}$  we use the PRG challenge as the resulting random tape, and return a random value as the RO output. Then we run  $A_1$ , which continues the proof (which should query  $k_{e*(i)}$  only with negligible probability since  $ks$  are chosen at random), and then runs the adversary. We pick a random  $j$ , and on the adversary's  $j$ -th RO query, we make a measurement and if it gives us a seed consistent with our challenge tape, we output 1, otherwise a random bit. If  $P_C$  is non-negligible then we will obtain the correct seed and distinguish with non-negligible probability.

**Game 4:** For each  $i$  instead of choosing random  $k_{e*(i)}$  and expanding it via the PRG to get the random tape used to compute the views, we choose those tapes directly at random.

Note that in Game 3,  $k_{e*(i)}$  are now only used as seeds for the PRG, so this follow from pseudo-randomness via a hybrid argument.

**Game 5:** We use the simulator to generate the views that will be opened, i.e.  $j \neq e*(i)$  for each  $i$ . We note that now the simulator no longer uses the witness.

This is identical by perfect privacy of the circuit decomposition.

**Game 6:** To allow for extraction in the simulation-extractability game we replace the random oracles with random polynomials whose degree is larger than the number of queries the adversary makes. The argument here identical to that in [Unr15].

**Online Extractability.** Before we prove online simulation-extractability, we define some notation to simplify the presentation:

For any proof  $\pi = e, \{b^{(i)}, g^{(i)}, z^{(i)}\}_{i=1\dots t}$ , let  $\text{hash-input}(\pi) = \{a^{(i)}, h^{(i)} = (g_1^{(i)}, g_2^{(i)}, g_3^{(i)})\}$  be the values that the verifier uses as input to  $H_{chal}$  in the verification of  $\pi$  as described in Figure 4.

For a proof  $\pi = (e, \{b^{(i)}, g^{(i)}, z^{(i)}\}_{i=1\dots t})$ , let  $\text{open}_0(z^{(i)})$ ,  $\text{open}_1(z^{(i)})$  denote the values derived from  $z^{(i)}$  and used to compute  $C_{e_i}^{(i)}$  and  $C_{e_i+1}^{(i)}$  respectively in the description of **Ver** in Figure 4.

We say a tuple  $(a, j, (o_1, o_2))$  is valid if  $a = (y_1, y_2, y_3, C_1, C_2, C_3)$ ,  $C_j = H_{\text{com}}(o_1)$ ,  $C_{j+1} = H_{\text{com}}(o_2)$  and  $o_1, o_2$  consist of  $k$ , **View** pairs for player  $j, j+1$  that are consistent according to the circuit decomposition. We say

$(a, j, (O_1, O_2))$  is set-valid if there exists  $o_1 \in O_1$  and  $o_2 \in O_2$  such that  $(a, j, (o_1, o_2))$  is valid and set-invalid if not.

We first restate lemma 16 from [Unr15] tailored to our application, in particular the fact that our proofs do not explicitly contain the commitment but rather the information the verifier needs to recompute it.

**Lemma 5.8.** *Let  $q_G$  be the number of queries to  $G$  made by the adversary  $A$  and the simulator  $S$  in the simulation-extractability game, and let  $n$  be the number of proofs generated by  $S$ . Then the probability that  $A$  produces  $x, \pi^* \notin \text{simproofs}$  where  $x, \pi^*$  is accepted by  $\text{Ver}^H$ , and  $\text{hash-input}(\pi^*) = \text{hash-input}(\pi')$  for a previous proof  $\pi'$  produced by the simulator, is at most  $n(n+1)/2(2^{-\kappa})^{3t} + O((q_G+1)^3 2^{-\kappa})$  (Call this event **MallSim**.)*

*Proof.* This proof follows almost exactly as in [Unr15].

First, we argue that  $G$  is indistinguishable from a random function exactly as in [Unr15].

Then, observe that there are only two ways **MallSim** can occur:

Let  $e'$  be the hash value in  $\pi'$ . Then either  $S$  reprograms  $H$  sometime after  $\pi'$  is generated so that  $H(\text{hash-input}(\pi'))$  is no longer  $e'$ , or  $\pi^*$  also contains the same  $e$  as  $\pi$ , i.e.  $e = e'$ .  $S$  only reprograms  $H$  if it chooses the same **hash-input** in a later proof - and **hash-input** includes  $g_j^{(i)}$ , i.e. a random function applied to an input which includes a randomly chosen seed. Thus, the probability that  $S$  chooses the same **hash-input** twice is at most  $n(n+1)/2(2^{-\kappa})^{3t} + O((q_G+1)^3 2^{-\kappa})$ , where the first term is the probability that two proofs use all the same seeds, and the second term is the probability that two different seeds result in a collision in  $G$ , where the latter follows from Theorem 8 in [Unr15].

The other possibility is that  $\text{hash-input}(\pi^*) = \text{hash-input}(\pi')$ , and  $e = e'$ , but  $b^{(i)}, g^{(i)}, z^{(i)} \neq b'^{(i)}, g'^{(i)}, z'^{(i)}$  for some  $i$ . First note, that if  $e = e'$  and  $\text{hash-input}(\pi^*) = \text{hash-input}(\pi')$ , then  $g^{(i)} = g'^{(i)}$  and  $b^{(i)} = b'^{(i)}$  for all  $i$ , by definition of **hash-input**. Thus, the only remaining possibility is that  $z^{(i)} \neq z'^{(i)}$  for some  $i$ . But since  $h^{(i)} = h'^{(i)}$  for all  $i$ , this implies a collision in  $G$ , which again by Theorem 8 in [Unr15] occurs with probability at most  $O((q_G+1)^3 2^{-\kappa})$ .

We conclude that **MallSim** occurs with probability at most

$$n(n+1)/2(2^{-\kappa})^{3t} + O((q_G+1)^3 2^{-\kappa}). \quad \square$$

Here, next we present our variant of lemma 17 from [Unr15]. Note that this is quoted almost directly from Unruh with two modifications to account

for the fact that our proofs do not explicitly contain the commitment but rather the information the verifier needs to recompute it, and the fact that our underlying  $\Sigma$ -protocol has only 3 challenges and satisfies 3-special soundness.  $H_0$  in this lemma will correspond in our final proof to the initial state of  $H_{chal}$ , before any reprogramming.

**Lemma 5.9.** *Let  $G, H_{com}$  be arbitrarily distributed functions, and let  $H_0 : \{0, 1\}^{\leq \ell} \rightarrow \{0, 1\}^{2t}$  be uniformly random (and independent of  $G$ ). Then, it is hard to find  $x$  and  $\pi = e, \{b^{(i)}, g^{(i)}, z^{(i)}\}_{i=1\dots t}$  such that for  $\{a^{(i)}, (g_1^{(i)}, g_2^{(i)}, g_3^{(i)})\} = \text{hash-input}(\pi)$  and  $J_1 || \dots || J_t := H_0(\text{hash-input}(\pi))$*

- (i)  $g_{J_i}^{(i)} = G(\text{open}_0(z^{(i)}))$  and  $g_{J_{i+1}}^{(i)} = G(\text{open}_1(z^{(i)}))$  for all  $i$ .
- (ii)  $(a^{(i)}, J_i, (\text{open}_0(z^{(i)}), \text{open}_1(z^{(i)})))$  is valid for all  $i$ .
- (iii) For every  $i$ , there exists a  $j$  such that  $(a^{(i)}, j, G^{-1}(g_{i,j}), G^{-1}(g_{i,j+1}))$  is set-invalid.

More precisely, if  $A^{G, H_0}$  makes at most  $q_H$  queries to  $H_0$ , it outputs  $(x, \pi)$  with these properties with probability at most  $2(q_H + 1)(\frac{2}{3})^{t/2}$

*Proof.* Without loss of generality, we can assume that  $G, H_{com}$  are fixed functions which  $A$  knows, so for this lemma we only treat  $H_0$  as a random oracle.

For any given value of  $H_0$ , we call a tuple  $c = (x, \{a^{(i)}\}_i, \{g_j^{(i)}\}_{i,j})$  a candidate iff: for each  $i$ , among the three transcripts,  $(a^{(i)}, 1, G^{-1}(g_1^{(i)}), G^{-1}(g_2^{(i)}))$ ,  $(a^{(i)}, 2, G^{-1}(g_2^{(i)}), G^{-1}(g_3^{(i)}))$ , and  $(a^{(i)}, 3, G^{-1}(g_3^{(i)}), G^{-1}(g_1^{(i)}))$  at least one is set-valid, and at least one is set-invalid. Let  $n_{\text{twovalid}}(c)$  be the number of  $i$ 's for which there are 2 set-valid transcripts. Let  $E_{\text{valid}}(c)$  be the set of challenge tuples which correspond to only set-valid conversations. (Note that  $|E_{\text{valid}}(c)| = 2^{n_{\text{twovalid}}(c)}$ .) We call a candidate an  $H_0$ -solution if the challenge produced by  $H_0$  only opens set-valid conversations, i.e. in lies in  $E_{\text{valid}}(c)$ . We now aim to prove that  $A^H$  outputs an  $H_0$  solution with negligible probability.

For any given candidate  $c$ , for uniformly random  $H_0$ , the probability that  $c$  is an  $H_0$ -solution is  $\leq (\frac{2}{3})^t$ . In particular, for candidate  $c$  the probability is  $(\frac{2}{3})^t * 2^{n_{\text{twovalid}}(c)-t}$ .

Let **Cand** be the set of all candidates. Let  $F : \text{Cand} \rightarrow \{0, 1\}$  be a random function such that for each  $c$   $F(c)$  is i.i.d. with  $\Pr[F(c) = 1] = (2/3)^t$ .

Given  $F$ , we construct  $H_F : \{0, 1\}^* \rightarrow \mathbb{Z}_3^t$  as follows:

- For each  $c \notin \text{Cand}$ ,  $H_F(c)$  is set to a uniformly random  $y \in \mathbb{Z}_3^t$ .
- For each  $c \in \text{Cand}$  such that  $F(c) = 0$ ,  $H_F(c)$  is set to a uniformly random  $y \in \mathbb{Z}_3^t \setminus \mathbf{E}_{\text{valid}}(c)$ .
- For each  $c \in \text{Cand}$  with  $F(c) = 1$ , with probability  $2^{n_{\text{twovalid}}-t}$ , choose a random challenge tuple  $e$  from  $\mathbf{E}_{\text{valid}}(c)$ , and set  $H_F(c) := e$ . Otherwise  $H_F(c)$  is set to a uniformly random  $y \in \mathbb{Z}_3^t \setminus \mathbf{E}_{\text{valid}}(c)$ .

Note that for each  $c$ , and  $e$  the probability of  $H(c)$  being set to  $e$  is  $3^{-t}$ . Suppose  $A_0^H$  outputs an  $H_0$ -solution with probability  $\mu$ , then since  $H_F$  has the same distribution as  $H_0$ ,  $A^{H_F}()$  outputs an  $H_F$  solution  $c$  with probability  $\mu$ . By our definition of  $H_F$ , if  $c$  is an  $H_F$  solution, then  $F(c) = 1$ . Thus,  $A^{H_F}()$  outputs  $c$  such that  $F(c) = 1$  with probability at least  $\mu$ .

As in [Unr15], we can simulate  $A^{H_F}()$  with another algorithm which generates  $H_F$  on the fly, and thus makes at most the same number of queries to  $F$  that  $A$  makes to  $H_F$ . Thus by applying Lemma 7 from [Unr15], we get

$$\mu \leq 2(q_H + 1)\left(\frac{2}{3}\right)^{t/2}.$$

□

Finally, as the sigma protocol underlying our proofs is only computationally sound (because we use  $H_{\text{com}}$  for our commitment scheme), we need to argue that an extractor can extract from 3 valid transcripts with all but negligible probability.

**Lemma 5.10.** *There exists an extractor  $E_\Sigma$  such that for any PPT quantum adversary  $A$ , the probability that  $A$  can produce  $(a, \{(\nu_{1,j}, \nu_{2,j})\}_{j=1,2,3})$  such that  $(a, j, (\nu_{1,j}, \nu_{2,j}))$  is a valid transcript for  $j = 1, 2, 3$ , but  $E_\Sigma(a, \{(\nu_{1,j}, \nu_{2,j})\}_{j=1,2,3})$  fails to extract a proof, is negligible.*

*Proof.* Recall that  $a = (y_1, y_2, y_3, C_1, C_2, C_3)$ , and if all three transcripts are valid,  $C_j = H_{\text{com}}(\nu_{1,j}) = H_{\text{com}}(\nu_{2,j-1})$  for  $j = 1, 2, 3$ . Thus, either we have  $\nu_{1,j} = \nu_{2,j-1}$  for all  $j$  or  $\mathcal{A}$  has found a collision in  $H_{\text{com}}$ . But, Theorem 8 in [Unr15] tells us that the probability of finding a collision in a random function with  $k$ -bit output using at most  $q$  queries is at most  $O((q+1)^3 2^{-k})$ , which is negligible. If  $\nu_{1,j} = \nu_{2,j-1}$  for all  $j$ , then we have  $3 k_j || \text{View}_j$  values, all of which are pairwise consistent, so we conclude by the correctness of the circuit decomposition, and the fact that  $(x = y, w) \in R$  iff  $\phi(w) = y$  that if we sum the input share in  $\text{View}_1, \text{View}_2, \text{View}_3$ , we get a witness such that  $(x, w) \in R$ . □

**Theorem 5.11.** *Our version of the Unruh protocol satisfies simulation-extractability against a quantum adversary.*

*Proof.* We define the following extractor:

1. On input  $\pi$ , compute  $\text{hash-input}(\pi) = \{a^{(i)}, h^{(i)} = (g_1^{(i)}, g_2^{(i)}, g_3^{(i)})\}$
2. For  $i \in 1, \dots, t$ : For  $j \in 1, 2, 3$ , check whether there exists  $\nu_{1,j} \in G^{-1}(g_j^{(i)}), \nu_{2,j} \in G^{-1}(g_{j+1}^{(i)})$  such that  $(a^{(i)}, j, (\nu_{1,j}, \nu_{2,j}))$  is a valid transcript. If there is a valid transcript for all  $j$ , output  $E_\Sigma(a^{(i)}, \{(\nu_{1,j}, \nu_{2,j})\}_{j=1,2,3})$  as defined by Lemma 5.10 and halt.
3. If no solution is found, output  $\perp$ .

First we define some notation, again borrowed heavily from [Unr15]:

Let  $\text{Ev}_i, \text{Ev}_{ii}, \text{Ev}_{iii}$  be events denoting that  $A$  in the simulation-extractability game produces a proof satisfying conditions (i), (ii), and (iii) from Lemma 5.9 respectively.

Let  $\text{SigExtFail}$  be the event that the extractor finds a successful  $(a, \{(\nu_{1,j}, \nu_{2,j})\}_{j=1,2,3})$ , but  $E_\Sigma$  fails to produce a valid witness.

Let  $\text{ShouldExt}$  denote the event that  $A$  produces  $x, \pi$  such that  $\text{Ver}^H$  accepts and  $(x, \pi) \notin \text{simproofs}$ .

Then our goal is to prove that the  $w$  produced by the extractor is such that  $(x, w) \in R$ . I.e., we want to prove that the following probability is negligible.

$$\begin{aligned}
& \Pr[\text{ShouldExt} \wedge (x, w) \notin R] \\
& \leq \Pr[\text{ShouldExt} \wedge (x, w) \notin R \wedge \neg \text{MallSim}] + \Pr[\text{MallSim}] \\
& = \Pr[\text{ShouldExt} \wedge (x, w) \notin R \wedge \neg \text{MallSim} \wedge \neg \text{Ev}_{iii}] \\
& \quad + \Pr[\text{ShouldExt} \wedge (x, w) \notin R \wedge \neg \text{MallSim} \wedge \text{Ev}_{iii}] + \Pr[\text{MallSim}] \\
& \leq \Pr[(x, w) \notin R \wedge \neg \text{Ev}_{iii}] \\
& \quad + \Pr[\text{ShouldExt} \wedge (x, w) \notin R \wedge \neg \text{MallSim} \wedge \text{Ev}_{iii}] + \Pr[\text{MallSim}] \\
& = \Pr[\text{SigExtFail}] \\
& \quad + \Pr[\text{ShouldExt} \wedge (x, w) \notin R \wedge \neg \text{MallSim} \wedge \text{Ev}_{iii}] + \Pr[\text{MallSim}] \\
& = \Pr[\text{SigExtFail}] \\
& \quad + \Pr[\text{ShouldExt} \wedge (x, w) \notin R \wedge \neg \text{MallSim} \wedge \text{Ev}_i \wedge \text{Ev}_{ii} \wedge \text{Ev}_{iii}] \\
& \quad + \Pr[\text{MallSim}] \\
& \leq \Pr[\text{SigExtFail}] + \Pr[\text{Ev}_i \wedge \text{Ev}_{ii} \wedge \text{Ev}_{iii}] + \Pr[\text{MallSim}]
\end{aligned}$$

Here, the second equality follows from the definition of  $\text{SigExtFail}$  and  $\text{Ev}_{iii}$ , and the description of the extractor. The third equality follows from the fact that  $\neg \text{MallSim}$  means that the hash function on  $\text{hash-input}(\pi)$  has not been reprogrammed, and the fact that  $\text{ShouldExt}$  means verification succeeds, which means that conditions (i) and (ii) are satisfied.

Finally, by Lemmas 5.10, 5.9, and 5.8, we conclude that this probability is negligible.

## 5.4 Strong Unforgeability of Picnic-FS and Picnic-UR

We have shown that Picnic-FS and Picnic-UR are a simulation-extractable NIZK proof systems in the classical (resp. quantum) random oracle model against classical (resp. quantum) adversaries. Strong unforgeability (sEUF-CMA security) follows directly: this is a well known result in the classical model, and shown in [Unr15] in the quantum setting. For completeness, we briefly sketch this result:

Suppose there exist an adversary  $\mathcal{A}$  who can break the strong unforgeability property (cf. Definition 2.8). Then we can construct an adversary against the ZK property of the NIZK, the simulation-extractability property of the NIZK, or the one-wayness of the one-way function. We proceed through a series of games. In the first transition, we switch the signature algorithm to use the ZK simulator rather than the prover. This is indistinguishable by ZK, so the adversary will still produce forgeries with high probability, or we have a distinguisher which breaks the ZK property. Then, when the adversary produces a valid forgery, we run the extractor to produce a pre-image of the one-way function. If this extractor does not succeed with high probability whenever the adversary produces a forgery, we break simulation-extractability. Note here, that our extractor is guaranteed to work as soon as either the statement or the proof is different from what the simulator produced, so we will be able to extract from new signatures on previously signed messages as required in strong unforgeability. Otherwise, we have produced a pre-image given only the output of the one-way function (recall that we use the simulator to sign, so we do not need the pre-image there), so we break the one-wayness property.

## 6 Formal Security Analysis of Picnic3

In this section we provide a formal analysis of the Picnic3 signature schemes, starting with the security of the underlying MPC protocol, then proving unforgeability in the random oracle model. We end with a brief analysis of two optimizations and discuss security in the quantum random oracle model.

### 6.1 Proof of Security of the Underlying MPC Protocol

The protocol is described in Section 2.9.1, here we prove it is secure against an all-but-one corruption in the semi-honest model. Here, the SHAKE XOF that is used to expand seeds to random tapes is modelled as a PRG.

**Lemma 6.1.** *Suppose there exists a  $(t, \epsilon_{PRG})$ -PRG. Then there exists a simulator for the MPC protocol of §2.9.1 such that no distinguisher running in time  $t$  can distinguish between the real-world execution and ideal-world execution defined by this simulator with better than  $\epsilon_{PRG}$  probability.*

*Proof.* We first describe a simulator  $\text{Sim}_P(1^\kappa, y, C)$  that outputs the view of all parties except for  $P$ . Denote the input and output sizes of  $C$  by  $m$  and  $l$  respectively. We use  $x \leftarrow X$  to denote choosing a value from  $X$  at random and assigning it to  $x$ . The simulator works as follows:

1. If  $P = n$ , set  $\text{state}_i \leftarrow \{0, 1\}^k$  for all  $i \neq P$ . Otherwise, set  $\text{state}_i \leftarrow \{0, 1\}^k$ , for  $i \notin \{n, P\}$  and set  $\text{state}_n \leftarrow \{0, 1\}^{k+|C|}$ .
2. Pick  $\hat{z} \leftarrow \{0, 1\}^m$ ,  $\text{msgs}_P \leftarrow \{0, 1\}^{|C|}$ .
3. Use  $\{\text{state}_i\}_{i \neq P}$ ,  $\hat{z}$  and  $\text{msgs}_P$  to simulate the online phase of the MPC protocol until the output reconstruction step, such that the simulator obtains the shares of outputs  $[y]$  for  $i \neq P$ , denoted as  $[y]_i$ . Compute  $[y]_P := \bigoplus_{i \neq P} [y]_i \oplus y$ . Append  $[y]_P$  to  $\text{msgs}_P$ .

**Hybrid<sub>1</sub>.** Same as the real-world protocol, except use true randomness, instead of seed-derived, for party  $P$ . String  $\text{aux}$  is computed as described in the protocol, based on the true randomness.

It is easy to see that the probability of distinguishing **Hybrid<sub>1</sub>** and the real-world protocol in running time  $t$  is no more than  $\epsilon_{PRG}$ .



**Hybrid<sub>2</sub>.** Replace **aux** in **Hybrid<sub>1</sub>** by a uniformly random string of the same length.

If  $P = n$ , then **aux** is not part of the view of the adversary; if  $P \neq n$ , then bits of **aux** are computed by XORing one bit of randomness from each seed from party  $i \neq P$ , then XORing one bit of randomness from party  $P$  (which is uniformly random in **Hybrid<sub>1</sub>**). Therefore **aux** is uniformly random in **Hybrid<sub>1</sub>**.

Therefore, **Hybrid<sub>1</sub>** and **Hybrid<sub>2</sub>** are identical.

**Hybrid<sub>3</sub>.** Same as **Hybrid<sub>2</sub>**, except that  $\hat{z}$  is changed to uniformly random string; The last message from party  $P$  is replaced by a message computed from the output as defined in the simulator. In more detail, use  $\{\text{state}_i\}_{i \neq P}$ ,  $\hat{z}$  and  $\text{msgs}_P$  to simulate the online phase of the MPC protocol locally, such that in the end, the simulator obtains share of outputs  $[y]_i$  for  $i \neq P$ . Compute  $[y]_P := \bigoplus_{i \neq P} [y]_i \oplus y$ . Replace the last message from party  $P$  for reconstructing the output to  $[y]_P$ .

It is easy to see that  $\hat{z}$  is uniformly random in both hybrids since the share of the mask held by party  $P$  is uniformly random.  $[y]_P$  is identically distributed in the two hybrids given the perfect correctness of the protocol: in both worlds,  $[y]_P$  is a deterministic function of the output  $y$  and the messages send by parties other than  $P$ .

Therefore, **Hybrid<sub>3</sub>** and **Hybrid<sub>2</sub>** are identical. □

## 6.2 Security Proof of the Signature Scheme

In this section, we give a dedicated proof of security for the signature scheme constructed by making the KKW proof protocol non-interactive with the Fiat-Shamir transform. In doing so, our goals are both to give a complete proof (taking into account certain optimizations mentioned in the text), as well as to highlight the concrete-security bound we obtain. The theorem below proves EUF-CMA security, we believe it can be generalized to strong unforgeability.

In this section  $\kappa$  is a security parameter, and  $G$  is a hash function modeled as a random oracle (different from the permutation  $G$  used in Picnic-UR). We abstract our scheme by assuming that the key-generation algorithm **Gen** outputs a pair  $(C, w)$  with  $C(w) = 1$ , where we view  $C$  as the public key and  $w$  as the private key. We assume  $|C| \geq \kappa$  and  $w \in \{0, 1\}^\kappa$ . Our hardness

assumption is that, given  $C$  as output by **Gen**, it is hard to find  $w'$  for which  $C(w') = 1$ . More formally, we say that **Gen** is  $(t, \epsilon)$ -one way if for all adversaries  $\mathcal{A}$  running in time at most  $t$  we have

$$\Pr[(C, w) \leftarrow \text{Gen}; w' \leftarrow \mathcal{A}(C) : C(w') = 1] \leq \epsilon.$$

**Theorem 6.2.** *Suppose the PRG used is  $(t, \epsilon_{PRG})$ -secure, **Gen** is  $(t, \epsilon_{OW})$ -one-way, and  $\Pi$  is the MPC protocol described in Section 2.9.1. Model  $H_0, H_1, H_2$ , and  $G$  as random oracles where  $H_0, H_1, H_2$  have  $2\kappa$ -bit output length. Then any attacker carrying out an adaptive chosen-message attack on Picnic3 (Figure 4), running in time  $t$ , making  $q_s$  signing queries, and making  $q_0, q_1, q_2, q_G$  queries, respectively, to the random oracles, succeeds in outputting a valid forgery with probability at most*

$$\Pr[\text{Forge}] \leq O(q_s \cdot \tau \cdot \epsilon_{PRG}) + O\left(\frac{(q_0 + q_1 + q_2 + M n q_s)^2}{2^\kappa}\right) + \epsilon_{OW} + q_G \cdot \epsilon(M, n, \tau),$$

where

$$\epsilon(M, n, \tau) = \max_{M-\tau \leq k \leq M} \left\{ \frac{\binom{k}{M-\tau}}{\binom{M}{M-\tau} \cdot n^{k-M+\tau}} \right\}.$$

In the Picnic3 specification, the parameters  $(M, n, \tau)$  are chosen such that  $\epsilon(M, n, \tau) \leq 2^{-\kappa}$ .

*Proof.* Fix some attacker  $\mathcal{A}$ . Let  $q_s$  denote the number of signing queries made by  $\mathcal{A}$ ; let  $q_0, q_1, q_2$ , respectively, denote the number of queries to  $H_0, H_1, H_2$  made by  $\mathcal{A}$ , and let  $q_G$  denote the number of queries to  $G$  made by  $\mathcal{A}$ . To prove security we define a sequence of experiments involving  $\mathcal{A}$ , where the first corresponds to the experiment in which  $\mathcal{A}$  interacts with the real signature scheme. We let  $\Pr_i[\cdot]$  refer to the probability of an event in experiment  $i$ . We let  $t$  denote the running time of the entire experiment, i.e., including both  $\mathcal{A}$ 's running time and the time required to answer signing queries and to verify  $\mathcal{A}$ 's output.

**Experiment 1.** This corresponds to the interaction of  $\mathcal{A}$  with the real signature scheme. In more detail: first **Gen** is run to obtain  $(C, w)$ , and  $\mathcal{A}$  is given the public key  $C$ . In addition, we assume the random oracles  $H_0, H_1, H_2$ , and  $G$  are chosen uniformly from the appropriate spaces.  $\mathcal{A}$  may make signing queries, which will be answered as in Figure 4;  $\mathcal{A}$  may also query

### Picnic3 Signing

**Keys:** The public key is a circuit  $C$ ; the private key is a value  $w$  for which  $C(w) = 1$ . Values  $M, n, \tau$  are parameters of the protocol.

To sign message  $m$ , the signer does the following.

**Step 1** For each  $j \in [M]$ :

1. Choose uniform  $\text{seed}_j^* \in \{0, 1\}^\kappa$  and use it to generate values  $\text{seed}_{j,1}, \dots, \text{seed}_{j,n}$  with a PRG. Also compute  $\text{aux}_j \in \{0, 1\}^{|C|}$  as described in the text. For  $i = 1, \dots, n-1$ , let  $\text{state}_{j,i} := \text{seed}_{j,i}$ ; let  $\text{state}_{j,n} := \text{seed}_{j,n} \parallel \text{aux}_j$ .
2. For  $i \in [n]$ , compute  $\text{com}_{j,i} := H_0(\text{state}_{j,i})$ .
3. The signer runs the online phase of the  $n$ -party protocol  $\Pi$  (as described in the text) using  $\{\text{state}_{j,i}\}_i$ , beginning by computing the masked inputs  $\{\hat{z}_{j,\alpha}\}$  (based on  $w$  and the  $\{\lambda_{j,\alpha}\}$  defined by the preprocessing). Let  $\text{msgs}_{j,i}$  denote the messages broadcast by  $S_i$  in this protocol execution.
4. Let  $h_j := H_1(\text{com}_{j,1}, \dots, \text{com}_{j,n})$  and let  $h'_j := H_2(\{\hat{z}_{j,\alpha}\}, \text{msgs}_{j,1}, \dots, \text{msgs}_{j,n})$ .

**Step 2** Compute  $(\mathcal{C}, \mathcal{P}) := G(m, h_1, h'_1, \dots, h_M, h'_M)$ , where  $\mathcal{C} \subset [M]$  is a set of size  $\tau$ , and  $\mathcal{P}$  is a list  $\{p_j\}_{j \in \mathcal{C}}$  with  $p_j \in [n]$ . The signature includes  $(\mathcal{C}, \mathcal{P})$ .

**Step 3** For each  $j \in [M] \setminus \mathcal{C}$ , the signer includes  $\text{seed}_j^*, h'_j$  in the signature. Also, for each  $j \in \mathcal{C}$ , the signer includes  $\{\text{state}_{j,i}\}_{i \neq p_j}$ ,  $\text{com}_{j,p_j}$ ,  $\{\hat{z}_{j,\alpha}\}$ , and  $\text{msgs}_{j,p_j}$  in the signature.

Figure 4: The signing algorithm in the Picnic3 signature scheme.

### Picnic3 Verification

A signature  $(\mathcal{C}, \mathcal{P}, \{\text{seed}_j^*, h'_j\}_{j \notin \mathcal{C}}, \{\{\text{state}_{j,i}\}_{i \neq p_j}, \text{com}_{j,p_j}, \{\hat{z}_{j,\alpha}\}, \text{msgs}_{j,p_j}\}_{j \in \mathcal{C}})$  on a message  $m$  is verified as follows:

1. For every  $j \in \mathcal{C}$  and  $i \neq p_j$ , set  $\text{com}_{j,i} := H_0(\text{state}_{j,i})$ ; then compute the value  $h_j := H_1(\text{com}_{j,1}, \dots, \text{com}_{j,n})$ .
2. For  $j \notin \mathcal{C}$ , use  $\text{seed}_j^*$  to compute  $h_j$  as the signer would.
3. For each  $j \in \mathcal{C}$ , run an execution of  $\Pi$  among the parties  $\{S_i\}_{i \neq p_j}$  using  $\{\text{state}_{j,i}\}_{i \neq p_j}$ ,  $\{\hat{z}_\alpha\}$ , and  $\text{msgs}_{j,p_j}$ ; this yields  $\{\text{msgs}_i\}_{i \neq p_j}$  and an output bit  $b$ . Check that  $b \stackrel{?}{=} 1$ . Then compute  $h'_j := H_2(\{\hat{z}_{j,\alpha}\} \text{msgs}_{j,1}, \dots, \text{msgs}_{j,n})$ .
4. Check that  $(\mathcal{C}, \mathcal{P}) \stackrel{?}{=} G(m, h_1, h'_1, \dots, h_M, h'_M)$ .

Figure 5: The verification algorithm in the Picnic3 signature scheme.

any of the random oracles. Finally,  $\mathcal{A}$  outputs a message/signature pair; we let **Forge** denote the event that the message was not previously queried by  $\mathcal{A}$  to its signing oracle, and the signature is valid. We are interested in upper-bounding  $\Pr_1[\text{Forge}]$ .

**Experiment 2.** We abort the experiment if, during the course of the experiment, a collision in  $H_0$ ,  $H_1$ , or  $H_2$  is found. Suppose  $q = \max\{q_0, q_1, q_2\}$ , then the number of queries to any oracle throughout the experiment (by either the adversary or the signing algorithm) is at most  $(q + Mnq_s)$ . Thus,

$$|\Pr_1[\text{Forge}] - \Pr_2[\text{Forge}]| \leq \frac{3(q + Mnq_s)^2}{2^{2\kappa}}.$$

**Experiment 3.** Here we modify the way signing is done. Specifically, when signing a message  $m$  we begin by choosing  $(\mathcal{C}, \mathcal{P})$  uniformly. Steps 1 and 3 of the signing algorithm are computed as before, but in step 2 we simply set the output of  $G$  equal to  $(\mathcal{C}, \mathcal{P})$ . Formally, a signature on a message  $m$  is now computed as follows:

**Step 0** Choose uniform  $(\mathcal{C}, \mathcal{P})$ , where  $\mathcal{C} \subset [M]$  is a set of size  $\tau$ , and  $\mathcal{P} = \{p_j\}_{j \in \mathcal{C}}$  with  $p_j \in [n]$ .

**Step 1** For each  $j \in [M]$ :

1. Choose uniform  $\text{seed}_j^* \in \{0, 1\}^\kappa$  and use it to generate values  $\text{seed}_{j,1}, \dots, \text{seed}_{j,n}$  and  $\text{aux}_j \in \{0, 1\}^{|C|}$ . For  $i = 1, \dots, n-1$ , let  $\text{state}_{j,i} := \text{seed}_{j,i}$ ; let  $\text{state}_{j,n} := \text{seed}_{j,n} \parallel \text{aux}_j$ .
2. For  $i \in [n]$ , compute  $\text{com}_{j,i} := H_0(\text{state}_{j,i})$ .
3. Run the online phase of the  $n$ -party protocol  $\Pi$  using  $\{\text{state}_{j,i}\}_i$ , beginning by computing the masked inputs  $\{\hat{z}_{j,\alpha}\}$  (based on  $w$  and the  $\{\lambda_{j,\alpha}\}$  defined by the preprocessing). Let  $\text{msgs}_{j,i}$  denote the messages broadcast by  $S_i$  in this protocol execution.
4. Let  $h_j := H_1(\text{com}_{j,1}, \dots, \text{com}_{j,n})$  and  $h'_j := H_2(\{\hat{z}_{j,\alpha}\}, \text{msgs}_{j,1}, \dots, \text{msgs}_{j,n})$ .

**Step 2** Set  $G(m, h_1, h'_1, \dots, h_M, h'_M)$  equal to  $(\mathcal{C}, \mathcal{P})$ . (I.e., if  $\mathcal{A}$  subsequently makes the query  $G(m, h_1, h'_1, \dots, h_M, h'_M)$ , return  $(\mathcal{C}, \mathcal{P})$  as the output.) Include  $(\mathcal{C}, \mathcal{P})$  in the signature.

**Step 3** For each  $j \in [M] \setminus \mathcal{C}$ , the signer includes  $\text{seed}_j^*, h'_j$  in the signature. Also, for each  $j \in \mathcal{C}$ , the signer includes  $\{\text{state}_{j,i}\}_{i \neq p_j}, \text{com}_{j,p_j}, \{\hat{z}_{j,\alpha}\}$ , and  $\text{msgs}_{j,p_j}$  in the signature.

The only difference between this experiment and the previous one occurs if, in the course of answering a signing query, the query to  $G$  in step 2 was ever made before (by either the adversary or as part of answering some other signing query). Letting  $\text{InputColl}_G$  denote this event, we have

$$|\Pr_3[\text{Forge}] - \Pr_2[\text{Forge}]| \leq \Pr_3[\text{InputColl}_G].$$

**Experiment 4.** Here we again modify the way signing is done. Now, the signer chooses uniform  $\{\text{seed}_{j,i}\}_{i=1}^n$  for all  $j \in \mathcal{C}$ . That is, signatures are now computed as follows:

**Step 0** Choose uniform  $(\mathcal{C}, \mathcal{P})$ , where  $\mathcal{C} \subset [M]$  is a set of size  $\tau$ , and  $\mathcal{P} = \{p_j\}_{j \in \mathcal{C}}$  with  $p_j \in [n]$ .

**Step 1** For each  $j \in [M]$ :

1. If  $j \notin \mathcal{C}$ , choose uniform  $\text{seed}_j^* \in \{0, 1\}^\kappa$  and use it to generate values  $\text{seed}_{j,1}, \dots, \text{seed}_{j,n}$ . If  $j \in \mathcal{C}$ , choose uniform  $\text{seed}_{j,1}, \dots, \text{seed}_{j,n} \in \{0, 1\}^\kappa$ .

2. Compute  $\mathbf{aux}_j \in \{0, 1\}^{|C|}$  based on  $\{\mathbf{seed}_{j,i}\}_i$ . For  $i = 1, \dots, n-1$ , let  $\mathbf{state}_{j,i} := \mathbf{seed}_{j,i}$ ; let  $\mathbf{state}_{j,n} := \mathbf{seed}_{j,n} \parallel \mathbf{aux}_j$ .
3. For  $i \in [n]$ , compute  $\mathbf{com}_{j,i} := H_0(\mathbf{state}_{j,i})$ .
4. Run the online phase of the  $n$ -party protocol  $\Pi$  using  $\{\mathbf{state}_{j,i}\}_i$ , beginning by computing the masked inputs  $\{\hat{z}_{j,\alpha}\}$  (based on  $w$  and the  $\{\lambda_{j,\alpha}\}$  defined by the preprocessing). Let  $\mathbf{msgs}_{j,i}$  denote the messages broadcast by  $S_i$  in this protocol execution.
5. Let  $h_j := H_1(\mathbf{com}_{j,1}, \dots, \mathbf{com}_{j,n})$  and  $h'_j := H_2(\{\hat{z}_{j,\alpha}\}, \mathbf{msgs}_{j,1}, \dots, \mathbf{msgs}_{j,n})$ .

**Step 2** Set  $G(m, h_1, h'_1, \dots, h_M, h'_M)$  equal to  $(\mathcal{C}, \mathcal{P})$ . (I.e., if  $\mathcal{A}$  subsequently makes the query  $G(m, h_1, h'_1, \dots, h_M, h'_M)$ , return  $(\mathcal{C}, \mathcal{P})$  as the output.) Include  $(\mathcal{C}, \mathcal{P})$  in the signature.

**Step 3** For each  $j \notin \mathcal{C}$ , include  $\mathbf{seed}_j^*, h'_j$  in the signature. For each  $j \in \mathcal{C}$ , include  $\{\mathbf{state}_{j,i}\}_{i \neq p_j}$ ,  $\mathbf{com}_{j,p_j}$ ,  $\{\hat{z}_{j,\alpha}\}$ , and  $\mathbf{msgs}_{j,p_j}$  in the signature.

It is easy to see that if the pseudorandom generator is  $(t, \epsilon_{PRG})$ -secure, then

$$|\Pr_4[\text{Forge}] - \Pr_3[\text{Forge}]| \leq q_s \cdot \tau \cdot \epsilon_{PRG}$$

and

$$|\Pr_4[\text{InputColl}_G] - \Pr_3[\text{InputColl}_G]| \leq q_s \cdot \tau \cdot \epsilon_{PRG}.$$

We now bound  $\Pr_4[\text{InputColl}_G]$ . Fix some previous query  $(m, h_1, h'_1, \dots, h_M, h'_M)$  to  $G$ , and look at a query  $G(\hat{m}, \hat{h}_1, \hat{h}'_1, \dots, \hat{h}_M, \hat{h}'_M)$  made while responding to some signing query. (In the rest of this discussion, we will use  $\hat{\cdot}$  to represent values computed as part of answering that signing query.) For some fixed  $j \in \hat{\mathcal{C}}$ , it is not hard to see that the probability of the event  $\hat{h}_j = h_j$  is maximized if  $h_j$  was output by a previous query  $H_1(\mathbf{com}_1, \dots, \mathbf{com}_n)$ , and each  $\mathbf{com}_i$  was output by a previous query  $H_0(\mathbf{state}_i)$ . (In all cases, the relevant prior query must be unique since the experiment is aborted if there is a collision in  $H_0$  or  $H_1$ .) In that case, the probability that  $\hat{h}_j = h_j$  is at most

$$(2^{-\kappa} + 2^{-2\kappa})^n + 2^{-2\kappa} \leq 2 \cdot 2^{-2\kappa}$$

(assuming  $n \geq 3$ ), and thus the probability that  $\hat{h}_j = h_j$  for all  $j \in \hat{\mathcal{C}}$  is at most  $2^{-\tau \cdot (2\kappa-1)}$ . Taking a union bound over all signing queries and all queries made to  $G$  (including those made during the course of answering signing queries), we conclude that

$$\Pr_4[\text{InputColl}_G] \leq q_s \cdot (q_s + q_G) \cdot 2^{-\tau \cdot (2\kappa-1)}.$$

**Experiment 5.** Here we again modify the way signing is done. Now:

- For each  $j \in \mathcal{C}$ , choose uniform  $\text{com}_{j,p_j}$  (i.e., without making the corresponding query to  $H_0$ ).
- For each  $j \notin \mathcal{C}$ , choose uniform  $h'_j$  (i.e., without making the corresponding query to  $H_2$ ).

So, signatures are now computed as follows:

**Step 0** Choose uniform  $(\mathcal{C}, \mathcal{P})$ , where  $\mathcal{C} \subset [M]$  is a set of size  $\tau$ , and  $\mathcal{P} = \{p_j\}_{j \in \mathcal{C}}$  with  $p_j \in [n]$ .

**Step 1** For each  $j \in [M]$ :

1. If  $j \notin \mathcal{C}$ , choose uniform  $\text{seed}_j^* \in \{0, 1\}^\kappa$  and use it to generate values  $\text{seed}_{j,1}, \dots, \text{seed}_{j,n}$ . If  $j \in \mathcal{C}$ , choose uniform  $\text{seed}_{j,1}, \dots, \text{seed}_{j,n} \in \{0, 1\}^\kappa$ .
2. Compute  $\text{aux}_j \in \{0, 1\}^{|\mathcal{C}|}$  based on  $\{\text{seed}_{j,i}\}_i$ . For  $i = 1, \dots, n-1$ , let  $\text{state}_{j,i} := \text{seed}_{j,i}$ ; let  $\text{state}_{j,n} := \text{seed}_{j,n} \parallel \text{aux}_j$ .
3. For  $j \in \mathcal{C}$ , choose uniform  $\text{com}_{j,p_j} \in \{0, 1\}^{2\kappa}$ . For all other  $j, i$ , set  $\text{com}_{j,i} := H_0(\text{state}_{j,i})$ .
4. Run the online phase of the  $n$ -party protocol  $\Pi$  using  $\{\text{state}_{j,i}\}_i$ , beginning by computing the masked inputs  $\{\hat{z}_{j,\alpha}\}$  (based on  $w$  and the  $\{\lambda_{j,\alpha}\}$  defined by the preprocessing). Let  $\text{msgs}_{j,i}$  denote the messages broadcast by  $S_i$  in this protocol execution.
5. Let  $h_j := H_1(\text{com}_{j,1}, \dots, \text{com}_{j,n})$ . If  $j \in \mathcal{C}$ , set  $h'_j := H_2(\{\hat{z}_{j,\alpha}\}, \text{msgs}_{j,1}, \dots, \text{msgs}_{j,n})$ ; otherwise, choose uniform  $h'_j \in \{0, 1\}^{2\kappa}$ .

**Step 2** Set  $G(m, h_1, h'_1, \dots, h_M, h'_M)$  equal to  $(\mathcal{C}, \mathcal{P})$ . (I.e., if  $\mathcal{A}$  subsequently makes the query  $G(m, h_1, h'_1, \dots, h_M, h'_M)$ , return  $(\mathcal{C}, \mathcal{P})$  as the output.) Include  $(\mathcal{C}, \mathcal{P})$  in the signature.

**Step 3** For each  $j \notin \mathcal{C}$ , include  $\text{seed}_j^*, h'_j$  in the signature. For each  $j \in \mathcal{C}$ , include  $\{\text{state}_{j,i}\}_{i \neq p_j}, \text{com}_{j,p_j}, \{\hat{z}_{j,\alpha}\}$ , and  $\text{msgs}_{j,p_j}$  in the signature.

The only difference between this experiment and the previous one occurs if, during the course of answering a signing query,  $\text{state}_{j,p_j}$  (for some  $j \in \mathcal{C}$ ) is queried to  $H_0$  at some other point in the experiment, or  $(\{\hat{z}_{j,\alpha}\}, \text{msgs}_{j,1}, \dots,$

$\text{msgs}_{j,n}$ ) (for some  $j \notin \mathcal{C}$ ) is ever queried to  $H_2$  at some other point in the experiment. Denoting this event by  $\text{InputColl}_H$ , we thus have

$$|\Pr_5[\text{Forge}] - \Pr_4[\text{Forge}]| \leq \Pr_5[\text{InputColl}_H].$$

**Experiment 6.** We again modify the signing algorithm. Now, for  $j \in \mathcal{C}$  the signer uses the simulator for  $\Pi$  (namely,  $\text{Sim}_\Pi$ ) to generate the views of the parties  $\{S_i\}_{i \neq p_j}$  in an execution of  $\Pi$  when evaluating  $C$  with output 1. This results in values  $\{\text{state}_{j,i}\}_{i \neq p_j}$ , masked input-wire values  $\{\hat{z}_{j,\alpha}\}$ , and  $\text{msgs}_{j,p_j}$ . From the respective views,  $\{\text{msgs}_{j,i}\}_{i \neq p_j}$  can be computed, and  $h_j, h'_j$  can be computed as well. Thus, signatures are now computed as follows:

**Step 0** Choose uniform  $(\mathcal{C}, \mathcal{P})$ , where  $\mathcal{C} \subset [M]$  is a set of size  $\tau$ , and  $\mathcal{P} = \{p_j\}_{j \in \mathcal{C}}$  with  $p_j \in [n]$ .

**Step 1** For  $j \notin \mathcal{C}$ :

1. Choose uniform  $\text{seed}_j^* \in \{0, 1\}^\kappa$  and use it to generate values  $\text{seed}_{j,1}, \dots, \text{seed}_{j,n}$ . Compute  $\text{aux}_j \in \{0, 1\}^{|C|}$  based on  $\{\text{seed}_{j,i}\}_i$ . For  $i = 1, \dots, n-1$ , let  $\text{state}_{j,i} := \text{seed}_{j,i}$ ; let  $\text{state}_{j,n} := \text{seed}_{j,n} \parallel \text{aux}_j$ .
2. For all  $i$ , set  $\text{com}_{j,i} := H_0(\text{state}_{j,i})$ .
3. Let  $h_j := H_1(\text{com}_{j,1}, \dots, \text{com}_{j,n})$ . Choose uniform  $h'_j \in \{0, 1\}^{2\kappa}$ .

For each  $j \in \mathcal{C}$ :

1. Compute  $(\{\text{state}_{j,i}\}_{i \neq p_j}, \{\hat{z}_{j,\alpha}\}, \text{msgs}_{j,p_j}) \leftarrow \text{Sim}_\Pi(p_j)$ . Compute  $\{\text{msgs}_{j,i}\}_{i \neq p_j}$  based on this information.
2. Choose uniform  $\text{com}_{j,p_j} \in \{0, 1\}^{2\kappa}$ . For all other  $i$ , set  $\text{com}_{j,i} := H_0(\text{state}_{j,i})$ .
3. Let  $h_j := H_1(\text{com}_{j,1}, \dots, \text{com}_{j,n})$  and  $h'_j := H_2(\{\hat{z}_{j,\alpha}\}, \text{msgs}_{j,1}, \dots, \text{msgs}_{j,n})$ .

**Step 2** Set  $G(m, h_1, h'_1, \dots, h_M, h'_M)$  equal to  $(\mathcal{C}, \mathcal{P})$ . (I.e., if  $\mathcal{A}$  subsequently makes the query  $G(m, h_1, h'_1, \dots, h_M, h'_M)$ , return  $(\mathcal{C}, \mathcal{P})$  as the output.) Include  $(\mathcal{C}, \mathcal{P})$  in the signature.

**Step 3** For each  $j \notin \mathcal{C}$ , the signer includes  $\text{seed}_j^*, h'_j$  in the signature. Also, for each  $j \in \mathcal{C}$ , the signer includes  $\{\text{state}_{j,i}\}_{i \neq p_j}$ ,  $\text{com}_{j,p_j}$ ,  $\{\hat{z}_{j,\alpha}\}$ , and  $\text{msgs}_{j,p_j}$  in the signature.



Observe that  $w$  is no longer used for generating signatures. Recall, the adversary in the underlying MPC protocol  $\Pi$  has distinguishing advantage  $\epsilon_{PRG}$  (see Lemma 6.1). It is immediate that

$$|\Pr_6[\text{Forge}] - \Pr_5[\text{Forge}]| \leq \tau \cdot q_s \cdot \epsilon_{PRG}$$

and

$$|\Pr_6[\text{InputColl}_H] - \Pr_5[\text{InputColl}_H]| \leq \tau \cdot q_s \cdot \epsilon_{PRG}.$$

We now bound  $\Pr_6[\text{InputColl}_H]$ . For any particular signing query and any  $j \in \mathcal{C}$ , the value  $\text{state}_{j,p_j}$  has min-entropy at least  $\kappa$  and is not used anywhere else in the experiment. Similarly, for any  $j \notin \mathcal{C}$ , the value  $\{\hat{z}_{j,\alpha}\}$  has min-entropy at least  $\kappa$ , since the input is  $\kappa$ -bit and they are all uniform according to the simulator defined in the next section. and is not used anywhere else in the experiment. Thus,

$$\Pr_6[\text{InputColl}_H] \leq M \cdot q_s \cdot (Mq_s + q_0 + q_2) \cdot 2^{-\kappa}.$$

**Experiment 7.** We first define some notation. At any point during the experiment, we classify a pair  $(h, h')$  in one of the following ways:

1. If  $h$  was output by a previous query  $H_1(\text{com}_1, \dots, \text{com}_n)$ , and each  $\text{com}_i$  was output by a previous query  $H_0(\text{state}_i)$  where the  $\{\text{state}_i\}$  form a valid preprocessing, then say  $(h, h')$  *defines correct preprocessing*.
2. If  $h$  was output by a previous query  $H_1(\text{com}_1, \dots, \text{com}_n)$ , and each  $\text{com}_i$  was output by a previous query  $H_0(\text{state}_i)$ , and  $h'$  was output by a previous query  $H_2(\{\hat{z}_\alpha\}, \text{msgs}_1, \dots, \text{msgs}_n)$  where  $\{\text{state}_i\}, \{\hat{z}_\alpha\}, \{\text{msgs}_i\}$  are consistent with an online execution of  $\Pi$  among all parties with output 1 (but the  $\{\text{state}_i\}$  may not form a valid preprocessing), then say  $(h, h')$  *defines correct execution*.
3. In any other case, say  $(h, h')$  is *bad*.

(Note that in all cases the relevant prior query, if it exists, must be unique since the experiment is aborted if there is ever a collision in  $H_0, H_1$ , or  $H_2$ .)

In Experiment 7, for each query  $G(m, h_1, h'_1, \dots, h_M, h'_M)$  made by the adversary (where  $m$  was not previously queried to the signing oracle), check if there exists an index  $j$  for which  $(h_j, h'_j)$  defines correct preprocessing and correct execution. We let **Invert** be the event that this occurs for some

query to  $G$ . Note that if that event occurs, the  $\{\text{state}_i\}, \{\hat{z}_\alpha\}$  (which can be determined from the oracle queries of the adversary) allow computation of  $w'$  for which  $C(w') = 1$ . Thus,  $\Pr_7[\text{Invert}] \leq \epsilon_{OW}$ .

We claim that

$$\Pr_7[\text{Forge} \wedge \overline{\text{Invert}}] \leq q_G \cdot \epsilon(M, n, \tau).$$

To see this, assume  $\text{Invert}$  does not occur. For any query  $G(m, h_1, h'_1, \dots, h_M, h'_M)$  made during the experiment (where  $m$  was not previously queried to the signing oracle), let  $\text{Pre}$  denote the set of indices for which  $(h_j, h'_j)$  defines correct preprocessing (but not correct execution), and let  $k = |\text{Pre}|$ . Let  $(\mathcal{C}, \mathcal{P})$  be the (random) answer from this query to  $G$ . The attacker can only possibly generate a forgery (using this  $G$ -query) if (1)  $[M] \setminus \mathcal{C} \subseteq \text{Pre}$ , and (2) for all  $j \in \text{Pre} \cap \mathcal{C}$ , the value  $p_j$  is chosen to be the unique party such that the views of the remaining parties  $\{S_i\}_{i \neq p_j}$  are consistent. Since  $|M \setminus \mathcal{C}| = M - \tau$ , the number of ways the first event can occur is  $\binom{k}{M-\tau}$ ; given this, there are  $k - (M - \tau)$  elements remaining  $\text{Pre} \cap \mathcal{C}$ . Thus, the overall probability with which the attacker can generate a forgery using this  $G$ -query is

$$\begin{aligned} \epsilon(M, n, \tau, k) &= \frac{\binom{k}{M-\tau} \cdot n^{M-k}}{\binom{M}{M-\tau} \cdot n^\tau} \\ &= \frac{\binom{k}{M-\tau}}{\binom{M}{M-\tau} \cdot n^{k-M+\tau}} \\ &\leq \epsilon(M, n, \tau) = \max_k \{\epsilon(M, n, \tau, k)\}. \end{aligned}$$

The final bound is obtained by taking a union bound over all queries to  $G$ .  $\square$

### 6.3 Tree-Based Optimizations

In this section we discuss standard constructions of seed tree and Merkle tree and their use in **Picnic3**. These are well-known cryptographic objects, which are used in a standard way. Nevertheless, in this section we briefly sketch security arguments corresponding to their use. We first consider some properties of the seed tree construction, then discuss the use of seed and Merkle trees in **Picnic3**.

### 6.3.1 Seed Tree

The beginning of this section is taken (with minor modifications) from the Picnic specification [Tea19a], Section 7.3.1, that describes how the tree is constructed, and how seeds are efficiently revealed.

When signing, the signer must generate a set of seeds, then reveal a subset of these based on the challenge. The seeds are then used by the verifier to check that the MPC protocol was setup or simulated correctly. By deriving seeds deterministically in a binary tree, then using the leaf seeds in the protocol, the signer can reveal large subsets of the seeds efficiently by revealing intermediate nodes in the tree. In Picnic3, the signer must reveal  $M - \tau$  of the  $M$  initial seeds and one of the  $n$  seeds in each of the  $\tau$  MPC instances that are checked by the verifier.

The tree is initialized with random data at the root node. For each non-leaf node in the tree, having seed `parent_seed`, compute the  $2\kappa$ -bit digest

$$d := H(\text{parent\_seed} \parallel \text{salt} \parallel j \parallel i)$$

where  $j$  is the MPC instance number, and  $i$  is the index of the parent node. The function  $H$  is a hash function with  $2\kappa$ -bit outputs. Then set the left child of the node to the leftmost  $\kappa$  bits of  $d$ , and the right child to the rightmost  $\kappa$  bits of  $d$ . The *salt* is a random, per-signature value that is included to prevent multi-target attacks (along with the counter). The values *salt*,  $j$  and  $i$  are always public, but only some of the `parent_seed` values are revealed.

**Collision-resistance** Non-malleability of Picnic3 signatures requires that the seed tree be collision-resistant. This means it must be hard to find distinct seeds that expand to the same set of leaf seeds. In practice 2nd preimage resistance should be sufficient, since an attacker must find a second seed that expands to the same set of seeds appearing in a valid signature. However, we use collision resistance because it allows us to prove that the proof protocol has computationally unique responses (also called quasi-unique responses) in Lemma 6.7.

**Theorem 6.3.** *The seed tree construction is collision-resistant if the hash function  $H$  is collision-resistant.*

*Proof.* The algorithm to reveal seeds ensures that any node without a sibling is revealed directly (as opposed to being re-derived from another seed). Therefore, both halves of the output bits of  $H$  are used, so if distinct seeds  $s$

and  $s'$  derive the same set of seeds, we have a collision,  $H(s) = H(s')$  where  $s$  is a parent of two leaf seeds.  $\square$

**Hiding and Pseudorandomness** The derived seeds must be pseudorandom for security of the protocol. Further, the unrevealed seeds must remain hidden, or else the signature will leak information about the private key. Therefore, given the revealed seeds, it must be difficult to distinguish the unrevealed seeds from random values. For this we will rely on the function  $H$  being a PRF keyed by the parent seed, so that given half the output bits of  $H(\text{parent\_seed} \parallel \text{salt} \parallel j \parallel i)$  the sibling value the other half of the output bits are indistinguishable from a random value.

**Theorem 6.4.** *Given the opening information for a set of revealed seeds, the unrevealed seeds remain indistinguishable from random values, assuming  $H$  is a secure PRF when keyed with an unrevealed seed.*

*Proof.* Consider a subtree with root  $p$ , having left child  $r = \text{leftHalf}(H(p \parallel \dots))$ , and right child  $s = \text{rightHalf}(H(p \parallel \dots))$ . We must show that if  $r$  is revealed,  $s$  remains indistinguishable from random (the same argument holds when  $r$  and  $s$  are reversed). Since  $H$  is assumed to be a PRF with key  $p$ , and  $p$  is unrevealed, the output  $H(p \parallel \dots)$  is indistinguishable from a random value. In a random value the two halves of the output are independent, so having one half gives no information about other, so the hidden half remains indistinguishable from random.

The value  $p$  is unrevealed, if it has a sibling that is revealed, the same argument is applied recursively up the tree (with height bounded by  $O(\log(M))$ , polynomial in  $\kappa$ ). The root is only revealed when there are no unrevealed seeds.  $\square$

In practice  $H$  is SHAKE, and inputs always have a fixed length.

### 6.3.2 Use in Picnic3

Here we review where the the seed tree and Merkle tree optimizations are used in Picnic3 and argue that the security properties above are sufficient to maintain security.

The seed tree construction is used in two places. First, in Step 1.1, instead of choosing  $\{\text{seed}_j^*\}_{j=1}^M$  at random, the signer derives them using a seed tree. Then in Step 3, the signer outputs seeds from the tree that allows the verifier to recompute the  $M - \tau$  revealed seeds.

- The privacy property of the seed tree construction (Theorem 6.4), ensures that the unrevealed seeds are indistinguishable from random.
- For malleability, if the seeds output by the signer are modified, collision resistance ensures that at least one of the  $M - \tau$  reconstructed seeds is different. Then verification will fail, since  $\mathbf{seed}_j^*$  is used as the root of a seed tree, which is collision-resistant, meaning the values  $\mathbf{seed}_{j,i}$  will be different.

The second use is in Step 1.1, in each instance  $j$ , the  $n$  seeds  $\{\mathbf{seed}_{j,i}\}$  are derived with the seed tree construction with root  $\mathbf{seed}_j^*$ . Then in Step 3, instead of revealing  $\{\mathbf{seed}_{j,i}\}_{i \neq p_j}$ , the signer reveals seeds from the tree that allow the verifier to recompute this set.

- The privacy property of the seed tree construction ensures that  $\mathbf{seed}_{j,p_j}$  is indistinguishable from random.
- For malleability, if the revealed seeds are modified such that the recomputed seeds  $\{\mathbf{seed}_{j,i}\}_{i \neq p_j}$  are different, then verification will fail when  $\{\mathbf{com}_{j,i}\}_{i \neq p_j}$  are recomputed (or there is a collision in  $H_0, H_1$  or  $G$ ). Collision resistance of the seed tree construction ensures that changing the seeds output by the signer will cause the seeds recomputed by the verifier to be different.

The Merkle tree is used as a commitment to a set of commitments, the values  $\{h'_j\}_{j=1}^M$ . Using a Merkle tree does not affect the hiding property of the commitments  $h'_j$  (which are public when a Merkle tree is not used), so we only need to consider whether the Merkle tree is a binding commitment. It can be shown that providing two openings for a left with respect to the same root gives a collision in  $H$  (the function used to form the Merkle tree).

## 6.4 QROM Security

In this section we discuss the QROM security of Picnic signatures instantiated with the KKW proof protocol.

Picnic3 is a Fiat-Shamir type signature scheme, and the recent work of Don, Fehr, Majenz and Schaffner [DFMS19a] proves that a large class of FS signature schemes are secure in the QROM. In [DFMS19a, Corollary 26], security of Picnic instantiated with the ZKBoo proof system is shown, and it can be checked that this also holds when ZKB++ is used, establishing

the QROM security of the Picnic-L1-FS, Picnic-L3-FS and Picnic-L5-FS parameter sets.

However, the published version of [DFMS19a] did not apply to instances of Picnic using the KKW protocol (i.e., Picnic3), since the KKW protocol did not meet the required definition of  $t$ -soundness. Informally, a  $\Sigma$ -protocol is said to be  $t$ -sound if we can extract a witness from any set of  $t$  accepting transcripts that have the same commitment, and different challenges and responses. Then the approach to show that a  $t$ -sound  $\Sigma$ -protocol is a proof of knowledge is to define an extractor that interacts with a malicious prover (in the random oracle model). The extractor runs the prover to obtain an accepting transcript, then rewinds her  $t - 1$  times to the point when the challenge is output by the random oracle, and outputs a different challenge. The extractor then obtains an additional  $t - 1$  accepting transcripts, then extracts the witness from these  $t$  transcripts.

While it is possible to extract a witness from three accepting KKW transcripts, the extractor does not work for **any** three. As shown in [KKW18, Theorem 2.1], the three transcripts must have the form  $(c, p)$ ,  $(c', *)$ ,  $(c, p')$ , where  $*$  denotes any value,  $c \neq c'$  and  $p \neq p'$ . In particular this means that there are many triples of challenges for which extraction will fail. For example, if  $c$  is fixed for the three challenges and  $p$  differs, extraction may fail because the preprocessing step may be incorrect. Similarly, if  $p$  is fixed and  $c$  varies, the extractor does not have the state of all  $n$  parties, and cannot recover the witness since the MPC protocol is secure.

The rewinding technique of [DFMS19a] chooses a new challenge at random. Therefore, the extractor for KKW is not guaranteed to succeed, since it may be that the malicious prover only succeeds when the three challenges do not satisfy the required property. In an updated full version, Don et al. [DFMS19b] generalize their result to allow the extractor to re-use the first part of a previous challenge, after first using a random challenge. In particular, this allows three random challenges of the form  $(C, P)$ ,  $(C, P')$  and  $(C'', P'')$  for random  $C, P, P', C''$ , and  $P''$ . Thus, with overwhelming probability, the transcripts will contain a challenge of the correct form, and extraction will succeed.

One caveat we note is that this generalization comes with a cost in tightness of the reduction. The reduction for the ZKB++ parameter sets loses a factor of  $q^2$ , and for KKW the loss is a factor  $q^6$ , where  $q$  is the number of hash queries. As the results are non-tight, and depend on the asymptotic analysis of [DFMS19a], we make no claims about the concrete security of

Picnic in the QROM.

We must also assume that the hash functions  $H$ ,  $H_0$ ,  $H_1$ , and  $H_2$  used for commitments and deriving seeds in KKW are *collapsing*, a quantum generalization of collision resistance. For a definition of collapsing, see [DFMS19b, Definition 23]. The hash function  $G$  (used to compute the challenge) is modeled as a quantum random oracle.

**Theorem 6.5.** *Picnic3 is strongly unforgeable under chosen message attacks in the QROM when  $H, H_0, H_1$  and  $H_2$  are instantiated with collapsing hash functions.*

*Proof.* The proof follows the proof of [DFMS19a, Corollary 26] for Picnic-FS. To apply the main theorems of [DFMS19a], We require that the KKW  $\Sigma$ -protocol:

1. Be non-abort honest-verifier zero-knowledge (naHVZK, or simply HVZK because the KKW protocol does not abort, we do not have to quantify the non-abort probability).
  - (a) This was first done in [KKW18, Theorem 2.2], but considers the protocol when commitments are randomized.
  - (b) In [AOTZ20, Lemma 18], a similar proof is given, but accounts for the optimization of using non-randomized commitments.
  - (c) The proofs assume only that the hash functions are modeled as random oracles.
2. Have min-entropy (denoted  $\alpha$ ) that is polynomial in the security parameter, where min-entropy of a  $\Sigma$ -protocol is as defined in [KLS18, Definition 2.6]. This was done in [AOTZ20, Lemma 17].
3. Has quantum computationally unique responses (CUR).
  - (a) We first show that it has classical CUR, assuming the hash functions are collision resistant, in Lemma 6.7.
  - (b) It then follows that KKW has quantum CUR under the further assumption that the hash functions are collapse-binding.

Now to prove Picnic3 is sEUF-CMA secure, we can use [DFMS19b, Theorem 25], which says that a  $t$ -sound protocol that has quantum CUR is a

computational proof of knowledge (PoK) in the QROM. Here we use the generalized version that uses [DFMS19b, Lemma 30], allowing the extractor to re-use the first part of the challenge when rewinding. Then we apply [DFMS19b, Theorem 22], which says that a PoK that is also zero-knowledge with  $\alpha$  bits of min-entropy and CUR is strongly unforgeable under chosen-message attacks.  $\square$

## 6.5 Computationally Unique Responses

A  $\Sigma$ -protocol has computationally unique responses, if it is hard to find two valid responses for a fixed commitment and challenge. In this section we show that the KKW protocol as used in Picnic3, has CUR (Lemma 6.7).

**Definition 6.6** (Computationally Unique Responses). For a  $\Sigma$ -protocol  $\Pi$  with transcripts  $(W, c, Z)$  and verification function  $V$ , if the following probability is negligible in  $\kappa$  for all polynomial-time adversaries  $A$ ,

$$\Pr[V(pk, W, c, Z) = V(pk, W, c, Z') = 1 | (pk, W, c, Z, Z') \leftarrow A(\kappa)]$$

then we say that  $\Pi$  has *computationally unique responses* (CUR).

Our definition of CUR is based on Fischlin’s [Fis05], which differs from the definition of [KLS18] where  $A$  is given the public key, rather than being allowed to choose it. In the context of signature schemes  $A$  is typically trying to find multiple responses with respect to a given public key, but since CUR holds for Picnic3 even when  $A$  is allowed to choose the key pair we use the stronger definition. The CUR property (as defined here) is also called *quasi-unique responses* in [FKMV12b].

Our proof of CUR applies to Picnic3 including the two optimizations: seeds are derived using the tree construction, and commitments to the views are done with a Merkle tree, as described in the Picnic specification. We use  $H$  to denote the hash function used to instantiate the Merkle tree and the seed tree. Hashes  $H_0, H_1, H_2$  and  $G$  are as defined in Figure 4.

**Lemma 6.7.** *If the hash functions  $H, H_0, H_1, H_2$  and  $G$  used to implement Picnic3 are collision-resistant, then Picnic3 has computationally unique responses.*

*Proof.* Using the notation of Definition 6.6, suppose that  $A$  outputs  $(W, c, Z, Z')$ . We will ignore the public key  $pk$  and some of the other checks that the verification function might make, however, these only add constraints to  $A$ ’s



output, making it strictly more difficult. We show that it is difficult for  $A$  to find distinct  $Z, Z'$  with respect to any  $W, c$ , that pass only a subset of the verification function (namely, the checks that commitment openings in  $Z$  correspond to the commitments in  $W$ ).

Recall that in Picnic3:

- $W$  is  $h_1, \dots, h_M$ , and  $\text{MerkleTreeRoot}(h'_1, \dots, h'_M)$
- $c$  is  $(\mathcal{C}, \mathcal{P})$ , and
- $Z$  is
  - for  $j \in \mathcal{C}$ :  $\{\text{state}_{j,i}\}_{i \neq p_j}$ ,  $\text{com}_{j,p_j}$ ,  $\{\hat{z}_{j,\alpha}\}$  and  $\text{msgs}_{j,p_j}$
  - for  $j \notin \mathcal{C}$ :  $\text{seed}_j^*$ ,  $h'_j$

We will argue that any difference in the values in  $Z$  and  $Z'$  gives a collision in one of the hash functions.

If  $\text{state}_i$  differs, then  $\text{com}_i$  will differ unless we have a collision in  $H_0$ . If  $\text{com}_i$  differs, then  $h_j$  will differ unless we have a collision in  $H_1$ . Finally if  $h_j$  differs,  $c$  will differ unless we have a collision in  $G$ . But  $A$  has output a single  $c$  value, so either there is a collision, or the  $\text{state}_i$  values are all equal.

If  $\text{com}_{p_j}$  differ, then  $h_j$  and  $c$  will differ, or we have a collision in  $H_0, H_1$  or  $G$ .

If any of the  $\{\hat{z}_{j,\alpha}\}$  differ, then  $h'_j$  will differ or we have a collision in  $H_2$ . The root of the Merkle tree will differ (or there is a collision in  $H$ ), causing  $c$  to differ, unless we have a collision in  $G$ .

The values  $\{\text{msgs}_{j,p_j}\}$  are hashed with  $\{\hat{z}_{j,\alpha}\}$  to form  $h'_j$ , so the same reasoning ensures they will be distinct.

The values  $\text{seed}_j^*$  are root seeds used to derive  $\{\text{seed}_{i,j}\}$  for the MPC instance  $j$ , using the seed-tree construction. The root seeds do not get hashed as part of computing  $c$ . If  $Z$  and  $Z'$  contain two different root seeds  $s$  and  $s'$  that expand to different seeds, then one or more of the  $\text{state}$  values will change, causing  $c$  to be different, or a collision in  $H_0, H_1, H_2$  or  $G$  as argued above. If  $s$  and  $s'$  expand to the same seeds, then there would be a collision in the seed tree construction. As shown in Theorem 6.3, this would imply a collision in  $H$ .

Finally, if any of the  $h'_j$  values differ but  $c$  is the same, we have a collision in  $H_2$ , as argued for  $\{\hat{z}_{j,\alpha}\}$ .  $\square$

## 7 Analysis with Respect to Known Attacks

In this section we analyze the Picnic signature scheme with respect to known attacks. First, we observe that in case we deal with ideal primitives, Corollary 5.4 already gives us a provable bound for EUF-CMA security. Since those primitives are, however, instantiated with concrete building blocks, we consider concrete attacks on those building blocks. In our scheme, we use the classical approach to turn  $\Sigma$ -protocols into signature schemes in the random oracle model. Based on the fact that, since the introduction of the random oracle model [BR93], no attack which arises from the assumption that a hash function behaves as a random oracle (except for some artificial counterexamples such as [CGH98]) was found [KM15], we claim that the best attacks against our scheme are attacks which also invalidate the claims made for the underlying symmetric primitives.

All cryptographic primitives, except the one-way function LOWMC, rely on the SHA3 function SHAKE [NIS15], a well established and standardized primitive, and we use it in a standard way. For those primitives, we have already gained substantial confidence regarding security due to extensive cryptanalysis efforts within the community. We therefore do not see these building blocks as a central attack surface and assume that the bounds given in [NIS15, Table 4] hold. We note that improvements in attacks against those primitives may also lead to improvements in the attacks against our signature scheme.

### 7.1 Usage and Security Margin of LowMC

Here we focus on attacks on the one-way function  $f$ . Essentially, the function  $f$  could be any one-way function, but since we found block ciphers—and, in particular the LOWMC family of block ciphers [ARS<sup>+</sup>15a, ARS<sup>+</sup>16]—gave the most efficient signatures, we decided to use them in our signature scheme. In particular, we assume that using LOWMC as described below yields a suitable family of one-way functions  $\{f_u\}_{u \in K_\kappa}$ . We use this function to establish a suitable relation between secret and public keys. In particular, let

$$f_u(x) := E(x, u),$$

and let  $E$  denote LOWMC encryption with respect to a single block  $u$  under key  $x$ . The keys in our signature scheme are generated as follows. First, one

chooses a LOWMC encryption key  $x$ , as well as a single block  $u$  uniformly at random. Then, the the public verification key  $\mathbf{pk}$ , as well as the secret signing key  $\mathbf{sk}$  are defined as follows

$$\mathbf{pk} := (y, u) = (f_u(x), u), \mathbf{sk} := (\mathbf{pk}, x).$$

**LowMC parameters with a partial S-box layer** The choice of the number of rounds within LOWMC comes with a significant security margin. For security level L1 with the specified 20 rounds, the best attack known is on 12 rounds. For security level L3 with the specified 30 rounds, the best attack known is on 19 rounds. For security level L5 with the specified 38 rounds, the best attack known is on 26 rounds. And even those attacks require an attacker to see two plaintext-ciphertext pairs for the same key, whereas within our signature scheme an attacker only ever sees a single input-output pair for every key.

**LowMC parameters with a full S-box layer** In addition to the existing cryptanalysis, we consider the security of these new LowMC instances in the setting of Picnic. Here, an attacker only ever gets to see a single plaintext-ciphertext pair for a given secret key. This restriction significantly limits the possible attacks on block ciphers (e.g., differential, linear and interpolation attacks all require more than one plaintext/ciphertext pair), leaving algebraic attacks that rely on solving a system of equations in the key variables (e.g., Gröbner basis attacks [Fau02] or attacks using SAT solvers [BCJ07]). However, using both of these attacks, we can attack a maximum of two rounds of LowMC with less than brute-force complexity.

We also used a general tool for finding meet-in-the-middle attacks [DF16] to evaluate full S-box layer instances of LowMC. The tool reported attacks on an instance with 2 rounds, but could not find any attacks for instances with 3 or more rounds.

A general approach to attack ciphers with low multiplicative complexity is given by Zajac [Zaj17]. There, the cipher is transformed into a multiple right-hand sides equation system and subsequently transformed into a syndrome decoding problem. Zajac argues that the complexity of decoding of random linear codes can be used to calculate the complexity of such an attack. For a three-round version of LowMC with a 129 bit state, this would result in an attack with time complexity of  $2^{118}$  and high memory complexity of  $2^{89}$ . However, for instances with 4 or more rounds, the time (and memory) complexity is expected to exceed brute-force complexity.

**LowMC Cryptanalysis Challenge** Some members of the Picnic and LowMC design teams have initiated the LowMC Cryptanalysis Challenge [Gra20], to encourage security analysis of LowMC, with a specific focus on the parameter sets of interest to Picnic. The challenge was started in May 2020 and is expected to run for about two years, with prizes totalling 100K USD. In [GKRS20], Grassi et al. survey existing attacks against LowMC in the restricted attack scenario of interest for Picnic, and experimentally investigate the costs and applicability of known attacks.

The challenge has already yielded interesting cryptanalysis results, announced at the rump session of CRYPTO 2020. The winner of round 1, for the best attack on -full instances (see the paper by Banik et al. [BBDV20]) confirms the analysis of [GKRS20], that the best attack applies to only two of four rounds.

A second paper from 2020 is by Liu, Isobe and Meier [LIM20], which describes new attacks on certain instances of LowMC, different from those used in Picnic. First are key recovery attacks when the block size is much larger than the key size. For Picnic, we always have equal block and key sizes. Second, they improve key recovery attacks on the -full instances when the attacker can choose two plaintexts, breaking four rounds. However, the attacks do not extend to the Picnic scenario, where chosen plaintext attacks are not in scope.

## 7.2 Attacks in the Single-User Setting

In the single-user setting, the attacker only ever sees a single key pair for the Picnic signature scheme, i.e., a single plaintext-ciphertext pair  $(f_u(x), u)$  of LOWMC with respect to a uniformly random key  $x$  and a uniformly random block  $u$ . Consequently, in this setting cryptanalytic results for LOWMC also directly apply to our scheme, and also the claimed bounds carry over to the Picnic signature scheme. Note that, one could even globally fix  $x$  to further shrink the size of the public verification key **pk**. However, we chose not to do so, as we also want to consider attacks in the multi-user setting (as discussed below).

## 7.3 Attacks in the Multi-User Setting

The multi-user setting more accurately models reality, in that there are multiple users, each with a public key, and the adversary is considered successful

if he can attack any one of the users.

**Multi-User EUF-CMA.** Even if single-user EUF-CMA security generically implies multi-user EUF-CMA (MU-EUF-CMA) security under a polynomial loss [GMS02], we put concrete focus on attack scenarios which become applicable by moving to the multi-user setting. Here, the adversary may see many signing key pairs and we need to be cautious with respect to more sophisticated attacks that might apply.

In particular—in contrast to the single-user setting—our decision to choose an independent and uniformly random block  $u$ , being the encryption function  $E(\cdot, u)$  of LOWMC, per signing key pair turns out to be important. This is because using the same, fixed block  $u$  with independent keys  $x_1, \dots, x_n$  for each of the  $n$  users would allow an adversary to apply multi-user key recovery attacks [Bih02] and generic time-memory trade-off attacks like [Hel80] and in particular time/memory/key trade-off attacks [BMS05]. In these attacks one of  $n$  block cipher keys may be recovered in less time than the cost of recovering a single key, and the attacks become more efficient for large  $n$ . Intuitively, the random block chooses a unique function per user, and work done to attack one user (function) can not be used to simultaneously attack another user (function). In addition, Banegas and Bernstein [BB17] have recently shown that parallel collision search attacks [vOW94] can also be applied in the quantum setting which also supports making a random choice of  $u$  per user. Finally, we note that one could choose a smaller value that is unique per user (with a potential decrease in security) to reduce the size of the public key. However, since public keys in our schemes are already small (at most 64 bytes), our design uses a full random block to be as conservative as possible.

**Key-Substitution Attacks.** These are attacks where an adversary who is given a signature  $\sigma_A$  on message  $M$  under  $A$ 's public key  $\text{pk}_A$  manages to come up with a public key  $\text{pk}_E$  (different from  $\text{pk}_A$ ) such that  $\sigma_A$  verifies under  $\text{pk}_E$  and message  $M$ . Menezes and Smart in [MS04] provide a formal model to cover such attacks, which are not covered by EUF-CMA security. We explicitly consider such attacks. Security against these types of attacks can be generically achieved. This has been shown in [MS04], and we recall their theorem below.

**Theorem 7.1** ([MS04], Theorem 6). *Let  $(\text{Gen}, \text{Sign}, \text{Verify})$  be an EUF-CMA secure signature scheme. Then,  $(\text{Gen}, \text{Sign}', \text{Verify})$  with  $\text{Sign}' := \text{Sign}(\text{sk}, \text{pk}||m)$*

and  $\text{pk}$  being an unambiguous encoding of the public key is a secure signature scheme in the multi-user setting.

We stress that the above result in particular holds for **s**EUF-CMA secure signature schemes.

As discussed in Section 3 (see also Section 4.3 of the specification), the public key is prepended to the message on signing, and the specification provides an unambiguous encoding (since the public key is a pair of bitstrings, the encoding is trivial). Consequently, we have the following corollary.

**Corollary 7.2.** *Picnic-FS, Picnic-UR and Picnic3 provide security in the multi-user setting in the sense of [MS04, Definition 6].*

## 7.4 Multi-Target Attacks

In [DKP<sup>+</sup>19a], Dinur and Nadler describe attacks against the version 1.0 specification of Picnic (the **Picnic-FS**, and **Picnic-UR** parameter sets). At the time, the **Picnic3** parameter sets were not specified, but the attack applies equally to a direct instantiation of the KKW protocol. Their attacks are multi-target attacks, where an attacker has a list of values of the form  $y_1 := H(x_1), \dots, y_S := H(x_S)$ , and recovering any of the  $x_i$  leads to a successful attack. Specifically, the  $x_i$  values are  $k$  bits long, and the attacker can recover the  $k$ -bit signing key. The  $y_i$  values can come from (i) a single signature (there are about  $2^7$  in a signature), (ii) from many signatures created by the same signer, or (iii) from signatures from multiple signers.

In Dinur and Nadler’s attack, the  $x_i$  values are the seeds used for each party, in each MPC instance. The function  $H$  expands  $x_i$  to a random tape, used during the MPC protocol simulation. In Picnic the seeds of two of three parties are revealed to the verifier, and in **Picnic3**  $n - 1$  of  $n$  are revealed. If an attacker learns the missing seed, they can recover the secret shared input, the signing key. Specifically  $H$  is the SHAKE XOF, and the output length depends on the parameter set, but is always above 600 bits.

What makes the attack non-obvious, is that the output of  $H$  (the random tapes) are not revealed directly. Dinur and Nadler show that given the states of the opened parties, it’s possible to solve for some of the unopened party’s random tape. They show this is a property of MPC protocols, that doesn’t affect the usual MPC security notion (Briefly, the MPC protocol must guarantee privacy of inputs, not the randomness, and leaking some randomness is

allowable provided it does not affect privacy). They also quantify the number of random bits that can be recovered from an MPC instance, and the cost. In all cases, it is possible to efficiently recover more than  $k$  bits, so that testing a candidate  $x_i$  value can be done. However, the bits from each target have different positions in the random tape, complicating an efficient implementation of the attack. In a typical multi-target attack, the attacker iterates over candidate values  $x'$ , computes  $y' = H(x')$ , then compares  $y'$  to the  $y_i$  efficiently using a data structure (e.g., a hash table or search tree). But here comparing the candidate tape  $y'$  (with all bits known) to the target tapes  $y_i$  (where a different subset of bits are known for each), is not obviously efficient. Dinur and Nadler show that it can be made efficient, and precisely quantify costs under various settings. In the best case, when all signatures are created by a single signer, their attacks cost  $2^{k-7}/S$  (information theoretically optimal). In other cases, the attack cost is higher, but still below the expected security level.

**Mitigation** The 2.0 version of the specification introduced a change to mitigate these attacks (in all parameter sets). The change adds additional information to the input of  $H$ , called a *salt*, so that a candidate seed  $x'$ , cannot be tested by comparing  $y'$  to all  $y_i$ , since each  $y_i$  is computed using a different salt. The salt ensures that  $y'$  would need to be recomputed with the correct salt before each comparison. To address all three variants of the attack, the salt must be unique per signature, per signer and per invocation of  $H$ . The first change is to choose a random per signature salt, 256 bits long. This ensures that (with high probability), the salt is unique across all signatures recorded by an adversary.

Then, to ensure that salts are unique within a signature, we also include a pair of counters, the first value corresponds to the MPC instance number, and the second corresponds to the invocation number of  $H$ . The specification already uses a domain separation technique, where different hash functions are created for different purposes, as follows,  $H_i(x) = H(i||x)$ . This mechanism also helps ensure salts are unique, e.g., when computing the seed tree we can use  $H_i$  and when computing the Merkle tree use  $H_j$ , and not worry about (salt, counters) pairs being repeated in both trees.

The main cost of the mitigation is a increase in signature size, of 256 bits, which is small relative to the overall size of the signature. Some additional data must also be hashed, but CPU costs in our benchmarks did not increase appreciably. This is likely due to the fact that hash inputs were short to begin

with, and remain short (smaller than the hash function block size), even with the salt.

## 8 Expected Security Strength

Since a Picnic keypair is a block cipher key and plaintext/ciphertext pair, we chose to define parameters at the L1/AES-128, L3/AES-192 and L5/AES-256 security levels. By the CFP, this means that L2 is implicitly defined by L3 and L4 is defined by L5.

We expect each parameter set to provide security equivalent to AES. For example, at L1, we expect 128-bits of security against classical attacks, and at least 64-bits against quantum attacks. Like AES, key recovery attacks using Grover’s algorithm against LOWMC are the best known quantum attacks on Picnic. The estimate of 64-bits comes from an idealized version of Grover’s algorithm capable of running for a long time. Like AES this may be conservative, even more so in the case of LOWMC, since the circuit is much larger than AES, due to the large amount of constant data required to implement it.<sup>6</sup> For example, at level L1, the Picnic LOWMC instance requires 34KBytes of constant data. These constants must either be encoded into the circuit, or the circuit must be expanded to recompute them on the fly. Thus the LOWMC circuit is orders of magnitude larger than AES, making attacking LOWMC with Grover’s algorithm at least as difficult as attacking AES.

In [JNRV20], Jaques et al. implement both AES and LowMC in Q#, and estimate the resources required to recover a key using Grover’s algorithm. They found that the gate cost for key recovery of an AES key is less than the cost a LowMC key (at the same security level). This confirms the above informal argument for the LowMC instances with partial S-box layers. It would be interesting to repeat their estimates with LowMC instances having a full S-box layer.

One exception is for Picnic3, at L5, where we chose 255-bit keys and block size, so that the block size is a multiple of three, as required to have a full S-box layer. This decreases security slightly when compared to AES, but simplifies implementations, and allows efficient use of 256-bit SIMD instructions.

---

<sup>6</sup>Previously, when we compared the LOWMC circuit size to AES, we were looking only at AND gates, but here we’re considering all gates.



As L5 is expected to provide a considerable amount of security margin, we see this choice as a reasonable tradeoff.

We similarly set the number of parallel iterations to provide 128-bit security against classical attacks, and 64-bits against quantum attacks using an idealized version of Grover’s algorithm. Like with LOWMC, the circuit required to break ZKB++ soundness with Grover’s algorithm is orders of magnitude larger than the AES circuit.

In more detail, suppose an attacker is trying to forge a ZKB++ proof as a generic search problem. In particular, if an attacker can find a permutation of a set of transcripts that hash to a challenge chosen in advance, he can forge a proof. Consider  $T$  parallel repetitions (Picnic-FS specifies 219, 329 and 438 rounds at levels L1, L3 and L5, resp.). Then there are  $3^T$  possible challenges that can come from hashing those  $3T$  transcripts (since there are 3 challenges).

Now consider an attacker who constructs invalid ZKB++ “proofs” such that for each parallel repetition, he can give a valid response to two of the challenges but not the third. If we model the hash function as a random oracle, the probability of getting a challenge for which he can respond is  $(\frac{2}{3})^T$ , and thus we expect that if the attacker searches a space of  $(\frac{3}{2})^T$  candidates (i.e., permutations of transcripts that are constructed in this manner) he can find one. Grover’s algorithm allows the attacker to search the space in time  $(\frac{3}{2})^{T/2}$ .

However, the items in the space are larger, and changes in one value (e.g., a seed value) requires re-computing many others (the random tape, the MPC transcript, and the commitments). Clearly this is far more expensive than a single AES evaluation, and so we assume that Grover’s algorithm applied to breaking ZKB++ soundness is at least as costly as AES key recovery.

**Using Larger LowMC Keys** We could reduce the feasibility of Grover key recovery attacks against LOWMC by increasing the keysize and keeping the block length fixed. There would still be a chance to use Grover’s algorithm to break the soundness of ZKB++ (unless the number of parallel iterations was increased). However, a quick inspection reveals that the computation of attacking soundness is computationally much more complex than attacking LOWMC. For example, checking whether a candidate secret key corresponds to a given public key requires one LOWMC evaluation, while checking whether a set of cheating commitments leads to a challenge that does not catch the cheating requires hundreds of SHA3 computations.

## 8.1 LowMC Parameter Selection

The choice of LowMC parameters may seem aggressive in the context of a general-purpose block cipher. The LowMC spec recommends an additional 1.3 times the number of rounds, as a security margin against unknown attacks. Picnic does not use these additional rounds.

The general block cipher security definition gives attackers as much power as possible, to model the worst case scenario. Consider the CPA security game, where attackers may choose plaintexts, query the encryption oracle many times, and must only distinguish encryptions of chosen plaintexts in order to break the cipher (as opposed to recovering plaintext or private keys). This strong security definition is sensible when the primitive will be used in a variety of (potentially unforeseen) applications.

By contrast, for the security of Picnic signatures, attackers are severely restricted. They are given a single plaintext/ciphertext pair, for a randomly chosen block and key, and succeed if they can recover the key. Signatures are zero-knowledge proofs, so by definition provide no additional useful information. Thus the success criteria for the LowMC attacker in our context is key recovery, which is more difficult than indistinguishability, while the capabilities are more limited. In this context, the parameters we have chosen for LowMC are arguably very conservative.

Further, it is difficult to quantitatively support parameter selection in our restricted attack model, since most research focuses on the standard security definition. Our claim is that the complexity of a key recovery attack against LowMC, given only a Picnic public key, is at least as difficult as attacking the CPA security of AES (for equivalent key size, and with Picnic the block size always matches the key size).

## 8.2 Hash Function Security

Picnic depends on secure hash functions when computing signatures, for commitments and the challenge. In our security analysis we have modeled these as random oracles. While choosing parameters we also took into account some more specific security properties (all implied by a random oracle). In practice, it is still unclear which properties are necessary and which are sufficient. For example, preimage resistance is clearly necessary (since we commit to the seed of the unopened party, and it must remain hidden for security), but collision resistance may not be.

All hash functions are implemented with SHAKE128 with 256-bit digests at security level L1, and SHAKE256 at levels L3 and L5, with 384 and 512-bit digests, respectively.<sup>7</sup> We expect the concrete security provided by SHAKE for collision and preimage resistance claimed in [NIS15], extended to quantum attacks.

For preimage resistance, in the classical case it is common to assume  $O(2^n)$  operations for standard hash functions. When considering quantum algorithms, Grover’s algorithm can find preimages with  $O(2^{n/2})$  operations. Therefore, we assume that our uses of SHAKE128 and SHAKE256 provide this level of preimage resistance.

When considering quantum algorithms, in theory it may be possible to find collisions using a generic algorithm of Brassard et al. [BHT98] with cost  $O(2^{n/3})$  (for  $n$ -bit digests). A detailed analysis of the costs of the algorithm in [BHT98] by Bernstein [Ber09] found that in practice the quantum algorithm is unlikely to outperform the  $O(2^{n/2})$  classical algorithm. Multiple cryptosystems have since made the assumption that standard hash functions with  $n$ -bit digests provide  $n/2$  bits of collision resistance against quantum attacks (for examples, see papers citing [Ber09]). We make this assumption as well, and in particular, that SHAKE128 with 256-bit digests provides 128 bits of PQ security, SHAKE256 with 384-bits provides 192-bits and SHAKE256 with 512-bit digests provides 256-bits.

## 9 Advantages and Limitations

### 9.1 Compatibility with Existing Protocols

Here we describe some work we did to demonstrate compatibility of Picnic signatures existing protocols protocols, TLS, and X.509. We also prototyped protecting Picnic private key operations on a commercial hardware security module.

---

<sup>7</sup>We initially considered using SHAKE256 at all three levels for simplicity, but L1 signing and verify times increased by roughly 10%. For Picnic3, the increase would be even more significant, as hashing accounts for a larger fraction of CPU time [KZ20b].

## 9.2 TLS and X.509 Compatibility

The optimized implementation of Picnic has been integrated with the Open Quantum Safe (OQS) project.<sup>8</sup> Then, using a modified version of OpenSSL<sup>9</sup> modified to use OQS, we were able to create X.509 certificates signed with Picnic and certificates certifying Picnic public keys. These keys and certificates were then used to establish TLS 1.2 connections, where the key exchange algorithm was one of the the LWE-Frodo or SIDH algorithms from OQS. To our knowledge, these were the world’s first TLS connections to use both post-quantum secure key exchange and authentication algorithms.

OpenSSL had to be patched in one place to handle larger signature sizes, since the TLS 1.2 standard has a limit of  $2^{16} - 1$  bytes. All Picnic3 signatures are below this limit, but Picnic-FS and Picnic-UR signatures at L3 and L5 exceed it and we would likely need an extension to support these signatures. Ideally a future version of TLS would allow larger signatures, but this is not pressing as we expect the L1 parameter set to provide sufficient security in the short and medium term, and the shorter signatures of Picnic3 are likely preferable in TLS. Otherwise the integration was smooth, and performance seemed acceptable in our limited experiments. In particular the certificate stack (X.509/ASN.1) worked unmodified with signatures this large, something we did not expect.

Earlier versions of this document included some benchmarks of file transfers over HTTPS, using OQS and the Apache web server. These experiments have been made obsolete by newer experiments published in the open literature, see e.g. [PST20].

## 9.3 Hardware Security Module Compatibility

Cryptographic keys are often protected by specialized hardware known as *hardware security modules* (HSMs). An HSM is a tamper-resistant device that stores keys and performs operations in response to calls to a limited API. The primary security goal is that private keys never leave the device (with exceptions for backup and export to other similar devices). Secondary goals are tamper proof logging and isolation of cryptographic operations from the rest of the system, forming a strong security boundary. If an attacker

---

<sup>8</sup><https://github.com/open-quantum-safe/liboqs>

<sup>9</sup><https://github.com/christianpaquin/openssl>

compromises the system, they can use the key, but cannot export it for use on another system, and ideally, cannot use it without leaving logs.

For signature keys, example operations are generating keys and signing (digests). Upon key generation, a key identifier is output, which can be used in sign calls to refer to the private key (that may be marked non-exportable). An example API is PKCS #11, standardized so that applications can be agnostic of the underlying hardware.

Many such devices are available on the market, with a range of features, performance and security hardening. They may be small peripheral devices similar to a smartcard, or standalone, network connected servers. Often the firmware on these devices is fixed by the manufacturer, and prototyping new algorithms is not possible. One device that does allow the owner to provide some custom firmware implementing new cryptographic algorithms is the Utimaco SecurityServer Se50 LAN V4. We added support for Picnic on this device and describe our experience here.

On the Utimaco HSM, the owner may (i) use the provided cryptographic modules (e.g., RSA, ECDSA, SHA-256, and RNG) in the firmware, (ii) write their own cryptographic module that doesn't depend on any of the provided modules, or (iii) write a module implementing a new primitive that uses some of the provided modules. To implement Picnic on the HSM, we experimented with option (iii). In particular we leveraged the RNG, ASN.1 and SHA3 modules from Utimaco, and implemented the remainder of our spec in a custom module (named PICNIC). The PICNIC module was a port of our reference implementation that replaced the RNG and SHA3 with calls to the Utimaco modules. Additionally we added a module named CERT, that uses the PICNIC and ASN1 modules to create X.509 certificates signed with Picnic.

Whether it is desirable to create certificates on the HSM, vs. creating them on the host application and using the HSM only to sign them is debatable. On one hand having more functionality in the HSM may allow CA policy to be better enforced in the event of a compromise of the host application. On the other hand, having more functionality in the HSM increases its attack surface.

Since having the HSM be a limited signing oracle is a special case of having it create certificates, in our demo we created certificates on the HSM. If this more complicated design can be demonstrated to work, then the simpler case should also work.

The certificates are standard X.509 v3 certificates, but with custom object

identifiers (OIDs) for Picnic keys and signatures. We confirmed that the resulting certificates parsed correctly in existing viewers (e.g., <http://www.lapo.it/asn1js/>).

Typically the sign API of an HSM accepts a hash of the data to be signed. In contrast, the sign API of the Picnic spec accepts the (unhashed) data to be signed, and hashes it internally. The Picnic spec allows digests to be signed, but this is intended only for very large messages, as it requires stronger security properties from the hash function (as with other Schnorr-like signatures, with EdDSA being another example, see RFC 8032). In our demo, we did not have issues sending larger amounts of data to the HSM (we tested up to roughly 10k bytes), so we expect the sign API without pre-hashing the message to work for most PKI scenarios.

**Scenario: Post-Quantum Public-Key Infrastructure.** We implemented a small demo to test Picnic in a public-key infrastructure (PKI) scenario, from the perspective of a certificate authority (CA).

The demo architecture has two main software components.

1. A host application, running on a Windows PC.
2. A pair of firmware modules, running on the HSM, housed in a machine room in a building across the street.

The application is connected to the HSM with a 100MBit connection. The latency of a no-op call between host and HSM was about 24 ms. We also tried running the host application from a laptop at home where latency was about 200ms, and saw larger variance in the roundtrip times for operations but similar mean times.

We implemented the following CA operations.

1. The HSM generates and stores new Picnic key pair, and creates a self-signed certificate. This is the root certificate of the PKI.
2. The host application generates and stores a new Picnic end-entity (EE) key, and then the CA issues a certificate for the EE key using the signing key from the previous step. This is not a typical CA operation, but was a useful stepping stone during development.
3. The host application sends a certificate signing request (CSR) to the HSM. The CSR is created with OpenSSL, and the subject public key

is an RSA key pair. The HSM issues a cert for the RSA public key, signed with Picnic.

For item 3, it would have been more realistic to use Picnic as the subject public key and sign the CSR with Picnic as well, but our version of OpenSSL (the OQS OpenSSL fork described above), did not support creating CSRs with post-quantum algorithms yet.

Our new modules were tested in the HSM simulator first, then cross-compiled for the HSM itself and uploaded.

**Software.** The software we used for this experiment is available at <https://microsoft.github.io/Picnic/> under the MIT license. Note that although our code is MIT licensed, building it and running it (even in the simulator), may require a license for software and tools from Utimaco.

The main effort involved was porting the reference implementation to build with the HSM's tools. It took between two and three person-weeks, and this included time to get familiar with the HSM tools and development process. Porting the Picnic library to the HSM required the following steps:

- Create an empty HSM module from the HSM's SDK.
- Add the Picnic source and header files to the module code and update the module's makefile for the c6000 cross-compiler.
- Replace the standard C libraries with the HSM provided libraries and update calls where the names differ. Most of these changes deal with memory management and string handling in error handling code.
- Update the RNG and SHA-3 calls in Picnic to use the HSM modules for these functions.
- The c6000 compiler is C89 compliant and lacks features of more modern C standards. Several small changes were required dealing mainly with variable declarations.
- Create a new public interface for the Picnic module that more closely resembles the other HSM's crypto modules for consistency.
- When the code is building and functioning properly in the HSM simulator, cross-compiled the code for the HSM; sign it; load it into the HSM and verify it is working correctly.

**Performance and Discussion.** The goal of this prototype was to demonstrate that using post-quantum signatures in a PKI scenario is practical, and that there are no major impediments to deployment even with existing commercially available HSM hardware. In particular, using new types of keys, and creating signatures with a new algorithm, having larger signatures than traditional algorithms, and hashing the message on the HSM was possible, and did not pose significant engineering challenges.

Some benchmarks are given in Table 3. We signed messages of three sizes, 100B, 1KB and 10KB, covering the range of message sizes we would expect in our PKI scenario. We measured the round trip time of a call to the HSM from the host application. There was no significant difference for the different message sizes. For the L1-FS parameter set, the round trip time is about half a second, and goes up to about four seconds for the L5-FS parameter set.

For many PKI applications the number of certificates created is relatively small, and this performance would be considered adequate. We stress that this is an unoptimized implementation, and we don’t have a breakdown of where the time was spent, i.e., network time vs. computation time on the HSM. For improving computation time, using our optimized implementation would be a first step.

	Sign 100B	Sign 1KB	Sign 10KB	Keygen
L1-FS	0.4474	0.4477	0.4504	0.0050
L3-FS	1.6478	1.6474	1.6509	0.0069
L5-FS	4.0854	4.0841	4.0860	0.0096

Table 3: Mean round trip times in seconds for calls to an HSM creating Picnic-FS signatures (average of 10 calls).

## 10 Additional Security Properties

### 10.1 Side-Channel Attacks

**Key Generation.** Key generation requires generating a random LOWMC key and plaintext, and computing the LOWMC block cipher. A fast implementation of the LOWMC block cipher may use precomputed data, and have



cache-timing side channels, because the access pattern depends on the secret key. However, there are a couple mitigations:

1. Since key generation happens infrequently, a slower LOWMC implementation with a constant access pattern can be used without performance penalty.
2. Even if a side-channel is present in key generation, since only one encryption with a given secret key is ever computed, and known attacks require observing multiple runs, a successful attack is unlikely.

**Signing.** Generally speaking, since the multi-party protocol simulated during signing is circuit-based, the same operations are performed, regardless of the values on the input wires of the circuit.

Signing is not constant time in the absolute sense, but is constant time with respect to the operations that depend on the ephemeral random values (that in turn depend on the signing key). The timing (and signature size) variation is due to the different operations performed depending on the (public) bits of the challenge. The optimized implementation is constant time relative to the secret key.

The Picnic design has what seems like a natural mitigation to some side-channel attacks, since the LOWMC secret keys used by each of the players is a randomized secret sharing of the actual key. The shares are only ever used once, and the randomization means that access pattern information learned by an attacker in one parallel iteration of ZKB++ or KKW can not be combined with information from other iterations. Since many side-channel attacks require multiple observations (called traces), we expect this mitigation to be effective.

Note that the derandomized variant of the sign algorithm (i.e., where the per-signature randomness is derived from the message to be signed and the secret key), may in fact use the same shares multiple times, when signing the same message. Therefore if a side-channel attacker can repeatedly cause a signature to be computed on the same message, while observing the signing device, they may be able to collect multiple traces. The Picnic spec recommends randomized signatures, and derives ephemeral random values from the secret key and additional entropy. Our implementations do not do this at the moment, to allow for easier testing.

The circuit decomposition technique is similar to the side channel countermeasure called masking, commonly used to protect block cipher imple-

mentations from side channel attacks. An early, well cited paper on the topic is Goubin and Patarin [GP99]. With further study, we may find that the ZKB++ circuit decomposition provides other types of side channel resistance.

This is not the case for differential power analysis (DPA) attacks which were recently demonstrated against the Picnic reference implementation by Gellersen et al. [GSE20]. At a high-level, these attacks probe internal values computed by the signer, for example, the unopened party’s share of the secret key created at the beginning of ZKB++, which leaks one key bit (when combined with the other shares revealed in the signature). After repeating the attack sufficiently many times the private key is recovered. While the demonstration used the reference implementation, it’s likely the attack could be mounted against an optimized implementation, without protections.

For protections against DPA attacks, one direction explored by Seker et al. [SBWE20] is to modify the MPC protocol so that more than one party remains unopened. This is elegant and can provide a strong security guarantee, however the resulting signatures are not compatible with the Picnic spec, are significantly larger, and are slower to create and to verify. We are investigating a mitigation for DPA attacks that complies with the specification and can be used transparently by the signer. The ability to randomize Picnic signatures is an important part of these protections.

**Fault Attacks** Fault attacks on deterministic signatures were described in [PSS<sup>+</sup>17, ABF<sup>+</sup>17]). The paper by Aranha et al. [AOTZ20], analyses the “hedged” construction recommended by the Picnic specification, i.e., where the per-signature random values are derived from the secret key and an additional nonce value. They show that this construction provides meaningful resistance to fault attacks (and consider Picnic2/Picnic3 in depth). The model used by [AOTZ20] does not consider arbitrary faults to any part of signature computation, but does rule out a large class of faults, including many fault attacks on Fiat-Shamir type signature schemes in the literature. Further analysis is required to determine the impact of faults in other parts of the computation, and if necessary, to design mitigations.

## 10.2 Security Impact of Using Weak Ephemeral Values

The specification allows the per-signature random values used when computing a signature be derived from the signing key and the message, to simplify

testing and to mitigate the security impact of a defective random number generator during signing. The goal is that signatures are secure, provided the random number generator was secure during key generation.

Like (EC)DSA, given the random values used when computing a signature, it is possible to recover the signer’s secret key. Recall that in each parallel iteration of ZKB++, three seed values are generated, one for each party in the MPC protocol. In normal operation, two of these seed values are revealed, and one is kept secret. The MPC protocol remains secure when two of the three parties are corrupted (i.e., have their seed exposed, which exposes their state, input and output shares). Given the entire third seed, it is possible to recover the input share of the third player, and recover the secret key.

Unlike (EC)DSA, slight biases in the random number generator do not allow the secret to be recovered from multiple signatures. This is because the seed values are never used directly; they are always expanded with an extendable output function (XOF) into a random tape, and the random tape values are used.

Regardless of how the seeds are generated, a bias in the XOF output may lead to an attack. For example, if the XOF/PRF used to derandomize (EC)DSA, EdDSA, and other ElGamal-like signatures was biased, the ephemeral value is biased, and the lattice attacks studied in the context of RNG biases apply [Ble00, HS01]. For Picnic, it’s not clear if this could be exploited.

### 10.3 Parameter Integrity

With some cryptographic primitives if the system parameters are changed, security is lost. For example, if an elliptic curve secret key is used on a weak curve, the primitive may still work, but leak the secret key.

In Picnic, the parameters are small integer values like the parallel repetition count that tend to be hard-coded in software, and the LOWMC matrices and constants. Using weak parameters for LOWMC could weaken the cipher to the point where key recovery attacks are feasible. If weak parameters are used for key generation it is possible to generate a weak keypair, however, it will not produce signatures that verify with respect to the correct parameters. So the common PKI practice of signing a certificate request with the subject key will catch keys generated with invalid parameters.

Creating signatures with invalid parameters can also be a security risk.

Suppose a set of weak LOWMC parameters were used, so that the signing algorithm proves knowledge of the signing key  $k$ , for an new circuit  $E'$ , i.e.,  $E'_k(x) = \text{pk}$ , where  $E'$  is LOWMC with invalid parameters. The signature would be invalid in most cases, unless key generation and verification also used  $E'$ . However, the invalid signature contains enough information to recover  $E'_k(x)$ , from which it may be possible to recover  $k$ .

## 11 Efficiency and Memory Usage

This section gives performance benchmarks of the Picnic signature scheme. A single core/thread was used for all benchmarks. All times are in milliseconds. We benchmark three implementations:

**Reference.** An expository C implementation. Makes no performance optimizations.

**Optimized-C.** A somewhat optimized implementation using C only.

**Optimized.** An optimized implementation that uses processor-specific compiler intrinsics for vector instructions, e.g. SSE2 and AVX2 on Intel x86-64 and NEON on ARM v8.

The **optimized-C** and **optimized** implementations are constant-time.

### 11.1 Description of the Benchmark Platforms

#### 11.1.1 Platform A

The primary benchmarking platform, **Platform A**, has the following specifications:

**CPU.** Intel(R) Xeon(R) W-2133 CPU @ 3.60GHz

**Memory.** 32 GB

**OS.** Ubuntu 18.04.5

**Compiler.** GCC 7.5.0

Intel Turbo Boost (dynamic frequency scaling) was disabled. For comparison, OpenSSL version 1.1.1 reports 0.02 ms for ECDSA signing and 0.06 ms for verification on Platform A<sup>10</sup>.

### 11.1.2 Platform B

The secondary benchmarking platform, **Platform B** (Raspberry Pi 3 Model B), has the following specifications:

**CPU.** Quad Core 1.2GHz Broadcom BCM2837 64-bit CPU, ARM Cortex A53 (ARMv8)

**Memory.** 1 GB RAM

**OS.** openSUSE Tumbleweed

**Compiler.** GCC 10

For comparison, OpenSSL version 1.0.2j reports 11 ms for ECDSA signing and 40 ms for verification on Platform B.

## 11.2 Description of the Benchmarking Methodology

Timing results for key generation, signing and signature verification were averaged over 100 runs.

On Platform A we measured CPU cycles using the `perf_event` performance monitoring subsystem of the Linux kernel. On Platform B CPU cycles were measured using the hardware performance counter available via the `MRS` instruction. The optimized implementations will use the GCC link time optimization (`-flto`) feature by default if it is available (it was available on both of our benchmarking platforms).

## 11.3 Benchmark Results: Sizes

In Table 4 we give the size of Picnic keys, and signatures. These are the same for all implementations.

---

<sup>10</sup>As reported by the command `openssl speed ecdsap256`

Parameter Set	Pub. key	Priv. key	Sig (max)	Sig (avg., std. dev.)
Picnic-L1-FS	32	16	34032	32838, 107
Picnic-L1-full	34	17	32061	30827, 115
Picnic-L1-UR	32	16	53961	
Picnic3-L1	34	17	14608	12437, 249
Picnic-L3-FS	48	24	76772	74134, 198
Picnic-L3-full	48	24	71179	68571, 227
Picnic-L3-UR	48	24	121845	
Picnic3-L3	48	24	35024	27474, 394
Picnic-L5-FS	64	32	132856	128176, 315
Picnic-L5-full	64	32	126286	121598, 338
Picnic-L5-UR	64	32	209506	
Picnic3-L5	64	32	61024	48452, 771

Table 4: Key and signature sizes (in bytes) by security level. For the FS variants, the signature length varies based on the challenge, therefore we give the maximum possible size, along with the average size and standard deviation computed over 100 signatures.

## 11.4 Benchmark Results: Timings

In this section we describe the time required for various operations, on the two benchmark platforms. In all tables presented in this section we give the timing information as milliseconds and CPU cycles.

In Tables 5, 6, and 7 we present the benchmark results of all implementations on Platform A. On this platform we observe speed improvements of the **optimized** implementation over the **optimized-C** implementation by a factor of about 2 for Picnic and 1.75 for Picnic3. This is because the **optimized** implementation can make use of the vector instructions (AVX2) available on this platform.

Next we give the results of the evaluation of our implementations on Platform B in Tables 8, 9, and 10. Again we observe improved performance figures for the **optimized** implementation, but they are not as significant as on Platform A. GCC’s optimizer vectorizes the code on Platform B more aggressively for **optimized-C** implementation and thus the performance gains of the **optimized** implementation are smaller.

Parameters	Keygen	Sign	Verify
Picnic-L1-FS	0.04	34.55	22.09
(cycles)	160129.49	124373587.72	79512545.88
Picnic-L1-full	0.02	20.18	13.12
(cycles)	81076.26	72659224.27	47232986.45
Picnic-L1-UR	0.04	41.93	27.77
(cycles)	147509.86	150954639.61	99967363.80
Picnic3-L1-FS	0.02	106.94	76.99
(cycles)	55480.13	384985571.38	277152107.39
Picnic-L3-FS	0.09	109.54	71.67
(cycles)	323583.45	394343487.96	258018933.56
Picnic-L3-full	0.03	41.08	26.78
(cycles)	112265.57	147874256.94	96398968.61
Picnic-L3-UR	0.09	133.54	88.82
(cycles)	326509.29	480737543.63	319766930.39
Picnic3-L3-FS	0.02	256.63	179.56
(cycles)	76592.54	923861427.79	646403653.90
Picnic-L5-FS	0.16	239.90	159.67
(cycles)	562732.01	863626029.60	574822319.79
Picnic-L5-full	0.13	184.81	121.74
(cycles)	464065.35	665321291.35	438266752.58
Picnic-L5-UR	0.16	271.05	183.52
(cycles)	571718.36	975779299.24	660673848.37
Picnic3-L5-FS	0.12	541.62	340.11
(cycles)	415263.59	1949836770.58	1224406678.22

Table 5: Benchmarks for the **reference** implementation, on benchmark Platform A.

Parameters	Keygen	Sign	Verify
Picnic-L1-FS	0.00	2.70	2.26
(cycles)	11036.87	9724857.39	8147317.71
Picnic-L1-full	0.00	2.21	1.68
(cycles)	6988.01	7942496.51	6039706.22
Picnic-L1-UR	0.00	3.37	2.79
(cycles)	11316.85	12130039.97	10054171.75
Picnic3-L1-FS	0.00	8.89	6.98
(cycles)	7447.73	32014574.17	25123857.04
Picnic-L3-FS	0.01	6.49	5.53
(cycles)	20974.63	23379304.18	19921502.30
Picnic-L3-full	0.00	4.30	3.35
(cycles)	10014.17	15463312.42	12048088.52
Picnic-L3-UR	0.01	8.38	6.96
(cycles)	21262.05	30159969.27	25066762.87
Picnic3-L3-FS	0.00	19.51	15.26
(cycles)	11119.27	70228447.71	54939616.58
Picnic-L5-FS	0.01	12.02	10.41
(cycles)	30621.56	43267651.79	37472161.57
Picnic-L5-full	0.00	7.45	6.05
(cycles)	15426.10	26819023.59	21783708.90
Picnic-L5-UR	0.01	14.53	12.48
(cycles)	31177.49	52290516.47	44928989.96
Picnic3-L5-FS	0.01	33.51	24.31
(cycles)	18353.81	120628101.19	87502559.15

Table 6: Benchmarks for the **optimized-C** implementation, on benchmark Platform A.



Parameters	Keygen	Sign	Verify
Picnic-L1-FS	0.00	1.37	1.11
(cycles)	5599.14	4924342.55	3982244.42
Picnic-L1-full	0.00	1.00	0.80
(cycles)	3680.38	3601664.60	2871980.78
Picnic-L1-UR	0.00	1.81	1.47
(cycles)	6418.02	6502156.22	5305718.37
Picnic3-L1-FS	0.00	5.07	3.84
(cycles)	4151.67	18252055.62	13811201.62
Picnic-L3-FS	0.00	3.20	2.63
(cycles)	10479.34	11509382.57	9452289.86
Picnic-L3-full	0.00	1.95	1.56
(cycles)	5045.94	7008817.04	5626166.82
Picnic-L3-UR	0.00	4.13	3.40
(cycles)	10694.91	14875862.02	12231304.32
Picnic3-L3-FS	0.00	10.44	8.12
(cycles)	6567.42	37595772.02	29243365.55
Picnic-L5-FS	0.00	5.58	4.65
(cycles)	15255.58	20085119.47	16722292.75
Picnic-L5-full	0.00	3.15	2.56
(cycles)	6701.83	11351041.44	9217982.76
Picnic-L5-UR	0.00	6.99	5.83
(cycles)	17118.46	25178763.70	20998784.44
Picnic3-L5-FS	0.00	18.21	13.02
(cycles)	9504.91	65555710.67	46887830.87

Table 7: Benchmarks for the **optimized** implementation, on benchmark Platform A.

Parameters	Keygen	Sign	Verify
Picnic-L1-FS	0.37	230.91	151.61
(cycles)	446279.44	277086642.56	181926064.21
Picnic-L1-full	0.23	142.68	93.18
(cycles)	271790.80	171214339.46	111810320.53
Picnic-L1-UR	0.34	284.00	191.07
(cycles)	404352.34	340800269.78	229279882.65
Picnic3-L1-FS	0.12	750.84	537.23
(cycles)	138367.17	901005731.88	644670514.63
Picnic-L3-FS	0.71	758.85	500.67
(cycles)	851056.51	910614342.57	600808483.90
Picnic-L3-full	0.36	291.01	189.67
(cycles)	430530.45	349210815.91	227609453.45
Picnic-L3-UR	0.76	913.92	615.44
(cycles)	912877.21	1096708491.56	738527069.98
Picnic3-L3-FS	0.15	1795.73	1253.93
(cycles)	183778.21	2154875893.05	1504721550.79
Picnic-L5-FS	1.33	1852.79	1238.73
(cycles)	1592241.17	2223352819.68	1486476927.61
Picnic-L5-full	1.14	1377.89	911.15
(cycles)	1370000.00	1653465266.19	1093375164.52
Picnic-L5-UR	1.42	2073.68	1405.47
(cycles)	1701187.10	2488417655.44	1686562214.88
Picnic3-L5-FS	0.86	3859.85	2408.84
(cycles)	1032647.72	4631815289.83	2890611251.56

Table 8: Benchmarks for the **reference** implementation, on benchmark Platform B.

Parameters	Keygen	Sign	Verify
Picnic-L1-FS	0.02	16.55	14.30
(cycles)	26565.18	19865298.18	17164909.19
Picnic-L1-full	0.01	13.28	11.04
(cycles)	17708.52	15940768.94	13253260.19
Picnic-L1-UR	0.02	19.93	17.06
(cycles)	26791.46	23912186.85	20470251.41
Picnic3-L1-FS	0.02	56.16	41.41
(cycles)	22005.76	67389337.91	49686779.77
Picnic-L3-FS	0.05	46.11	43.25
(cycles)	59549.95	55328821.49	51903410.39
Picnic-L3-full	0.02	23.31	22.15
(cycles)	23449.58	27974688.43	26581559.82
Picnic-L3-UR	0.06	55.56	50.71
(cycles)	67502.38	66669926.47	60857222.54
Picnic3-L3-FS	0.02	120.27	88.71
(cycles)	27936.83	144324621.25	106456854.10
Picnic-L5-FS	0.07	83.19	80.37
(cycles)	85897.69	99830339.56	96448612.16
Picnic-L5-full	0.03	40.45	43.54
(cycles)	39607.97	48537053.70	52249412.54
Picnic-L5-UR	0.11	96.49	91.21
(cycles)	128217.43	115792544.02	109456869.64
Picnic3-L5-FS	0.04	207.81	146.90
(cycles)	50803.04	249372729.64	176276189.51

Table 9: Benchmarks for the **optimized-C** implementation, on benchmark Platform B.

Parameters	Keygen	Sign	Verify
Picnic-L1-FS	0.02	15.34	13.00
(cycles)	25928.85	18413331.23	15599045.71
Picnic-L1-full	0.01	12.18	9.76
(cycles)	15718.34	14619369.89	11709251.51
Picnic-L1-UR	0.02	18.76	15.78
(cycles)	26859.73	22508150.15	18934098.72
Picnic3-L1-FS	0.02	56.51	41.31
(cycles)	18903.88	67812435.00	49575453.45
Picnic-L3-FS	0.05	45.19	40.22
(cycles)	58904.18	54228071.33	48265388.83
Picnic-L3-full	0.02	23.65	19.25
(cycles)	22815.95	28384012.31	23104784.11
Picnic-L3-UR	0.06	54.69	47.74
(cycles)	75804.25	65626022.88	57290433.77
Picnic3-L3-FS	0.02	120.80	89.66
(cycles)	26956.40	144965473.00	107586444.00
Picnic-L5-FS	0.07	77.54	69.56
(cycles)	82322.36	93046271.92	83470511.50
Picnic-L5-full	0.02	38.91	32.18
(cycles)	29298.35	46690401.09	38619506.05
Picnic-L5-UR	0.11	90.99	80.65
(cycles)	132437.46	109187306.51	96785474.50
Picnic3-L5-FS	0.04	206.54	143.38
(cycles)	48456.05	247846585.82	172052611.28

Table 10: Benchmarks for the **optimized** implementation, on benchmark Platform B.

## 11.5 Memory Requirements

In this section we give the memory requirements for our implementations. The memory requirements of an implementation are assumed to be the same for all platforms.

We note that at this time, our implementations have not been optimized to reduce memory consumption.

### 11.5.1 Reference Implementation Detailed Memory Usage

Memory usage was benchmarked using the Valgrind<sup>11</sup> tool Massif. Massif was run on an example program, that generates a key pair, creates a signature, then verifies it, using the API in `picnic.h`. Then the tool `massifcherrypick`<sup>12</sup> was used to determine the peak memory usage of specific functions.

Massif was invoked with the command:

```
valgrind --tool=massif --stacks=yes ./example
```

When creating a signature, peak memory usage ranged from about 182K to 5.3M bytes, as shown in Table 11, while verification ranged from about 132K to 5.2M bytes.

Massif measures memory usage by sampling so there is some variability in these measurements. The variance was low so these are a reasonable estimate, since we're only interested in peak usage.

### 11.5.2 Optimized Implementation Detailed Memory Usage

With the same methodology as used for the reference implementation, peak memory usage data was generated with `valgrind`. The data is presented in Table 12.

## 11.6 Size of Precomputed Constants and Data

The LOWMC block cipher uses a large amount of constant data when compared to traditional block ciphers. This data may be computed on-the-fly as needed, or precomputed and stored in advance. All our implementations compile this data into the binary.

---

<sup>11</sup><http://valgrind.org/docs/manual/ms-manual.html>

<sup>12</sup><https://github.com/lnishan/massif-cherrypick>

Parameter set	Sign	Verify
Picnic-L1-FS	190,148	132,247
Picnic-L1-full	182,953	135,123
Picnic-L1-UR	268,080	239,367
Picnic3-L1	1,206,160	1,164,112
Picnic-L3-FS	380,380	290,507
Picnic-L3-full	368,568	265,088
Picnic-L3-UR	552,752	492,779
Picnic3-L3	2,873,728	2,748,832
Picnic-L5-FS	632,544	468,162
Picnic-L5-full	606,264	435,648
Picnic-L5-UR	984,696	733,426
Picnic3-L5	5,387,928	5,294,136

Table 11: Peak memory usage (stack and heap combined) of reference implementation, in bytes. This excludes memory used for the LOWMC constants, which was stored as static data in program binary (see §11.6).

Parameter set	Sign	Verify
Picnic-L1-FS	140,792	86,448
Picnic-L1-UR	200,243	127,722
Picnic-L1-full	133,010	81,951
Picnic3-L1	1,282,910	497,974
Picnic-L3-FS	284,664	166,688
Picnic-L3-UR	427,779	264,234
Picnic-L3-full	260,976	153,032
Picnic3-L3	2,994,400	1,441,048
Picnic-L5-FS	463,040	252,768
Picnic-L5-UR	707,006	434,100
Picnic-L5-full	442,016	245,760
Picnic3-L5	5,618,712	3,011,856

Table 12: Peak memory usage (stack and heap combined) of the optimized implementation, in bytes. This excludes memory used for the LOWMC constants, which was stored as static data in program binary (see §11.6).

We did not investigate the cost of re-computing the LOWMC constants at runtime. The output of the Grain LSFR is used as a self-shrinking generator to create the constants.

The **optimized-C** and **optimized** implementations use an alternative but equivalent representation of LOWMC, for instances with a partial S-box layer. They implement an optimized linear layer [DKP<sup>+</sup>19a], which allows to greatly reduced the size of the LOWMC constants and also gives a significant performance boost. The sizes of those matrices are are given in Tables 14 and 14 in the Key Matrices and Linear Matrices columns. The total reduction in size of precomputed constants ranges from 2.4 to 4.9 times. For the parameters with a full S-box layer, the optimization from [DKP<sup>+</sup>19a] only helps at security level L3.

Params	Linear Matrices	Round Constants	Key Matrices	Total
L1	40,960	320	43,008	84,288
L1-full	8,320	64	10,400	18,784
L3	138,240	720	142,848	281,808
L3-full	27,648	144	34,560	62,352
L5	311,296	1,216	319,488	632,000
L5-full	32,768	128	40,960	73,856

Table 13: Size of constant data (in bytes) required by LOWMC as used by the reference implementation, with the parameters used in Picnic at security levels L1, L3 and L5.

Params	Linear Matrices	Round Constants	Key Matrices	Total
L1	20,288	128	14,336	34,752
L1-full				18,786
L3	61,824	160	30,720	92,704
L3-full				41,568
L5	79,232	192	49,152	128,576
L5-full				73,281

Table 14: Size of constant data (in bytes) required by LOWMC as used by the optimized implementations, with the parameters used in Picnic at security levels L1, L3 and L5.

## References

- [ABF<sup>+</sup>17] Christopher Ambrose, Joppe W. Bos, Bjorn Fay, Marc Joye, Manfred Lochter, and Bruce Murray. Differential attacks on deterministic signatures. Cryptology ePrint Archive, Report 2017/975, 2017.
- [AGR<sup>+</sup>16] Martin R. Albrecht, Lorenzo Grassi, Christian Rechberger, Arnab Roy, and Tyge Tiessen. MiMC: Efficient encryption and cryptographic hashing with minimal multiplicative complexity. In *ASIACRYPT*, pages 191–219, 2016.
- [AOTZ20] Diego F. Aranha, Claudio Orlandi, Akira Takahashi, and Greg Zaverucha. Security of hedged Fiat-Shamir signatures under fault attacks. To appear at EUCROCRYPT, 2020. <https://eprint.iacr.org/2019/956>.
- [ARS<sup>+</sup>15a] Martin R. Albrecht, Christian Rechberger, Thomas Schneider, Tyge Tiessen, and Michael Zohner. Ciphers for MPC and FHE. In *EUROCRYPT*, 2015.
- [ARS<sup>+</sup>15b] Martin R. Albrecht, Christian Rechberger, Thomas Schneider, Tyge Tiessen, and Michael Zohner. Ciphers for MPC and FHE. In Elisabeth Oswald and Marc Fischlin, editors, *EUROCRYPT 2015, Part I*, volume 9056 of *LNCS*, pages 430–454. Springer, Heidelberg, April 2015.
- [ARS<sup>+</sup>16] Martin R. Albrecht, Christian Rechberger, Thomas Schneider, Tyge Tiessen, and Michael Zohner. Ciphers for MPC and FHE. *IACR Cryptology ePrint Archive*, 2016:687, 2016.
- [BB17] Gustavo Banegas and Daniel J. Bernstein. Low-communication parallel quantum multi-target preimage search. Cryptology ePrint Archive, Report 2017/789, 2017. <http://eprint.iacr.org/2017/789>.
- [BBDV20] Subhadeep Banik, Khashayar Barooti, F. Betül Durak, and Serge Vaudenay. Solving lowmc challenge. Manuscript, 2020. <https://github.com/lowmcchallenge/lowmcchallenge-material/tree/master/docs>.



- [BCG<sup>+</sup>12] Julia Borghoff, Anne Canteaut, Tim Güneysu, Elif Bilge Kavun, Miroslav Knezevic, Lars R. Knudsen, Gregor Leander, Ventzislav Nikov, Christof Paar, Christian Rechberger, Peter Rombouts, Søren S. Thomsen, and Tolga Yalçın. PRINCE - a low-latency block cipher for pervasive computing applications - extended abstract. In *ASIACRYPT*, 2012.
- [BCJ07] Gregory V. Bard, Nicolas T. Courtois, and Chris Jefferson. Efficient methods for conversion and solution of sparse systems of low-degree multivariate polynomials over  $\text{GF}(2)$  via SAT-Solvers. Cryptology ePrint Archive, Report 2007/024, 2007. <http://eprint.iacr.org/2007/024>.
- [BDP<sup>+</sup>18] Guido Bertoni, Joan Daemen, Michaël Peeters, Gilles Van Assche, Ronny Van Keer, and Benoît Viguier. KangarooTwelve: Fast hashing based on Keccak-p. In Bart Preneel and Frederik Vercauteren, editors, *ACNS 18*, volume 10892 of *LNCS*, pages 400–418. Springer, Heidelberg, July 2018.
- [Ber09] Daniel J. Bernstein. Cost analysis of hash collisions: Will quantum computers make SHARCS obsolete? 2009. <http://cr.yp.to/hash/collisioncost-20090823.pdf>.
- [BHT98] Gilles Brassard, Peter Høyer, and Alain Tapp. Quantum cryptanalysis of hash and claw-free functions. In Claudio L. Lucchesi and Arnaldo V. Moura, editors, *LATIN 1998*, volume 1380 of *LNCS*, pages 163–169. Springer, Heidelberg, April 1998.
- [Bih02] Eli Biham. How to decrypt or even substitute des-encrypted messages in  $2^{28}$  steps. *Inf. Process. Lett.*, 84(3):117–124, 2002.
- [Ble00] Daniel Bleichenbacher. On the generation of one-time keys in dl signature schemes. Presentation at IEEE P1363 Working Group meeting, November 2000. Unpublished., 2000.
- [BMP13] Joan Boyar, Philip Matthews, and René Peralta. Logic minimization techniques with applications to cryptology. *Journal of Cryptology*, 26(2):280–312, 2013.

- [BMS05] Alex Biryukov, Sourav Mukhopadhyay, and Palash Sarkar. Improved time-memory trade-offs with multiple data. In *Selected Areas in Cryptography, 12th International Workshop, SAC 2005, Kingston, ON, Canada, August 11-12, 2005, Revised Selected Papers*, pages 110–127, 2005.
- [BPW12] David Bernhard, Olivier Pereira, and Bogdan Warinschi. How not to prove yourself: Pitfalls of the Fiat-Shamir heuristic and applications to Helios. In Xiaoyun Wang and Kazue Sako, editors, *ASIACRYPT 2012*, volume 7658 of *LNCS*, pages 626–643. Springer, Heidelberg, December 2012.
- [BR93] Mihir Bellare and Phillip Rogaway. Random oracles are practical: A paradigm for designing efficient protocols. In *ACM CCS*, 1993.
- [CCF<sup>+</sup>16] Anne Canteaut, Sergiu Carpov, Caroline Fontaine, Tancrede Lepoint, María Naya-Plasencia, Pascal Paillier, and Renaud Sirdey. Stream ciphers: A practical solution for efficient homomorphic-ciphertext compression. In *FSE*, 2016.
- [CDG<sup>+</sup>17] Melissa Chase, David Derler, Steven Goldfeder, Claudio Orlandi, Sebastian Ramacher, Christian Rechberger, Daniel Slamanig, and Greg Zaverucha. Post-quantum zero-knowledge and signatures from symmetric-key primitives. In Bhavani M. Thuraisingham, David Evans, Tal Malkin, and Dongyan Xu, editors, *ACM CCS 2017*, pages 1825–1842. ACM Press, October / November 2017.
- [CDS94] Ronald Cramer, Ivan Damgård, and Berry Schoenmakers. Proofs of partial knowledge and simplified design of witness hiding protocols. In *CRYPTO*, 1994.
- [CGH98] Ran Canetti, Oded Goldreich, and Shai Halevi. The random oracle methodology, revisited (preliminary version). In *Proceedings of the Thirtieth Annual ACM Symposium on the Theory of Computing, Dallas, Texas, USA, May 23-26, 1998*, pages 209–218, 1998.

- [CGP<sup>+</sup>12] Claude Carlet, Louis Goubin, Emmanuel Prouff, Michaël Quisquater, and Matthieu Rivain. Higher-order masking schemes for s-boxes. In *FSE*, 2012.
- [Dam10] Ivan Damgård. On  $\Sigma$ -protocols. 2010. <http://www.cs.au.dk/~ivan/Sigma.pdf>.
- [dDOS19] Cyprien de Saint Guilhem, Lauren De Meyer, Emmanuela Orsini, and Nigel P. Smart. BBQ: Using AES in picnic signatures. In Kenneth G. Paterson and Douglas Stebila, editors, *SAC 2019*, volume 11959 of *LNCS*, pages 669–692. Springer, Heidelberg, August 2019.
- [DEM16] Christoph Dobraunig, Maria Eichlseder, and Florian Mendel. Higher-order cryptanalysis of LowMC. In Soonhak Kwon and Aaram Yun, editors, *ICISC 15*, volume 9558 of *LNCS*, pages 87–101. Springer, Heidelberg, November 2016.
- [DF16] Patrick Derbez and Pierre-Alain Fouque. Automatic search of meet-in-the-middle and impossible differential attacks. In Matthew Robshaw and Jonathan Katz, editors, *CRYPTO 2016, Part II*, volume 9815 of *LNCS*, pages 157–184. Springer, Heidelberg, August 2016.
- [DFMS19a] Jelle Don, Serge Fehr, Christian Majenz, and Christian Schaffner. Security of the Fiat-Shamir transformation in the quantum random-oracle model. In Alexandra Boldyreva and Daniele Micciancio, editors, *CRYPTO 2019, Part II*, volume 11693 of *LNCS*, pages 356–383. Springer, Heidelberg, August 2019.
- [DFMS19b] Jelle Don, Serge Fehr, Christian Majenz, and Christian Schaffner. Security of the Fiat-Shamir transformation in the quantum random-oracle model. Cryptology ePrint Archive, Report 2019/190, 2019. <https://eprint.iacr.org/2019/190>.
- [DKP<sup>+</sup>19a] Itai Dinur, Daniel Kales, Angela Promitzer, Sebastian Rasmacher, and Christian Rechberger. Linear Equivalence of Block Ciphers with Partial Non-Linear Layers: Application to LowMC. In *EUROCRYPT*, 2019. To appear.

- [DKP<sup>+</sup>19b] Itai Dinur, Daniel Kales, Angela Promitzer, Sebastian Rammacher, and Christian Rechberger. Linear equivalence of block ciphers with partial non-linear layers: Application to LowMC. In Yuval Ishai and Vincent Rijmen, editors, *EUROCRYPT 2019, Part I*, volume 11476 of *LNCS*, pages 343–372. Springer, Heidelberg, May 2019.
- [DLMW15] Itai Dinur, Yunwen Liu, Willi Meier, and Qingju Wang. Optimized interpolation attacks on LowMC. In Tetsu Iwata and Jung Hee Cheon, editors, *ASIACRYPT 2015, Part II*, volume 9453 of *LNCS*, pages 535–560. Springer, Heidelberg, November / December 2015.
- [DP08] Christophe De Cannière and Bart Preneel. Trivium. In *New Stream Cipher Designs - The eSTREAM Finalists*. 2008.
- [DPVAR00] Joan Daemen, Michaël Peeters, Gilles Van Assche, and Vincent Rijmen. Nessie proposal: Noekeon. In *First Open NESSIE Workshop*, 2000.
- [Fau02] Jean Charles Faugère. A new efficient algorithm for computing gröbner bases without reduction to zero (f5). In *ISSAC 2002*, pages 75–83. ACM, 2002.
- [Fis05] Marc Fischlin. Communication-efficient non-interactive proofs of knowledge with online extractors. In Victor Shoup, editor, *CRYPTO 2005*, volume 3621 of *LNCS*, pages 152–168. Springer, Heidelberg, August 2005.
- [FKMV12a] Sebastian Faust, Markulf Kohlweiss, Giorgia Azzurra Marson, and Daniele Venturi. On the non-malleability of the fiat-shamir transform. In *INDOCRYPT*, 2012.
- [FKMV12b] Sebastian Faust, Markulf Kohlweiss, Giorgia Azzurra Marson, and Daniele Venturi. On the non-malleability of the Fiat-Shamir transform. In Steven D. Galbraith and Mridul Nandi, editors, *INDOCRYPT 2012*, volume 7668 of *LNCS*, pages 60–79. Springer, Heidelberg, December 2012.

- [FS86] Amos Fiat and Adi Shamir. How to prove yourself: Practical solutions to identification and signature problems. In *CRYPTO*, pages 186–194, 1986.
- [GKRS20] Lorenzo Grassi, Daniel Kales, Christian Rechberger, and Markus Schofnegger. Survey of key-recovery attacks on lowmc in a single plaintext/ciphertext scenario. Manuscript, 2020. <https://github.com/lowmcchallenge/lowmcchallenge-material/tree/master/docs>.
- [GLSV14] Vincent Grosso, Gaëtan Leurent, François-Xavier Standaert, and Kerem Varici. Ls-designs: Bitslice encryption for efficient masked software implementations. In *FSE*, 2014.
- [GMO16a] Irene Giacomelli, Jesper Madsen, and Claudio Orlandi. ZKBoo: Faster zero-knowledge for Boolean circuits. Cryptology ePrint Archive, Report 2016/163, 2016. <http://eprint.iacr.org/2016/163>.
- [GMO16b] Irene Giacomelli, Jesper Madsen, and Claudio Orlandi. ZKBoo: Faster zero-knowledge for boolean circuits. In *USENIX Security*, 2016.
- [GMS02] Steven D. Galbraith, John Malone-Lee, and Nigel P. Smart. Public key signatures in the multi-user setting. *Inf. Process. Lett.*, 83(5):263–266, 2002.
- [GP99] Louis Goubin and Jacques Patarin. DES and differential power analysis (the “duplication” method). In Çetin Kaya Koç and Christof Paar, editors, *CHES’99*, volume 1717 of *LNCS*, pages 158–172. Springer, Heidelberg, August 1999.
- [Gra20] TU Graz. Lowmc cryptanalysis challenge, 2020. <https://lowmcchallenge.github.io/>.
- [Gro96] Lov K. Grover. A fast quantum mechanical algorithm for database search. In *STOC*, 1996.
- [GSE20] Tim Gellersen, Okan Seker, and Thomas Eisenbarth. Differential power analysis of the picnic signature scheme. Cryptology ePrint Archive, Report 2020/267, 2020. <https://eprint.iacr.org/2020/267>.

- [Hel80] Martin Hellman. A cryptanalytic time-memory trade-off. *IEEE transactions on Information Theory*, 26(4):401–406, 1980.
- [HS01] Nick Howgrave-Graham and Nigel P. Smart. Lattice attacks on digital signature schemes. *Des. Codes Cryptography*, 23(3):283–290, 2001.
- [IKOS09] Yuval Ishai, Eyal Kushilevitz, Rafail Ostrovsky, and Amit Sahai. Zero-knowledge proofs from secure multiparty computation. *SIAM Journal on Computing*, 39(3):1121–1152, 2009.
- [JNRV20] Samuel Jaques, Michael Naehrig, Martin Roetteler, and Fernando Virdia. Implementing grover oracles for quantum key search on AES and LowMC. In Anne Canteaut and Yuval Ishai, editors, *EUROCRYPT 2020, Part II*, volume 12106 of *LNCS*, pages 280–310. Springer, Heidelberg, May 2020.
- [Kat10] Jonathan Katz. *Digital Signatures*. Springer, 2010.
- [KKW18] Jonathan Katz, Vladimir Kolesnikov, and Xiao Wang. Improved non-interactive zero knowledge with applications to post-quantum signatures. In David Lie, Mohammad Mannan, Michael Backes, and XiaoFeng Wang, editors, *ACM CCS 2018*, pages 525–537. ACM Press, October 2018.
- [KLS18] Eike Kiltz, Vadim Lyubashevsky, and Christian Schaffner. A concrete treatment of Fiat-Shamir signatures in the quantum random-oracle model. In Jesper Buus Nielsen and Vincent Rijmen, editors, *EUROCRYPT 2018, Part III*, volume 10822 of *LNCS*, pages 552–586. Springer, Heidelberg, April / May 2018.
- [KM15] Neal Koblitz and Alfred J. Menezes. The random oracle model: a twenty-year retrospective. *Des. Codes Cryptography*, 77(2-3):587–610, 2015.
- [KRR<sup>+</sup>20] Daniel Kales, Sebastian Ramacher, Christian Rechberger, Roman Walch, and Mario Werner. Efficient FPGA implementations of LowMC and Picnic. In *CT-RSA: Cryptographers Track at the RSA Conference. Springer LNCS volume 12006*, pages 417–441, 2020. Source code: <https://github.com/IAIK/Picnic-FPGA>.

- [KZ20a] Daniel Kales and Greg Zaverucha. An attack on some signature schemes constructed from five-pass identification schemes. To appear at CANS 2020: 19th International Conference on Cryptology and Network Security, 2020. <https://eprint.iacr.org/2020/837>.
- [KZ20b] Daniel Kales and Greg Zaverucha. Improving the performance of the picnic signature scheme. IACR TCHES, Volume 2020, Issue 4., 2020. Please refer to the full version: <https://eprint.iacr.org/2020/427>.
- [LIM20] Fukang Liu, Takanori Isobe, and Willi Meier. Cryptanalysis of full lowmc and lowmc-m with algebraic techniques. Cryptology ePrint Archive, Report 2020/1034, 2020. <https://eprint.iacr.org/2020/1034>.
- [MJSC16] Pierrick Méaux, Anthony Journault, François-Xavier Standaert, and Claude Carlet. Towards stream ciphers for efficient FHE with low-noise ciphertexts. In *EUROCRYPT*, 2016.
- [MR18] Payman Mohassel and Peter Rindal. ABY<sup>3</sup>: A mixed protocol framework for machine learning. In David Lie, Mohammad Mannan, Michael Backes, and XiaoFeng Wang, editors, *ACM CCS 2018*, pages 35–52. ACM Press, October 2018.
- [MS04] Alfred Menezes and Nigel P. Smart. Security of signature schemes in a multi-user setting. *Des. Codes Cryptography*, 33(3):261–274, 2004.
- [MZ17] Payman Mohassel and Yupeng Zhang. SecureML: A system for scalable privacy-preserving machine learning. In *2017 IEEE Symposium on Security and Privacy*, pages 19–38. IEEE Computer Society Press, May 2017.
- [NIS15] NIST. SHA-3 Standard: Permutation-Based Hash and Extendable-Output Functions. National Institute of Standards and Technology (NIST), FIPS PUB 202, U.S. Department of Commerce, 2015.
- [NNL01] Dalit Naor, Moni Naor, and Jeffery Lotspiech. Revocation and tracing schemes for stateless receivers. In Joe Kilian, editor,

- CRYPTO 2001*, volume 2139 of *LNCS*, pages 41–62. Springer, Heidelberg, August 2001.
- [PSS<sup>+</sup>17] Damian Poddebniak, Juraj Somorovsky, Sebastian Schinzel, Manfred Lochter, and Paul Rosler. Attacking deterministic signature schemes using fault attacks. Cryptology ePrint Archive, Report 2017/1014, 2017.
  - [PST20] Christian Paquin, Douglas Stebila, and Goutam Tamvada. Benchmarking post-quantum cryptography in TLS. In Jintai Ding and Jean-Pierre Tillich, editors, *Post-Quantum Cryptography - 11th International Conference, PQCrypto 2020*, pages 72–91. Springer, Heidelberg, 2020.
  - [RST18] Christian Rechberger, Hadi Soleimany, and Tyge Tiessen. Cryptanalysis of low-data instances of full LowMCv2. *IACR Trans. Symm. Cryptol.*, 2018(3):163–181, 2018.
  - [SBWE20] Okan Seker, Sebastian Berndt, Luca Wilke, and Thomas Eisenbarth. Sni-in-the-head: Protecting mpc-in-the-head protocols against side-channel analysis. Cryptology ePrint Archive, Report 2020/544, 2020. <https://eprint.iacr.org/2020/544>. To appear at CCS’20.
  - [Tea19a] The Picnic Design Team. The Picnic signature algorithm specification, March 2019. Version 2.1, Available at <https://microsoft.github.io/Picnic/>.
  - [Tea19b] The Picnic Design Team. The Picnic website, March 2019. <https://microsoft.github.io/Picnic/>.
  - [Tie17] Tyge Tiessen. LowMC reference implementation, September 2017. Available at <https://github.com/LowMC/lowmc>, HEAD was 3994bc857661ac33134b36163b131a215f0fe9c3 when constants were generated.
  - [Unr12] Dominique Unruh. Quantum proofs of knowledge. In David Pointcheval and Thomas Johansson, editors, *EUROCRYPT 2012*, volume 7237 of *LNCS*, pages 135–152. Springer, Heidelberg, April 2012.



- [Unr15] Dominique Unruh. Non-interactive zero-knowledge proofs in the quantum random oracle model. In Elisabeth Oswald and Marc Fischlin, editors, *EUROCRYPT 2015, Part II*, volume 9057 of *LNCS*, pages 755–784. Springer, Heidelberg, April 2015.
- [Unr16] Dominique Unruh. Computationally binding quantum commitments. In *EUROCRYPT*, 2016.
- [vOW94] Paul C. van Oorschot and Michael J. Wiener. Parallel collision search with application to hash functions and discrete logarithms. In *CCS '94, Proceedings of the 2nd ACM Conference on Computer and Communications Security, Fairfax, Virginia, USA, November 2-4, 1994.*, pages 210–218, 1994.
- [Zaj17] Pavol Zajac. Upper bounds on the complexity of algebraic cryptanalysis of ciphers with a low multiplicative complexity. *Des. Codes Cryptogr.*, 82(1-2):43–56, 2017.

## A Change History

**Version 1.0** November 2017. Initial version, part of the first round submission to the NIST Post-Quantum Cryptography Standardization Process.

**Version 2.0** March 2019. Update for the second round submission to the NIST Post-Quantum Cryptography Standardization Process. The main change is to add the background material and security analysis of the Picnic2 parameter sets.

**Version 2.1** May 2019. Correct a typo in the statement of Theorem 6.2. The denominator of the 2nd term in the bound should be  $2^\kappa$  instead of  $2^{2\kappa}$ .

**Version 2.2** April 2020. Revise the QROM section (Section 6.4) of the Picnic2 analysis to include proof using results of [\[DFMS19a\]](#)

**Version 3.0** October 2020. Updated for Round 3 submission. Main changes are to replace Picnic2 with Picnic3, and add LowMC instances with a full S-box layer.